

RV: A Unified Region Vectorizer for LLVM

Simon Moll

Compiler Design Lab
Saarland Informatics Campus
Saarbrücken Germany



```

#pragma omp simd
for (int y = 0; y < height; ++y) {
  for (int x = 0; x < width; ++x) {
    complex<double> c = (startX+x*step) + (startY-y*step) * I;
    complex<double> z = 0.0;

    for (int n = 0; n < MAX_ITER; ++n) {
      z = z * z + c;
      if (hypot(z.real, z.imag) >= ESCAPE)
        break;
    }
    buffer[y][x] = colorMap(z);
  }
}

```

```

#pragma omp simd
for (int y = 0; y < height; ++y) {
  for (int x = 0; x < width; ++x) {
    complex<double> c = (startX+x*step) + (startY-y*step) * I;
    complex<double> z = 0.0;

    [
      for (int n = 0; n < MAX_ITER; ++n) {
        z = z * z + c;
        if (hypot(z.real, z.imag) >= ESCAPE)
          break;
      }
    ]
    buffer[y][x] = colorMap(z);
  }
}

```

```

#pragma omp simd
for (int y = 0; y < height; ++y) {
  for (int x = 0; x < width; ++x) {
    complex<double> c = (startX+x*step) + (startY-y*step) * I;
    complex<double> z = 0.0;

    [
      for (int n = 0; n < MAX_ITER; ++n) {
        z = z * z + c;
        if (hypot(z.real, z.imag) >= ESCAPE)
          break;
      }
      buffer[y][x] = colorMap(z);
    ]
  }
}

```

LoopVectorizer can not handle outer loops.

```

#pragma omp simd
for (int y = 0; y < height; ++y) {
  for (int x = 0; x < width; ++x) {
    complex<double> c = (startX+x*step) + (startY-y*step) * I;
    complex<double> z = 0.0;

    for (int n = 0; n < MAX_ITER; ++n) {
      z = z * z + c;
      if (hypot(z.real, z.imag) >= ESCAPE)
        break;
    }
    buffer[y][x] = colorMap(z);
  }
}

```

LoopVectorizer can not handle outer loops.

→ **RV can vectorize it.**

```
#pragma omp simd
for (int y = 0; y < height; ++y) {
    for (int x = 0; x < width; ++x) {
        complex<double> c = (startX+x*step) + (startY-y*step) * I;
        complex<double> z = 0.0;

        for (int n = 0; n < MAX_ITER; ++n) {
            z = z * z + c;
            if (hypot(z.real, z.imag) >= ESCAPE)
                break;
        }
        buffer[y][x] = colorMap(z);
    }
}
```

```
#pragma omp declare simd
float min (float a, float b)
{
    if (a < b) return a; else return b;
}
```

```
float min_v8 (<8 x float> a, <8 x float> b) {
    return select(a < b, a, b);
}
```

```
#pragma omp simd
for (int y = 0; y < height; ++y) {
  for (int x = 0; x < width; ++x) {
    complex<double> c = (startX+x*step) + (startY-y*step) * I;
    complex<double> z = 0.0;

    for (int n = 0; n < MAX_ITER; ++n) {
      z = z * z + c;
      if (hypot(z.real, z.imag) >= ESCAPE)
        break;
    }
    buffer[y][x] = colorMap(z);
  }
}
```

```
#pragma omp declare simd
float min (float a, float b)
{
  if (a < b) return a; else return b;
}
```

```
float min_v8 (<8 x float> a, <8 x float> b) {
  return select(a < b, a, b);
}
```

```
#pragma omp simd
for (int y = 0; y < height; ++y) {
  for (int x = 0; x < width; ++x) {
    complex<double> c = (startX+x*step) + (startY-y*step) * I;
    complex<double> z = 0.0;

    for (int n = 0; n < MAX_ITER; ++n) {
      z = z * z + c;
      if (hypot(z.real, z.imag) >= ESCAPE)
        break;
    }
    buffer[y][x] = colorMap(z);
  }
}
```

```
#pragma omp declare simd
float min (float a, float b)
{
  if (a < b) return a; else return b;
}
```

```
float min_v8 (<8 x float> a, <8 x float> b) {
  return select(a < b, a, b);
}
```



```
#pragma omp simd
for (int y = 0; y < height; ++y) {
  for (int x = 0; x < width; ++x) {
    complex<double> c = (startX+x*step) + (startY-y*step) * I;
    complex<double> z = 0.0;




    for (int n = 0; n < MAX_ITER; ++n) {
      z = z * z + c;
      if (hypot(z.real, z.imag) >= ESCAPE)
        break;
    }
    buffer[y][x] = colorMap(z);
  }
}
```

```
#pragma omp declare simd
float min (float a, float b)
{
  if (a < b) return a; else return b;
}
```






```
float min_v8 (<8 x float> a, <8 x float> b) {
  return select(a < b, a, b);
}
```


■ `rv::analyze(rv::Region& , ...)`






■ `rv::analyze(rv::Region& , ...)`



■ `rv::linearize`

- ▶ *If-Conversion / Loop predication.*
- ▶ *Preserves uniform control.*

■ `rv::analyze(rv::Region& , ..)`



■ `rv::linearize`

- ▶ *If-Conversion / Loop predication.*
- ▶ *Preserves uniform control.*

■ `rv::vectorize`

Vector IR generation.

RV: A Unified Region Vectorizer for LLVM

Simon Moll / Saarland University / Saarland Informatics Campus

Introduction

The Region Vectorizer provides a single, unified API to vectorize code regions.

- RV is a generalization of the Whole-Function Vectorizer
R. Karrenberg, S. Hack, "Whole Function Vectorization" (CGO '11)

Applications

- Outer-Loop Vectorizer** An "unroll-and-jam" vectorizer based on RV's analysis and transformations
- pragma omp simd** Emit vector code for SIMD regions right from Clang
- Vectorizer Cost Model** How much predication? Which memory accesses vectorize well?
- Polly** Directly vectorize loops during Polly code generation
- PIR** Parallel region vectorizer

```
rv::VectorizationInfo vi;
// region set up
rv::Region R(xLoop);
vi.setVectorShape(rfn,
                  VectorShape::consecutive());

// Vectorization analysis
rv::analyze(R, vi, domTree, loopInfo);

// Control conversion
rv::linearize(R, vi, domTree, loopInfo);

// Vector IR generation
rv::vectorize(R, vi, domTree);
```

rv::Region Region

A **region** can be a subset of the basic blocks in a function or an entire function (omp declare simd).

```
#pragma omp simd
for (int x = 0; x < width; ++x) {
    for (int y = 0; y < height; ++y) {
        complex<double> c = (startX+x*step) + (startY-y*step) * I;
        complex<double> z = 0.0;

        for (int n = 0; n < MAX_ITER; ++n) {
            z = z * z + c;
            if (hypot(z.real, z.imag) >= ESCAPE)
                break;
        }
        buffer[y][x] = colorMap(z);
    }
}
```

divergent loop

```
#pragma omp declare simd
float min (float a, float b)
{
    if (a < b) return a; else return b;
}

float min_v8 (<8 x float> a, <8 x float> b) {
    return select(a < b, a, b);
}
```

rv::analyze Vectorization Analysis

(stride alignment) or T