



# Scalable, Robust, and “Regression-Free” Loop Optimizations for Scientific Fortran and Modern C++

Michael Kruse, Tobias Grosser

Albert Cohen, Sven Verdoolaege,  
Oleksandre Zinenko  
Polly Labs, ENS Paris

Johannes Doerfert  
Uni. Saarbruecken


Siddharth Bhat  
IIIT Hyderabad (Intern at ETH)

Roman Gereev,  
Ural Federal University


Hongin Zheng, Alexandre Isonard  
Xilinx

Swiss Universities / PASC  
Qualcomm, ARM, Xilinx

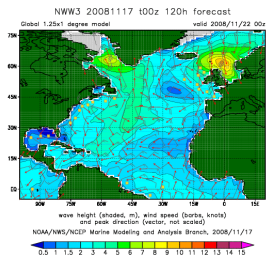
... many others






# Physics Simulations




# Weather



# Graphics



# Machine Learning



# COSMO: Weather and Climate Model

- 
- 500.000 Lines of Fortran
  - 18.000 Loops
  - 19 Years of Knowledge
- 
- Used in Switzerland, Russia, Germany, Poland, Italy, Israel, Greece, Romania, ...

# COSMO – Climate Modeling



Piz Daint, Lugano, Switzerland




Rank	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.450GHz, Sunway . NRCCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371
2	Tianhe-2 (MilkyWay-2) - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P, NUOT National Super Computer Center in Guangzhou China	3,120,000	33,862.7	54,902.4	17,808
3	Piz Daint - Cray XC50, Xeon E5-2690v3 12C 2.60GHz, Aries interconnect, NVIDIA Tesla P100, Cray Inc. Swiss National Supercomputing Centre (CSCS) Switzerland	361,760	19,590.0	25,326.3	2,272
4	Titan - Cray XK7, Opteron 6274 16C 2.200GHz, Cray Gemini	560,640	17,590.0	27,112.5	8,209

- Global (low-resolution Model)
- Up to 5000 Nodes
- Run ~monthly

# COSMO – Weather Forecast

## The range of the COSMO models



- Regional model
- High-resolution
- Runs hourly  
(20 instances in parallel)

- Today: 40 Nodes \* 8 GPU
- Manual Translation to GPUs
- 3 year multi-person project

Can LLVM do  
this automatically?



## Statistics - COSMO

### ■ Number of Loops

- 18,093 Total
- 9,760 Static Control Loops (Modeled precisely by Polly)
- 15,245 Non-Affine Memory Accesses (Approximated by Polly)

- 11.154 Loops after precise modeling, less e.g. due to:
  - *Infeasible assumptions taken, or modeling timeouts*

### ■ Largest set of loops: 72 loops

### ■ Reasons why loops cannot be modeled

- Function calls with side-effects
- Uncomputable loops bounds (data-dependent loop bounds?)



Siddharth Bhat

# Interprocedural Loop Interchange for GPU Execution

Pulled out parallel loop for  
OpenACC Annotations

```
#ifdef _OPENACC
    !$acc parallel
    !$acc loop gang vector
    DO j1 = ki1sc, ki1ec
        CALL coe_th_gpu(pduh2oc (j1, ki3sc), pduh2of(j1, ki3sc), pduco2(j1, ki3sc),
                        pduo3(j1, ki3sc), ..., pa2f(j1), pa3c(j1), pa3f(j1))
    ENDDO
    !$acc end parallel
#else
    CALL coe_th (pduh2oc, pduh2of, pduco2, pduo3, palogp, palogt, podsc, podsf, podac, podaf,
                ..., pa3c, pa3f)
#endif
```



# Optical Effect on Solar Layer

Outer loop is sequential

```
DO j3 = ki3sc+1, ki3ec
  CALL coe_th (j3) { ! Determine effect of the layer in *coe_th*
```

```
! Optical depth of gases
```

```
DO j1 = ki1sc, ki1ec
```

Inner loop is parallel

```
...
IF (kco2 /= 0) THEN
```

```
  zodgf = zodgf + pduco2(j1 ,j3)* (cobi(kco2,kspec,2)* EXP ( coali(kco2,kspec,2) *
    palogp(j1 ,j3) -cobti(kco2,kspec,2) * palogt(j1 ,j3)))
```

```
ENDIF
```

```
...
zeps=SQRT(zodgf*zodgf)
```

## Sequential Dependences

```
...
ENDDO
```

```
}
```

Inner loop is parallel

```
DO j1 = ki1sc, ki1ec ! Set RHS
```

```
...
```

```
ENDDO
```

```
DO j1 = ki1sc, ki1ec ! Elimination and storage of utility variables
```

Inner loop is parallel

```
...
```

```
ENDDO
```

```
ENDDO ! End of vertical loop over layers
```

# Optical Effect on Solar Layer – After interchange

```

!> Turn loop structure with multiple ip loops inside a
!> single k loop into perfectly nested k-in loop on GPU
#ifdef _OPENACC
!$acc parallel
!$acc loop gang vector
DO j1 = ki1sc, ki1ec
!$acc loop seq
DO j3 = ki3sc+1, ki3ec      ! Loop over vertical

! Determine effects of layer in *coe_so*
CALL coe_so_gpu(pduh2oc (j1,j3) , pduh2of (j1,j3) , ..., pa4c(j1), pa4f(j1), pa5c(j1), pa5f(j1))

! Elimination
...
ztd1 = 1.0_dp/(1.0_dp-pa5f(j1)*(pca2(j1,j3)*ztu6(j1,j3-1)+pcc2(j1,j3)*ztu8(j1,j3-1)))
ztu9(j1,j3) = pa5c(j1)*pcd1(j1,j3)+ztd6*ztu3(j1,j3) + ztd7*ztu5(j1,j3)
ENDDO
END DO      ! Vertical loop
!$acc end parallel

```

Outer loop is parallel

Inner loop is sequential

# Life Range Reordering (IMPACT'16 Verdoolaege)

Privatization needed  
for parallel execution

False dependences  
prevent interchange

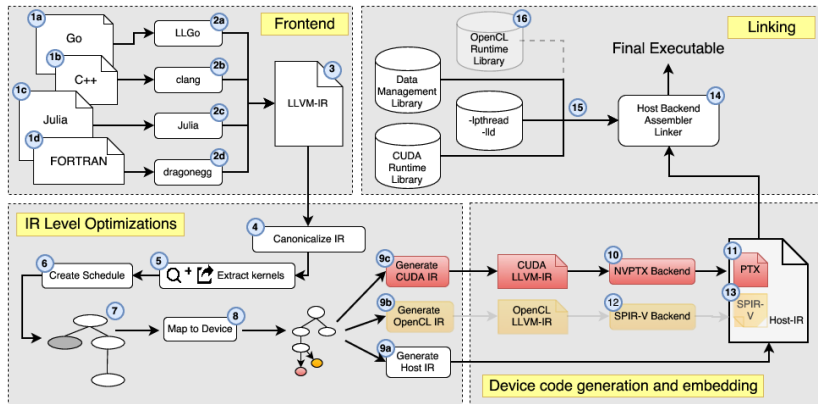
Scalable  
Scheduling

sequential

parallel

parallel  
sequential

# Polly-ACC: Architecture



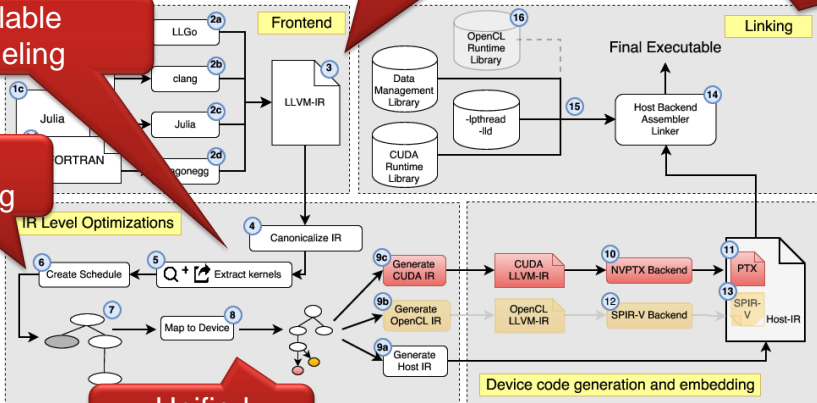
# Polly-ACC: Architecture

Intrinsics to model  
Multi-dimensional  
strided arrays

Better ways to link  
with NVIDIA  
libdevice

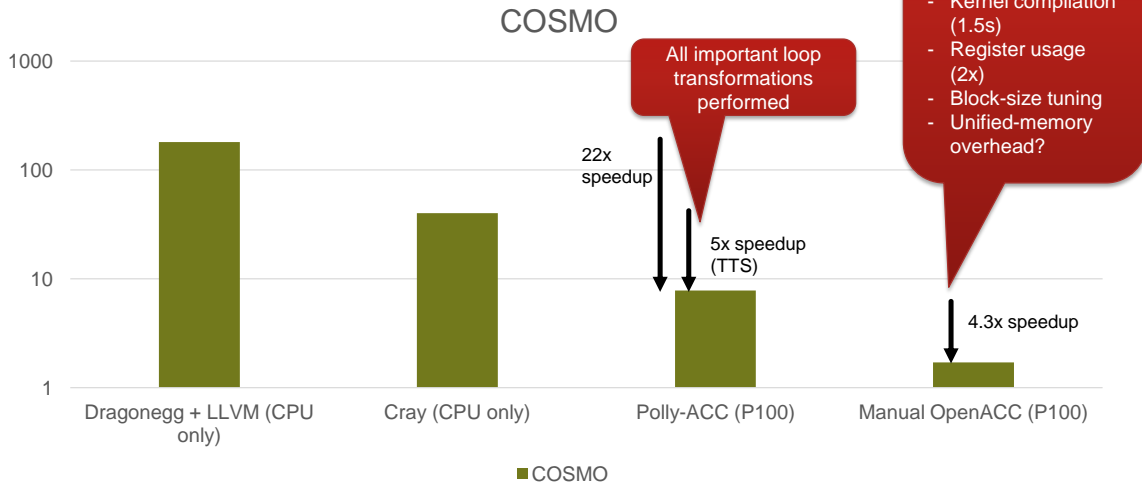
Scalable  
Modeling

Scalable  
Scheduling



Unified  
Memory

# Performance



# Expression Templates (in a nutshell)

```
class Vec : public VecExpression<Vec> {
    std::vector<double> elems;

public:
    double operator[](size_t i) const { return elems[i]; }
    double &operator[](size_t i)      { return elems[i]; }
    size_t size() const               { return elems.size(); }

    Vec(size_t n) : elems(n) {}

    // A Vec can be constructed from any VecExpression, forcing its evaluation.
    template <typename E>
    Vec(VecExpression<E> const& vec) : elems(vec.size()) {
        for (size_t i = 0; i != vec.size(); ++i) {
            elems[i] = vec[i];
        }
    }
};
```

# Expression Templates (in a nutshell)

```
class Vec : public VecExpression<Vec> {
    std::vector<double> elems;

public:
    double operator[](size_t i) const { return elems[i]; }
    double &operator[](size_t i)      { return elems[i]; }
    size_t size() const               { return elems.size(); }

    Vec(size_t n) : elems(n) {}

    // A Vec can be constructed from any VecExpression, forcing its evaluation.
    template <typename E>
    Vec(VecExpression<E> const& vec) : elems(vec.size()) {
        for (size_t i = 0; i != vec.size(); ++i) {
            elems[i] = vec[i];
        }
    }
};
```



## Expression Templates (in a nutshell) - II

```
template <typename E1, typename E2>
class VecSum : public VecExpression<VecSum<E1, E2> > {
    E1 const& _u;
    E2 const& _v;

public:
    VecSum(E1 const& u, E2 const& v) : _u(u), _v(v) {
        assert(u.size() == v.size());
    }

    double operator[](size_t i) const { return _u[i] + _v[i]; }
    size_t size() const { return _v.size(); }
};

template <typename E1, typename E2>
VecSum<E1,E2> const
operator+(E1 const& u, E2 const& v) {
    return VecSum<E1, E2>(u, v);
}
```

```
Vec a, b, c;
```

```
auto Sum = a + b + c;
```

```
assert(sizeof(sum) ==
VecSum<VecSum<Vec, Vec>, Vec>)
```

```
// evaluation only happens on
// assignment to type Vec
```

```
Vec evaluate = Sum;
```

# “Modern C++” -- boost::ublas and Expression Templates

## 1. Detect operations on tropical semi-rings

- SGEMM/DGEMM (+, \*)
- Floyd-Warshall (min, +)

Data-Layout  
Transformations in Polly

## 2. Apply GOTO Algorithm (1)

- L2 Tiling
- Cache Transposed Submatrices
- Register Tiling

## 3. Chose optimal Cache and Register Tile Sizes (2)

TargetData:

- L1/L2 Cache Sizes
- L2/L2 Cache Latencies

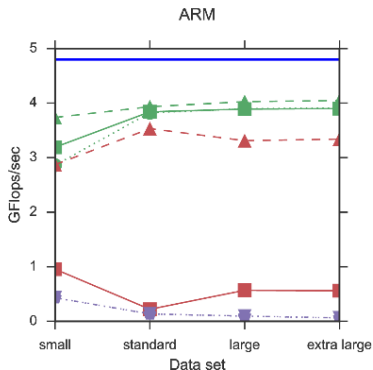
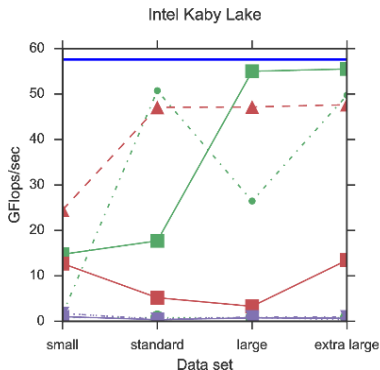


Roman Gareev

(1) High-performance implementation of the level-3 BLAS, Goto et al

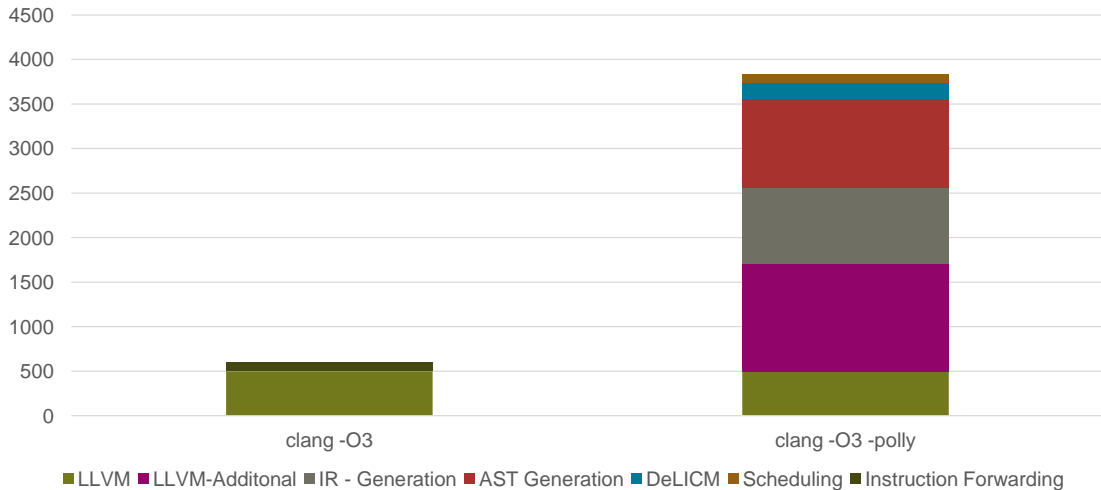
(2) A Analytical Modeling is Enough for High Performance BLIS, Tzem Low et al

# DGEMM Performance




Thanks  
@ARMHPC  
(Florian Hahn)  
for ARM codegen  
improvements!

## 3MM Compile Time



# “Provable” Correct Types for Loop Transformations

```
for (int32 i = 1; i < N; i++)  
  for (int32 j = 1; j <= M; j++)  
    A(i,j) = A(i-1,j) + A(i,j-1)
```





Maximilian  
Falkenstein

# “Provable” Correct Types for Loop Transformations

```
for (int32 i = 1; i < N; i++)  
  for (int32 j = 1; j <= M; j++)  
    A(i,j) = A(i-1,j) + A(i,j-1)
```



```
for (intX c = 2; c < N+M; c++)  
  #pragma simd  
  for (intX i = max(1, c-M); i <= min(N, c-1); i++)  
    A(i,c-i) = A(i-1,c-1) + A(i,c-i-1)
```





Maximilian  
Falkenstein

# “Provable” Correct Types for Loop Transformations

```
for (int32 i = 1; i < N; i++)  
  for (int32 j = 1; j <= M; j++)  
    A(i,j) = A(i-1,j) + A(i,j-1)
```

What is X?

```
for (intX c = 2; c < N+M; c++)  
  #pragma simd  
  for (intX i = max(1, c-M); i <= min(N, c-1); i++)  
    A(i,c-i) = A(i-1,c-1) + A(i,c-i-1)
```



Maximilian  
Falkenstein

## Precise Solution

```
for (intX c = 2; c < N+M; c++)  
  # simd  
  for (intX i = max(1, c-M); i <= min(N, c-1); i++)  
    A(i, c-i) = A(i-1, c-1) + A(i, c-i-1)
```


**Domain:**  $\{ [c] : 2 \leq c < N + M$   
 $\text{INT\_MIN} \leq N, M \leq \text{INT\_MAX} \}$

$f0() = c - i$


$f1() = c - i - 1$

1) **calc:**  $\min(fX()), \max(fX())$  under **Domain**

2) **choose type accordingly**







- Do you still target 32 bit?
- GPUs are faster in 32 bit
- FPGA?!

### ILP Solver

- Minimal Types
- Potentially Costly

### Approximations\*


- $s(a+b) \leq \max(s(a), s(b)) + 1$
- Good, if smaller than native type

\* Earlier uses in GCC and Polly

### Preconditions


- Assume values fit into 32 bit
- Derive required pre-conditions

# Type Distribution for LNT SCOPS




32 + epsilon is  
almost always  
enough!

# Compile Time Overhead

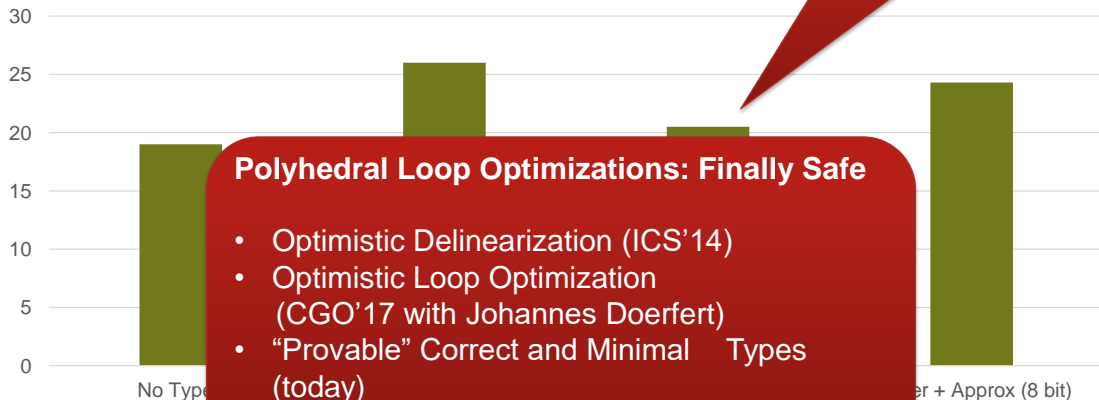


# Compile Time Overhead



# Compile Time Overhead

GPU Code Generation (5000 lines of code)



# Scalar Dependencies

## Virtual Registers and PHI-Nodes

```
    for (int i=0; i<N; i++) {  
S:    A[i] = ...;  
T:    ... = A[i];  
    }
```

- $T(i)$  depends on  $S(i)$  Read-After-Write/Flow-dependency
- $S(0), S(1), \dots, S(N-1), T(0), T(1), \dots$  is a valid execution
- Parallel loop (OpenMP, OpenCL/PTX, tiling, vectorization, etc.)

# Scalar Dependencies

Virtual Registers and PHI-Nodes

```
    for (int i=0; i<N; i++) {  
S:    A[i] = ...;  
T:    ... = A[i];  
    }
```



```
    for (int i=0; i<N; i++) {  
S:    tmp = ...;  
T:    ... = tmp;  
    }
```

“0-dimensional array”

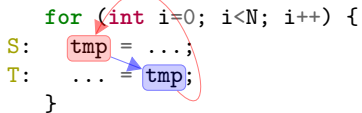
# Scalar Dependencies

## Virtual Registers and PHI-Nodes

```
for (int i=0; i<N; i++) {  
S:   A[i] = ...;  
T:   ... = A[i];  
}
```



```
for (int i=0; i<N; i++) {  
S:   tmp = ...;  
T:   ... = tmp;  
}
```



“0-dimensional array”

- $S(i)$  “depends” on  $S(i-1)$
- $S(i)$  “depends” on  $T(i-1)$
- $S(0), S(1), \dots, T(0), \dots$  is **no** valid execution

Write-After-Write/Output-dependency

Write-After-Read/Anti-dependency



# Scalar Dependencies

Virtual Registers and PHI-Nodes

```
for (int i=0; i<N; i++) {  
S:   A[i] = ...;  
T:   ... = A[i];  
}
```



mem2reg/SROA

```
for (int i=0; i<N; i++) {  
S:   tmp = ...;  
T:   ... = tmp;  
}
```

“0-dimensional array”

- $S(i)$  “depends” on  $S(i-1)$
- $S(i)$  “depends” on  $T(i-1)$
- $S(0), S(1), \dots, T(0), \dots$  is **no** valid execution

Write-After-Write/Output-dependency

Write-After-Read/Anti-dependency

# Loop-Invariant Code Motion

```
for (int i = 0; i < N; i += 1)
  for (int k = 0; k < K; k += 1)
S:    C[i] += A[i] * B[k];
```

# Loop-Invariant Code Motion

```
for (int i = 0; i < N; i += 1)
    for (int k = 0; k < K; k += 1)
S:    C[i] += A[i] * B[k];
```



```
for (int i = 0; i < N; i += 1) {
T:    double tmp = A[i];
    for (int k = 0; k < K; k += 1)
S:    C[i] += tmp * B[k];
}
```

# Loop-Invariant Code Motion

```
for (int i = 0; i < N; i += 1)
  for (int k = 0; k < K; k += 1)
S:    C[i] += A[i] * B[k];
```



GVN/LICM

```
for (int i = 0; i < N; i += 1) {
T:   double tmp = A[i];
    for (int k = 0; k < K; k += 1)
S:    C[i] += tmp * B[k];
}
```

The diagram illustrates the transformation of the code. A red oval highlights the expression `A[i]` in the assignment `double tmp = A[i];` and the variable `tmp` in the expression `C[i] += tmp * B[k];`. A red arrow points from the `A[i]` to the `tmp` variable, indicating that the value of `A[i]` is stored in `tmp` and then used in the subsequent calculation, which is a form of Loop-Invariant Code Motion (LICM).

## Scalar Promotion in Loops

```
    for (int i = 0; i < N; i += 1) {  
T:    C[i] = 0;  
        for (int k = 0; k < K; k += 1)  
S:    C[i] += A[i][k];  
    }
```

## Scalar Promotion in Loops

```
for (int i = 0; i < N; i += 1) {  
T:   C[i] = 0;  
    for (int k = 0; k < K; k += 1)  
S:   C[i] += A[i][k];  
}
```



```
for (int i = 0; i < N; i += 1) {  
T:   double tmp = 0;  
    for (int k = 0; k < K; k += 1)  
S:   tmp += A[i][k];  
U:   C[i] = tmp;  
}
```

Diagram illustrating scalar promotion in loops. The variable `tmp` is introduced and used to store the result of the inner loop's accumulation, avoiding repeated updates to `C[i]` and promoting the scalar `C[i]` to a double.

## Scalar Promotion in Loops

```
    for (int i = 0; i < N; i += 1) {  
T:    C[i] = 0;  
        for (int k = 0; k < K; k += 1)  
S:    C[i] += A[i][k];  
    }
```



LICM

```
    for (int i = 0; i < N; i += 1) {  
T:    double tmp = 0;  
        for (int k = 0; k < K; k += 1)  
S:    tmp += A[i][k];  
U:    C[i] = tmp;  
    }
```

The diagram illustrates the transformation of the original code using Loop-Invariant Code Motion (LICM). In the transformed code, the variable `tmp` is declared as a `double` and initialized to 0 at the start of the inner loop. The accumulation of values from `A[i][k]` is performed on `tmp` instead of directly on `C[i]`. Finally, the value of `tmp` is assigned to `C[i]` once the inner loop completes. Red and blue boxes highlight the `tmp` variable and its usage, with arrows showing the flow of data between the assignment, the accumulation, and the final store.

# Speculative Execution

```
    for (int i = 0; i < N; i += 1) {  
        if (i > 5)  
S1:      C[i] = 5 + 2*x;  
        else  
S2:      C[i] = 7 + 2*x;  
    }
```



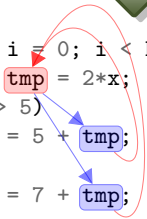
# Speculative Execution

```
for (int i = 0; i < N; i += 1) {  
    if (i > 5)  
S1:    C[i] = 5 + 2*x;  
    else  
S2:    C[i] = 7 + 2*x;  
}
```



EarlyCSE/GVN/NewGVN/GVNHoist

```
for (int i = 0; i < N; i += 1) {  
T:    double tmp = 2*x;  
    if (i > 5)  
S1:    C[i] = 5 + tmp;  
    else  
S2:    C[i] = 7 + tmp;  
}
```



## (Partial) Redundancy Elimination

```
    for (int i = 0; i < N; i += 1) {  
T:    C[i] = 0;  
        for (int k = 0; k < K; k += 1)  
S:    C[i] += A[i][k];  
    }
```

## (Partial) Redundancy Elimination

```
for (int i = 0; i < N; i += 1) {  
T:   C[i] = 0;  
    for (int k = 0; k < K; k += 1)  
S:   C[i] += A[i][k];  
}
```



GVN Load PRE

```
for (int i = 0; i < N; i += 1) {  
T:   double tmp = 0;  
    for (int k = 0; k < K; k += 1)  
S:   C[i] = (tmp += A[i][k]);  
}
```

# Loop Idiom

“doitgen” – Multiresolution Kernel from MADNESS

```
for (int r = 0; r < R; r++)
  for (int q = 0; q < Q; q++) {
    for (int p = 0; p < P; p++) {
      sum[p] = 0;
      for (int s = 0; s < P; s++)
        sum[p] += A[r][q][s] * C4[s][p];
    }
    for (int p = 0; p < P; p++)
      A[r][q][p] = sum[p];
  }
```

# Loop Idiom

“doitgen” – Multiresolution Kernel from MADNESS

```
for (int r = 0; r < R; r++)  
  for (int q = 0; q < Q; q++) {  
    for (int p = 0; p < P; p++) {  
      sum[p] = 0;  
      for (int s = 0; s < P; s++)  
        sum[p] += A[r][q][s] * C4[s][p];  
    }  
    for (int p = 0; p < P; p++)  
      A[r][q][p] = sum[p];  
  }
```




LoopIdiom

```
for (int r = 0; r < R; r++)  
  for (int q = 0; q < Q; q++) {  
    for (int p = 0; p < P; p++) {  
      sum[p] = 0;  
      for (int s = 0; s < P; s++)  
        sum[p] += A[r][q][s] * C4[s][p];  
    }  
    memcpy(A[r][q], sum, sizeof(sum[i]) * P);  
  }
```

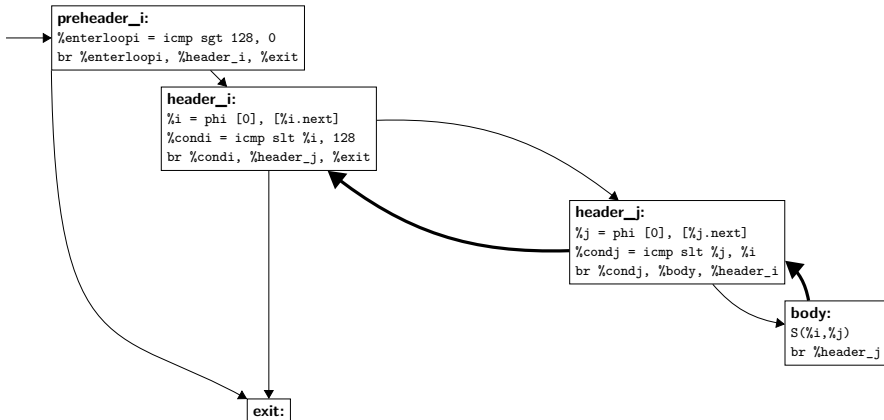
# Loop Rotation

```
for (int i = 0; i < 128; i += 1)
  for (int j = 0; j < i; j += 1)
    S(i,j);
```



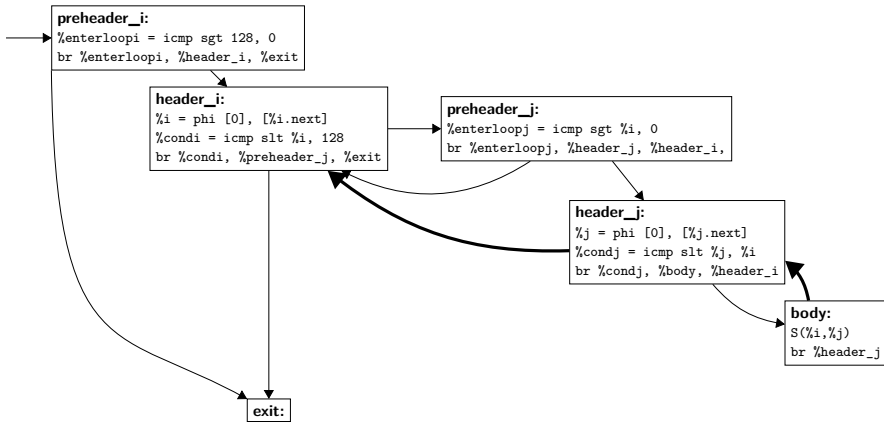
# Loop Rotation

```
for (int i = 0; i < 128; i += 1)
  for (int j = 0; j < i; j += 1)
    S(i,j);
```



# Loop Rotation

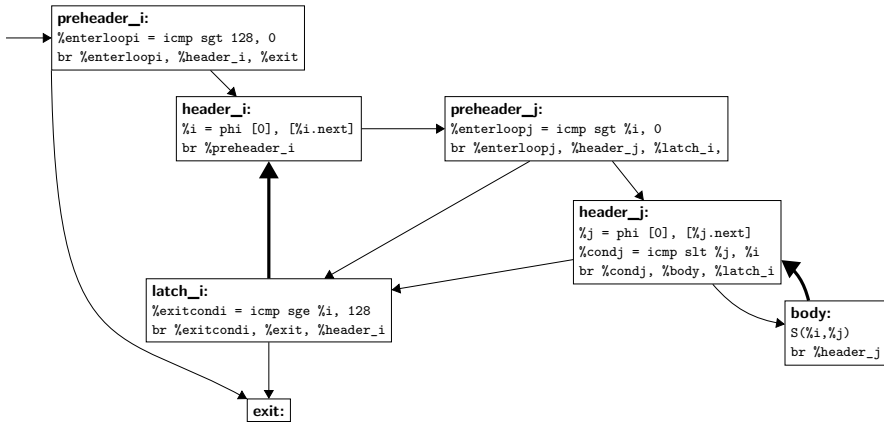
```
for (int i = 0; i < 128; i += 1)
  for (int j = 0; j < i; j += 1)
    S(i,j);
```





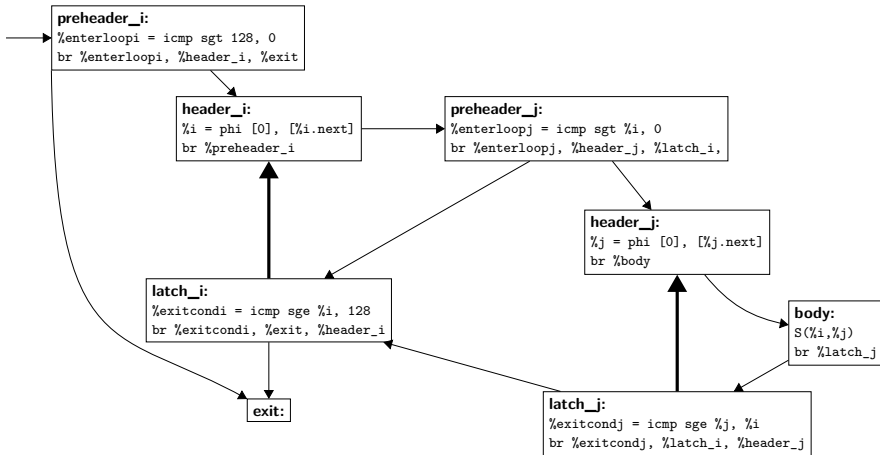
# Loop Rotation

```
for (int i = 0; i < 128; i += 1)
  for (int j = 0; j < i; j += 1)
    S(i,j);
```



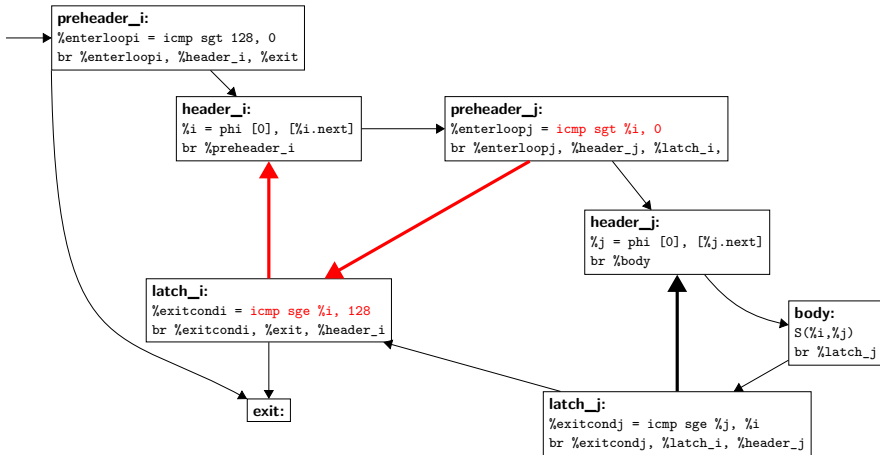
# Loop Rotation

```
for (int i = 0; i < 128; i += 1)
  for (int j = 0; j < i; j += 1)
    S(i,j);
```



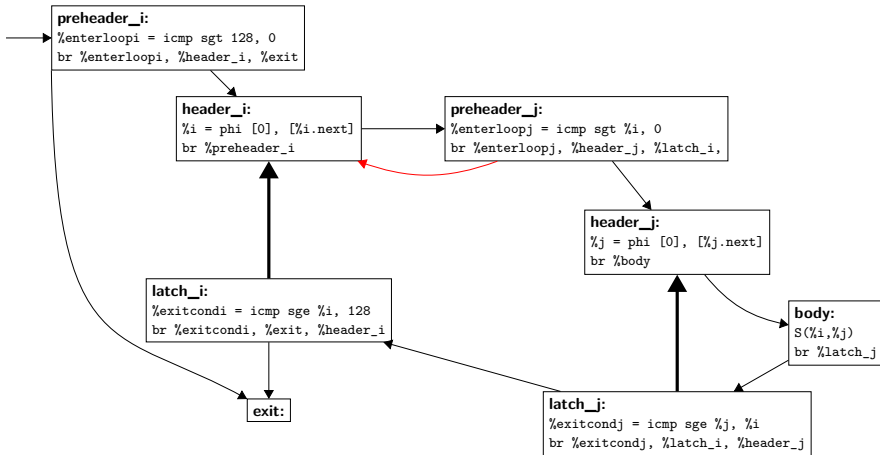
# Jump Threading

```
for (int i = 0; i < 128; i += 1)
  for (int j = 0; j < i; j += 1)
    S(i,j);
```



# Jump Threading


```
for (int i = 0; i < 128; i += 1)
  for (int j = 0; j < i; j += 1)
    S(i,j);
```



# Chapter Summary

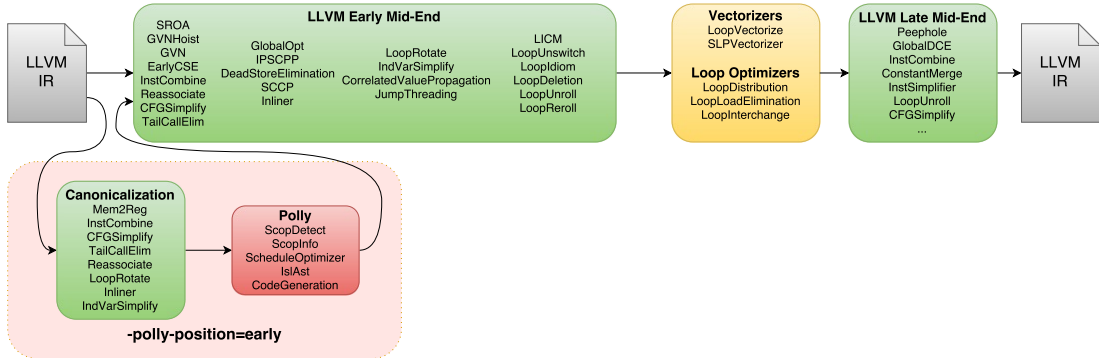
- Semantically identical IR can be harder to optimize
- Possible causes:
  - Static Single Assignment form (e.g. mem2reg)
  - Non-Polyhedral transformation passes (e.g. GVN, LICM)
  - C++ abstraction layers (e.g. Boost uBLAS)
  - Manual source optimizations (e.g. loop hoisting)
  - Code generators (e.g. TensorFlow XLA)

# LLVM Pass Pipeline -03



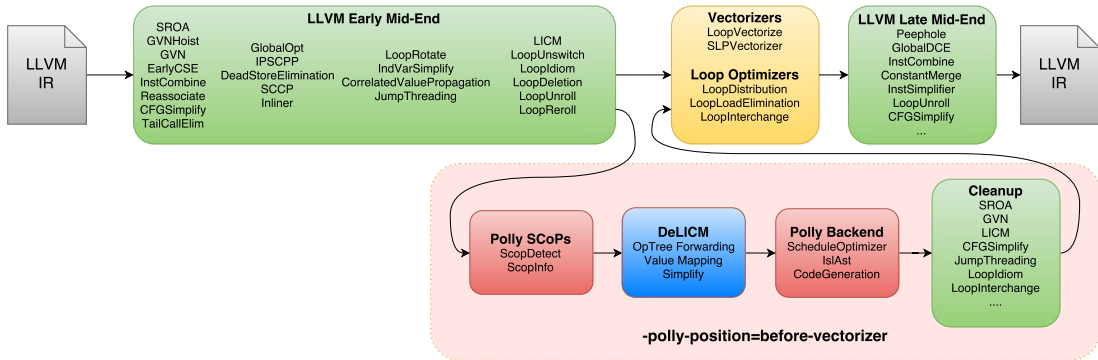
# LLVM Pass Pipeline

-O3 -polly -polly-position=early




# LLVM Pass Pipeline

-O3 -polly -polly-position=before-vectorizer





# Effects of GVN, LICM, Loopldiom



# Value Mapping

“DeLICM”

```
double c;  
for (int i = 0; i < 3; i += 1) {  
T:   c = 0;  
    for (int k = 0; k < 3; k += 1)  
S:   c += A[i] * B[k];  
U:   C[i] = c;  
}
```

# Value Mapping

“DeLICM”

```

double c;
for (int i = 0; i < 3; i += 1) {
T:   c = 0;
    for (int k = 0; k < 3; k += 1)
S:   c += A[i] * B[k];
U:   C[i] = c;
}
```

	$T(0)$	$S(0,0)$	$S(0,1)$	$S(0,2)$	$U(0)$	$T(1)$	$S(1,0)$	$S(1,1)$	$S(1,2)$	$U(1)$	$T(2)$	$S(2,0)$	$S(2,1)$	$S(2,2)$	$U(2)$
c	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

# Value Mapping

“DeLICM”

```

double c;
for (int i = 0; i < 3; i += 1) {
T:   c = 0;
    for (int k = 0; k < 3; k += 1)
S:   c += A[i] * B[k];
U:   C[i] = c;
}
```

	$T(0)$	$S(0,0)$	$S(0,1)$	$S(0,2)$	$U(0)$	$T(1)$	$S(1,0)$	$S(1,1)$	$S(1,2)$	$U(1)$	$T(2)$	$S(2,0)$	$S(2,1)$	$S(2,2)$	$U(2)$
c	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

# Value Mapping

“DeLICM”

```

double c;
for (int i = 0; i < 3; i += 1) {
T:   c = 0;
    for (int k = 0; k < 3; k += 1)
S:   c += A[i] * B[k];
U:   C[i] = c;
}
```

	$T(0)$	$S(0,0)$	$S(0,1)$	$S(0,2)$	$U(0)$	$T(1)$	$S(1,0)$	$S(1,1)$	$S(1,2)$	$U(1)$	$T(2)$	$S(2,0)$	$S(2,1)$	$S(2,2)$	$U(2)$
c	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

C[0]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[1]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[2]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

# Value Mapping

“DeLICM”

```

double c;
for (int i = 0; i < 3; i += 1) {
T:   c = 0;
    for (int k = 0; k < 3; k += 1)
S:   c += A[i] * B[k];
U:   C[i] = c;
}
```

	T(0)	S(0,0)	S(0,1)	S(0,2)	U(0)	T(1)	S(1,0)	S(1,1)	S(1,2)	U(1)	T(2)	S(2,0)	S(2,1)	S(2,2)	U(2)
c	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

C[0]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[1]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[2]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

# Value Mapping

“DeLICM”

```

double c;
for (int i = 0; i < 3; i += 1) {
T:   C[2] = 0;
    for (int k = 0; k < 3; k += 1)
S:   C[2] += A[i] * B[k];
U:   C[i] = C[2];
}
  
```

	T(0)	S(0,0)	S(0,1)	S(0,2)	U(0)	T(1)	S(1,0)	S(1,1)	S(1,2)	U(1)	T(2)	S(2,0)	S(2,1)	S(2,2)	U(2)
c	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

C[0]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[1]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[2]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

# Value Mapping

“DeLICM”

```

double c;
for (int i = 0; i < 3; i += 1) {
T:   C[i] = 0;
      for (int k = 0; k < 3; k += 1)
S:     C[i] += A[i] * B[k];
U:     C[i] = C[i];
}

```

	T(0)	S(0,0)	S(0,1)	S(0,2)	U(0)	T(1)	S(1,0)	S(1,1)	S(1,2)	U(1)	T(2)	S(2,0)	S(2,1)	S(2,2)	U(2)
c	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[0]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[1]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[2]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...



# Value Mapping

“DeLICM”


```

double c;
for (int i = 0; i < 3; i += 1) {
T:   C[i] = 0;
    for (int k = 0; k < 3; k += 1)
S:   C[i] += A[i] * B[k];
}


```

	T(0) S(0,0) S(0,1) S(0,2)					T(1) S(1,0) S(1,1) S(1,2)					T(2) S(2,0) S(2,1) S(2,2)				
c	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[0]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[1]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...
C[2]	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14 ...

# Experiments



# Experiments



# Chapter Summary

- LLVM mid-end canonicalization inhibits polyhedral optimization
- Can undo scalar optimizations on the polyhedral representation (“DeLICM”)
- Reasons to run Polly after canonicalization:
  - More optimizations, especially the inliner
  - More canonicalized, less dependent on input code
  - Avoid running canonicalization passes redundantly
  - No IR-modification when no polyhedral transformation was done

# SPEC CPU 2006 456.hmmmer

```
for (k = 1; k <= M; k++) {  
    mc[k] = mpp[k-1] + tpmm[k-1];  
    if ((sc = ip[k-1] + tpim[k-1]) > mc[k]) mc[k] = sc;  
    if ((sc = dpp[k-1] + tpdm[k-1]) > mc[k]) mc[k] = sc;  
    if ((sc = xmb + bp[k]) > mc[k]) mc[k] = sc;  
    mc[k] += ms[k];  
    if (mc[k] < -INFTY) mc[k] = -INFTY;  
    dc[k] = dc[k-1] + tpdd[k-1];  
    if ((sc = mc[k-1] + tpmd[k-1]) > dc[k]) dc[k] = sc;  
    if (dc[k] < -INFTY) dc[k] = -INFTY;  
    if (k < M) {  
        ic[k] = mpp[k] + tpmi[k];  
        if ((sc = ip[k] + tpim[k]) > ic[k]) ic[k] = sc;  
        ic[k] += is[k];  
        if (ic[k] < -INFTY) ic[k] = -INFTY;  
    }  
}
```

## SPEC CPU 2006 456.hmmmer

```
for (k = 1; k <= M; k++) {  
    mc[k] = mpp[k-1] + tpmm[k-1];  
    if ((sc = ip[k-1] + tpim[k-1]) > mc[k]) mc[k] = sc;  
    if ((sc = dpp[k-1] + tpdm[k-1]) > mc[k]) mc[k] = sc;  
    if ((sc = xmb + bp[k]) > mc[k]) mc[k] = sc;  
    mc[k] += ms[k];  
    if (mc[k] < -INFTY) mc[k] = -INFTY;  
    dc[k] = dc[k-1] + tpdd[k-1];  
    if ((sc = mc[k-1] + tpmd[k-1]) > dc[k]) dc[k] = sc;  
    if (dc[k] < -INFTY) dc[k] = -INFTY;  
    if (k < M) {  
        ic[k] = mpp[k] + tpmi[k];  
        if ((sc = ip[k] + tpim[k]) > ic[k]) ic[k] = sc;  
        ic[k] += is[k];  
        if (ic[k] < -INFTY) ic[k] = -INFTY;  
    }  
}
```

# SPEC CPU 2006 456.hmmmer

```
for (k = 1; k <= M; k++) {
```

```
mc[k] = mpp[k-1] + tpmm[k-1];
if ((sc = ip[k-1] + tpim[k-1]) > mc[k]) mc[k] = sc;
if ((sc = dpp[k-1] + tpdm[k-1]) > mc[k]) mc[k] = sc;
if ((sc = xmb + bp[k]) > mc[k]) mc[k] = sc;
mc[k] += ms[k];
if (mc[k] < -INFTY) mc[k] = -INFTY;
```

Compute mc[k] (*vectorizable*)

```
dc[k] = dc[k-1] + tpdd[k-1];
if ((sc = mc[k-1] + tpmd[k-1]) > dc[k]) dc[k] = sc;
if (dc[k] < -INFTY) dc[k] = -INFTY;
```

Compute dc[k] (**not** vectorizable)

```
if (k < M) {
```

```
ic[k] = mpp[k] + tpmi[k];
if ((sc = ip[k] + tpai[k]) > ic[k]) ic[k] = sc;
ic[k] += is[k];
if (ic[k] < -INFTY) ic[k] = -INFTY;
```

Compute ic[k] (*vectorizable*)

```
}
```

```
}
```

# LoopDistribution/LoopVectorizer

-enable-loop-distribute

■ Gerolf Hoflehner, *LLVM Performance Improvements and Headroom*, LLVM DevMtg 2015

```
for (k = 1; k <= M; k++) {  
    mc[k] = mpp[k-1] + tpmm[k-1];  
    if ((sc = ip[k-1] + tpim[k-1]) > mc[k]) mc[k] = sc;  
    if ((sc = dpp[k-1] + tpdm[k-1]) > mc[k]) mc[k] = sc;  
    if ((sc = xmb + bp[k]) > mc[k]) mc[k] = sc;  
    mc[k] += ms[k];  
    if (mc[k] < -INFTY) mc[k] = -INFTY;  
}
```

*loop-vectorized*

```
for (k = 1; k <= M; k++) {  
    dc[k] = dc[k-1] + tpdd[k-1];  
    if ((sc = mc[k-1] + tpmd[k-1]) > dc[k]) dc[k] = sc;  
    if (dc[k] < -INFTY) dc[k] = -INFTY;  
    if (k < M) {  
        ic[k] = mpp[k] + tpmi[k];  
        if ((sc = ip[k] + tpim[k]) > ic[k]) ic[k] = sc;  
        ic[k] += is[k];  
        if (ic[k] < -INFTY) ic[k] = -INFTY;  
    }  
}
```

**not** vectorized



# LoopDistribution/LoopVectorizer

-loop-distribute-non-if-convertible

```
for (k = 1; k <= M; k++) {
    mc[k] = mpp[k-1] + tpmm[k-1];
    if ((sc = ip[k-1] + tpim[k-1]) > mc[k]) mc[k] = sc;
    if ((sc = dpp[k-1] + tpdm[k-1]) > mc[k]) mc[k] = sc;
    if ((sc = xmb + bp[k]) > mc[k]) mc[k] = sc;
    mc[k] += ms[k];
    if (mc[k] < -INFTY) mc[k] = -INFTY;
}
```

*loop-vectorized*

```
for (k = 1; k <= M; k++) {
    dc[k] = dc[k-1] + tpdd[k-1];
    if ((sc = mc[k-1] + tpmd[k-1]) > dc[k]) dc[k] = sc;
    if (dc[k] < -INFTY) dc[k] = -INFTY;
}
```

**not** vectorized

```
for (k = 1; k <= M; k++) {
    if (k < M) {
        ic[k] = mpp[k] + tpmi[k];
        if ((sc = ip[k] + tpil[k]) > ic[k]) ic[k] = sc;
        ic[k] += is[k];
        if (ic[k] < -INFTY) ic[k] = -INFTY;
    }
}
```

vectorized with if-conversion

# Polly

-polly-stmt-granularity=bb

```
for (k = 1; k <= M; k++) {  
  Stmt1: mc[k] = mpp[k-1] + tpmm[k-1];  
  if ((sc = ip[k-1] + tpim[k-1]) > mc[k]) mc[k] = sc;  
  if ((sc = dpp[k-1] + tpdm[k-1]) > mc[k]) mc[k] = sc;  
  if ((sc = xmb + bp[k]) > mc[k]) mc[k] = sc;  
  mc[k] += ms[k];  
  if (mc[k] < -INFTY) mc[k] = -INFTY;  
  dc[k] = dc[k-1] + tpdd[k-1];  
  if ((sc = mc[k-1] + tpmd[k-1]) > dc[k]) dc[k] = sc;  
  if (dc[k] < -INFTY) dc[k] = -INFTY;  
  
  if (k < M) {  
    Stmt2: ic[k] = mpp[k] + tpmi[k];  
    if ((sc = ip[k] + tpri[k]) > ic[k]) ic[k] = sc;  
    ic[k] += is[k];  
    if (ic[k] < -INFTY) ic[k] = -INFTY;  
  }  
}
```

# Polly

-polly-stmt-granularity=scalar-indep

```
for (k = 1; k <= M; k++) {
```

```
Stmt1: mc[k] = mpp[k-1] + tpmm[k-1];
if ((sc = ip[k-1] + tpim[k-1]) > mc[k]) mc[k] = sc;
if ((sc = dpp[k-1] + tpdm[k-1]) > mc[k]) mc[k] = sc;
if ((sc = xmb + bp[k]) > mc[k]) mc[k] = sc;
mc[k] += ms[k];
if (mc[k] < -INFTY) mc[k] = -INFTY;
```

```
Stmt2: dc[k] = dc[k-1] + tpdd[k-1];
if ((sc = mc[k-1] + tpmd[k-1]) > dc[k]) dc[k] = sc;
if (dc[k] < -INFTY) dc[k] = -INFTY;
```

```
if (k < M) {
```

```
Stmt3: ic[k] = mpp[k] + tpmi[k];
if ((sc = ip[k] + tpim[k]) > ic[k]) ic[k] = sc;
ic[k] += is[k];
if (ic[k] < -INFTY) ic[k] = -INFTY;
}
}
```


# Loop Distribution by Polyhedral Scheduler

```
$ opt -polly-stmt-granularity=scalar-indep -polly-invariant-load-hoisting -polly-use-llvm-names \
    fast_algorithms.ll -polly-opt-isl -polly-ast -analyze
```

```
[...]
```

```
{
    for (int c0 = 0; c0 < _lcssa; c0 += 1)
        Stmt_for_body72(c0);
    for (int c0 = 0; c0 < _lcssa; c0 += 1)
        Stmt_for_body721(c0);
    for (int c0 = 0; c0 < _lcssa - 1; c0 += 1)
        Stmt_if_then167(c0);
    if (_lcssa >= 1)
        Stmt_for_end204_loopexit();
}
```

# Experiments



# Chapter Summary

- Finer-grained statements
- One basic block  $\Rightarrow$  multiple statement if no computation is shared
- Enables loop distribution by Polly
- Speed-up of 80% in 456.hmmer
- With support from Nandini Singhal



*That's all Folks!*

# Summary

- 1 COSMO weather forecasting on GPGPUs
- 2 Life Range Reordering (Verdoolaege IMPACT'16)
- 3 DGEMM detection also with C++ expression templates (Roman Gareev)
- 4 Correct types for loop transformations (Maximilian Falkenstein)
- 5 Some LLVM passes make polyhedral optimization harder
- 6 `-polly-position=early` vs. `-polly-position=before-vectorizer`
- 7 DeLICM: Avoiding scalar dependencies
- 8 `-polly-stmt-granularity` and loop-distribution in 456.hmmmer (with Nandini Singhal)