



VECTORIZING LOOPS WITH VPLAN – CURRENT STATE AND NEXT STEPS

Ayal Zaks and Gil Rapaport, Vectorization Team, Intel Corporation

October 18th, 2017 US LLVM Developers' Meeting, San Jose, CA

Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © 2017, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

Optimization Notice

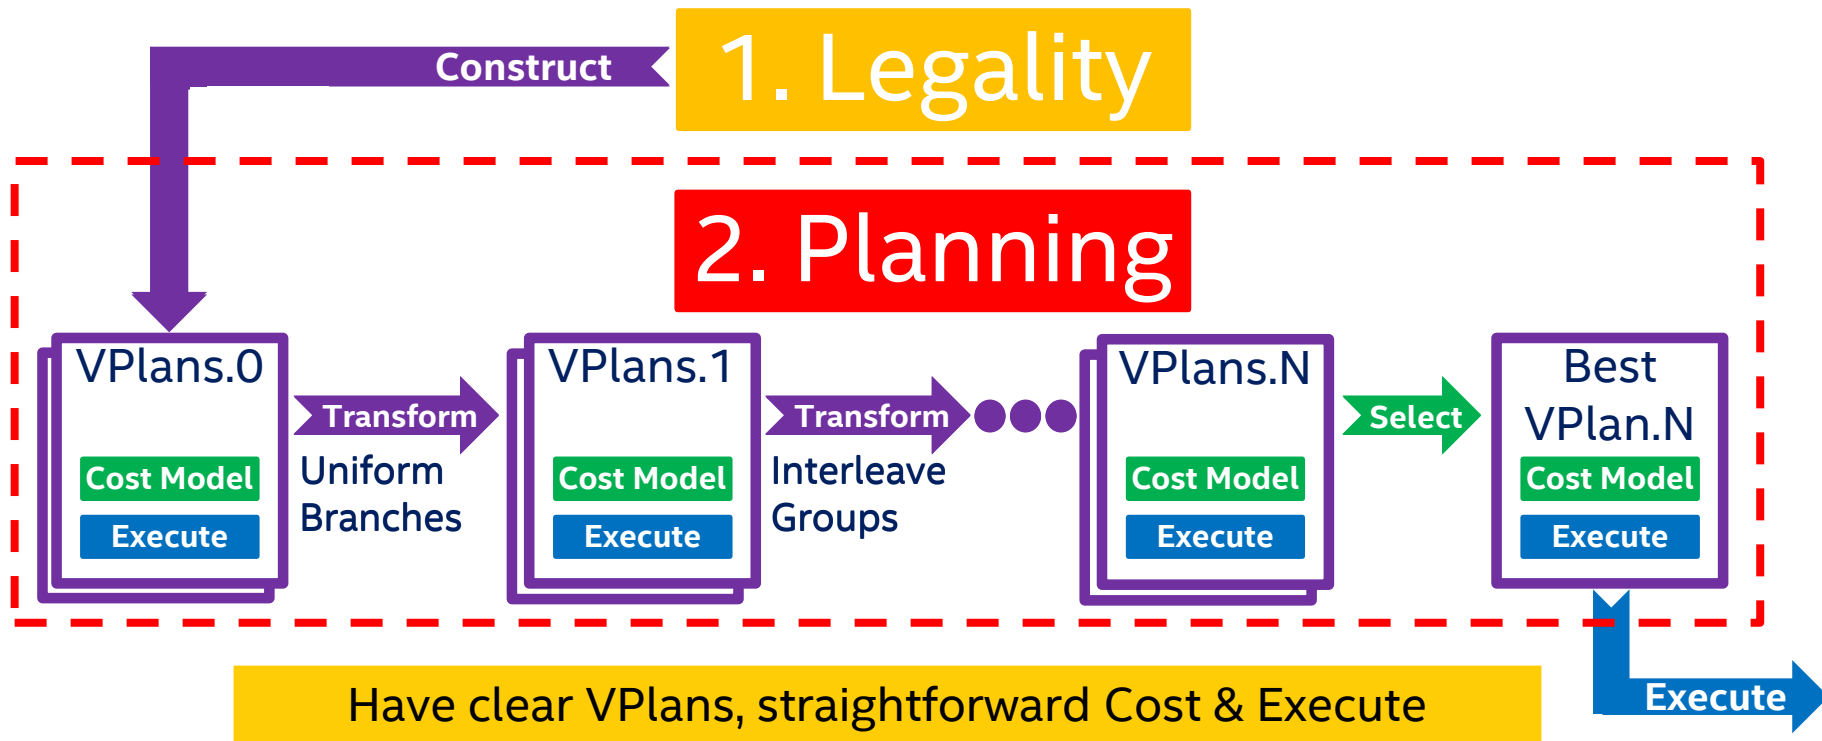
Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

Key Takeaways

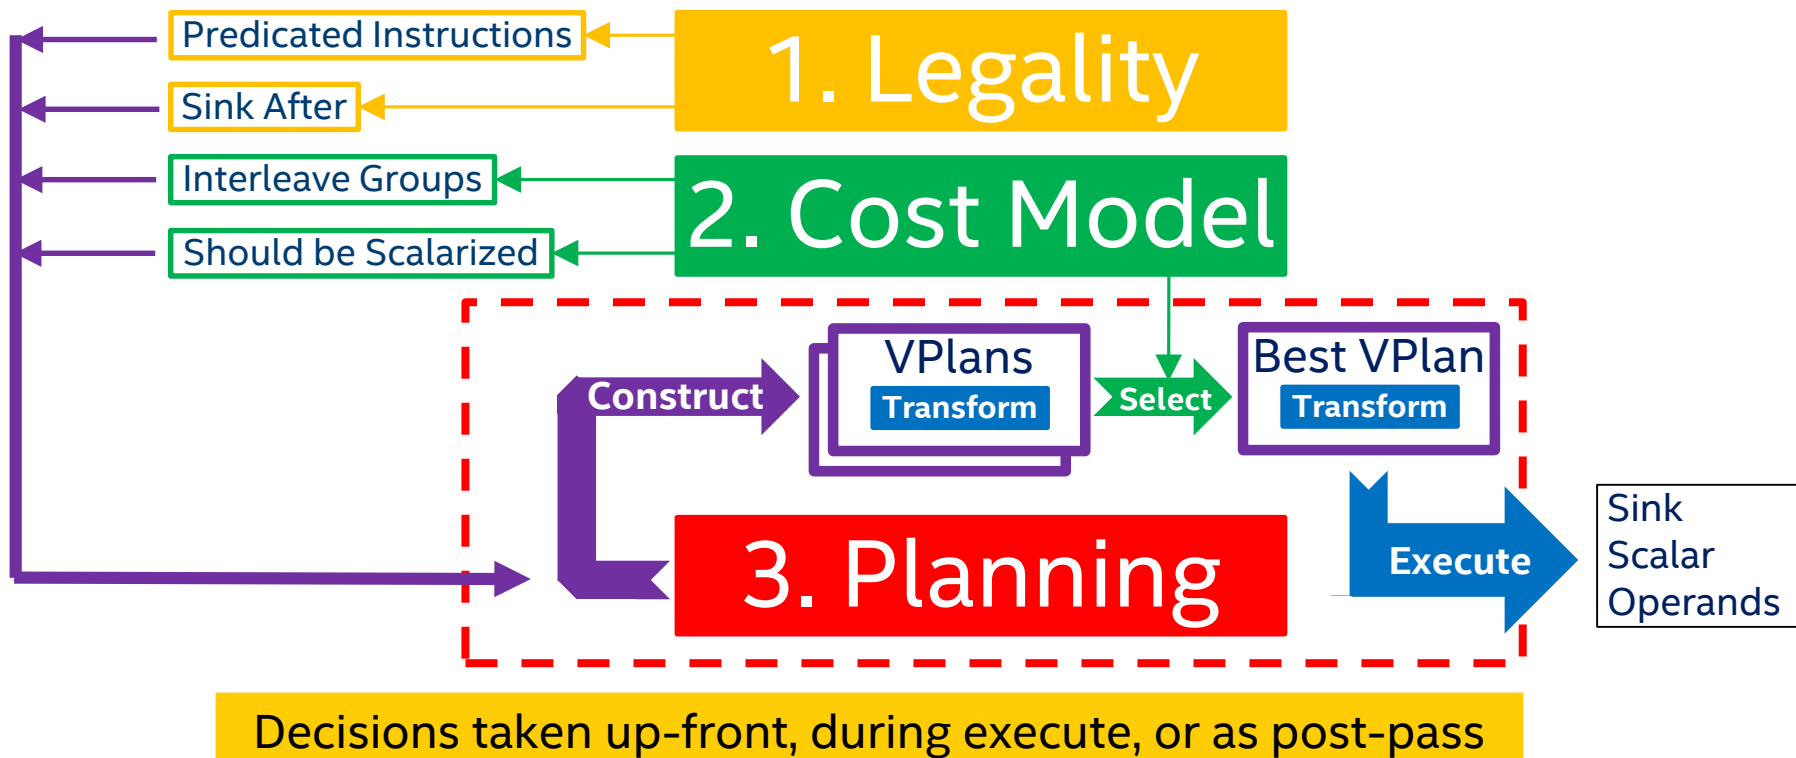
- A. Current State: 1st step introducing VPlan to Loop Vectorizer – committed
 - 1. Records vectorization decisions in VPlan
 - 2. Drives vector code generation by executing a VPlan
- B. Going Forward: shift vectorization process to be VPlan-based
 - 1. Refine the model, include masking and break Recipes into VPInstructions
 - 2. Carry out decisions based on VPlan, in addition to recording them
 - 3. Make decisions based on VPlan, including legal and cost-based analyses

Recap: Loop Vectorization Plan



A. CURRENT STATE OF VPLAN

1st Step Committed: VPlan Refactors Transform



VPlan Model: Current State

```
void foo(int *a, int b, int *c) {
  for (int i = 0; i < 10000; ++i)
    if (a[i] > 777)
      a[i] = b - (c[100*i] * 7 + a[i]) / b;
}
```



LLVM-IR Before Vectorizer

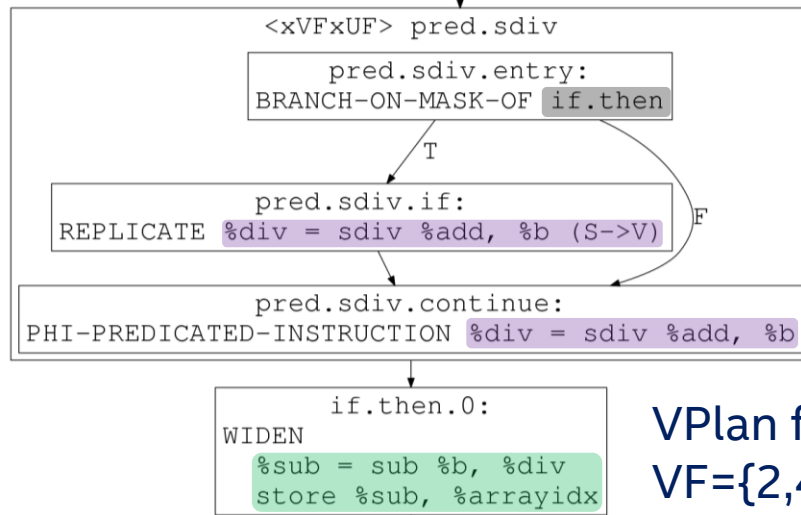
```
for.body:
  %indvars.iv = phi i64 [ 0, %entry ], [ %indvars.iv.next, %for.inc ]
  %arrayidx = getelementptr inbounds i32, i32* %a, i64 %indvars.iv
  %0 = load i32, i32* %arrayidx, align 4
  %cmp1 = icmp sgt i32 %0, 777
  br i1 %cmp1, label %if.then, label %for.inc

if.then:
  %1 = mul nsw i64 %indvars.iv, 100
  %arrayidx3 = getelementptr inbounds i32, i32* %c, i64 %1
  %2 = load i32, i32* %arrayidx3, align 4
  %mul4 = mul nsw i32 %2, 7
  %add = add nsw i32 %mul4, %0
  %div = sdiv i32 %add, %b
  %sub = sub nsw i32 %b, %div
  store i32 %sub, i32* %arrayidx, align 4
  br label %for.inc
```



```
for.body:
WIDEN-INDUCTION %indvars.iv = phi 0, %indvars.iv.next
CLONE %arrayidx = getelementptr %a, %indvars.iv
WIDEN
  %0 = load %arrayidx
  %cmp1 = icmp %0, 777
```

```
if.then:
REPLICATE %1 = mul %indvars.iv, 100
REPLICATE %arrayidx3 = getelementptr %c, %1
REPLICATE %2 = load %arrayidx3
REPLICATE %mul4 = mul %2, 7
REPLICATE %add = add %mul4, %0
```



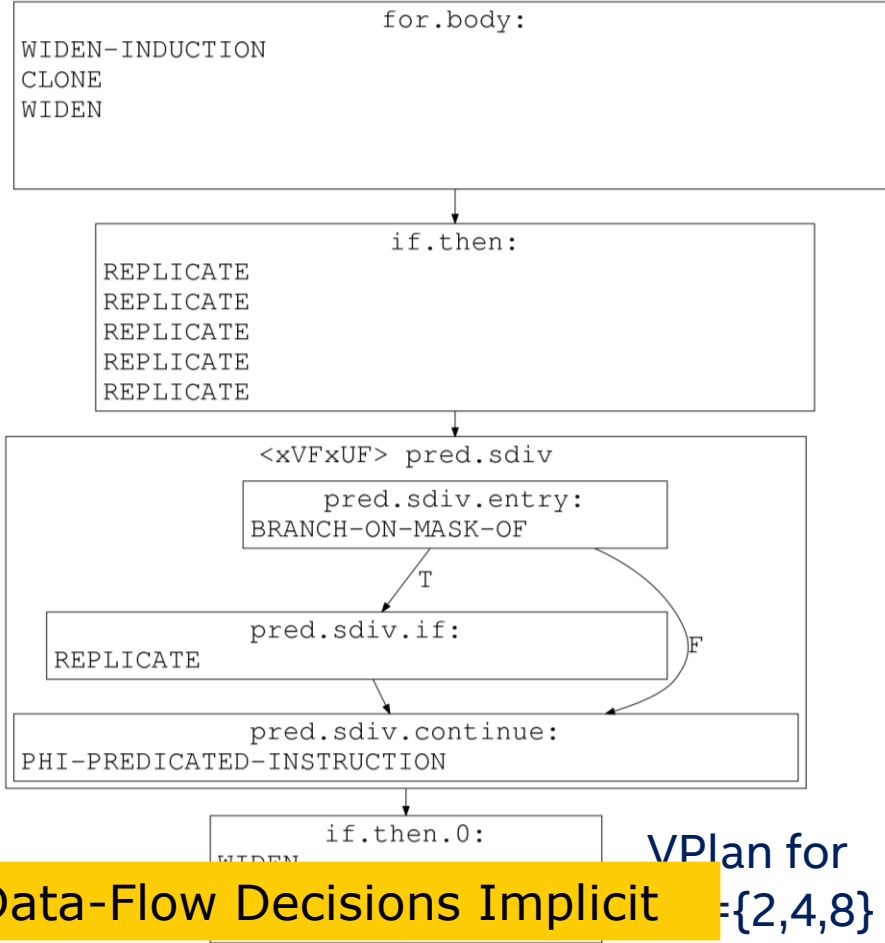
VPlan for
VF={2,4,8}

VPlan Model: Current State

Recipe: models a sequence of instructions to appear in the vectorized code.
May refer to Ingredients.
Ingredient: element of the original scalar loop, such as an existing instruction.

VPRecipeBase
void execute()= 0
VPBasicBlock *getParent()

VPWidenRecipe
void execute()



Control-Flow Decisions Explicit, Data-Flow Decisions Implicit

VPlan for
={2,4,8}

B.1. MODEL MASKING AND INSTRUCTIONS

VPlan Model: Next Step

```
void foo(int *a, int b, int *c) {  
    for (int i = 0; i < 10000; ++i)  
        if (a[i] > 777)  
            a[i] = b - (c[100*i] * 7 + a[i]) / b;  
}
```



```
for.body:  
WIDEN-INDUCTION %i.015 = phi 0, %inc  
CLONE %arrayidx = getelementptr %a, %i.015  
WIDEN %0 = load %arrayidx  
WIDEN  
(%vp63408) %cmp1 = icmp %0, 777
```

```
if.then:  
REPLICATE %mul = mul %i.015, 100  
REPLICATE %arrayidx2 = getelementptr %c, %mul  
WIDEN %1 = load %arrayidx2, %vp63408  
REPLICATE %mul3 = mul %1, 7  
REPLICATE %add = add %mul3, %0
```

```
<xVFxUF> pred.sdiv  
pred.sdiv.entry:  
BRANCH-ON-MASK %vp63408  
T  
pred.sdiv.if:  
REPLICATE %div = sdiv %add, %b (S->V)  
F  
pred.sdiv.continue:  
PHI-PREDICATED-INSTRUCTION %div = sdiv %add, %b
```

```
if.then.0:  
WIDEN  
3408
```

VPValue

VPUsers users()

VPUser

VPValues operands()

VPRecipeBase

void execute()= 0
VPBasicBlock *getParent()

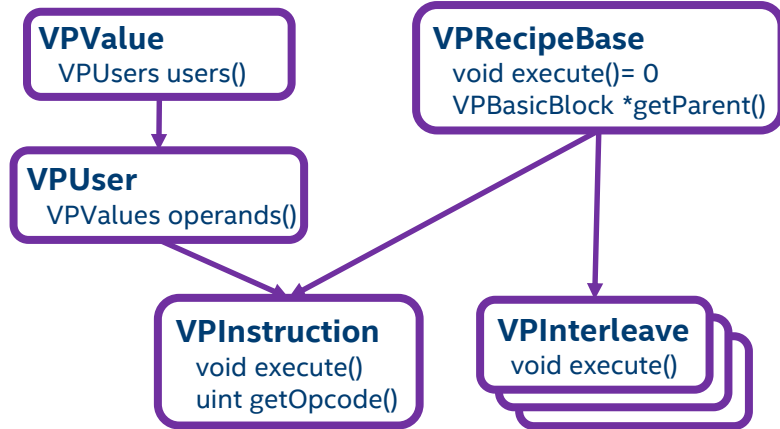
VPInterleave

void execute()

Model Masking in VPlan using Def/Use Relations [D38676]

VPlan Model: Next Step (cont'd)

```
void foo(int* a, int b, int* c) {  
    for (int i = 0; i < 10000; ++i)  
        if (a[i] > 777) {  
            c[i] = b;  
            if (a[i] > 888)  
                a[i] = b;  
        }  
}
```



VPlan for VF={2,4,8,16}

```
for.body:  
WIDEN-INDUCTION %i.017 = phi 0, %inc  
CLONE %arrayidx = getelementptr %a, %i.017  
WIDEN %0 = load %arrayidx  
WIDEN  
(%vp27696) %cmp1 = icmp %0, 777
```

```
if.then:  
CLONE %arrayidx2 = getelementptr %c, %i.017  
WIDEN store %b, %arrayidx2, %vp27696  
WIDEN %1 = load %arrayidx  
WIDEN  
(%vp30784) %cmp4 = icmp %1, 888
```

```
if.then5:  
EMIT %vp58664 = and %vp30784 %vp27696  
WIDEN store %b, %arrayidx, %vp58664
```

VPInstruction: Instruction-level Modeling in VPlan [D38676]

B.2. FROM RECORDING DECISIONS TO CARRYING THEM OUT

Taking Decision (1/4): Interleave Groups

```
void foo(int *a, int n, int *c) {  
    for (int i = 0; i < n; ++i)  
        a[i] = 3*c[2*i+1] + c[2*i];  
}
```

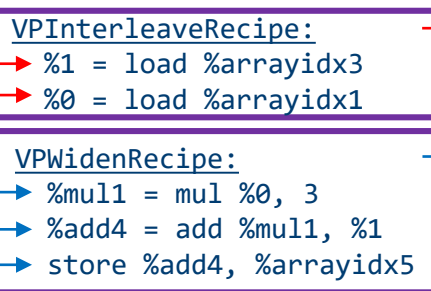


IR Before Vectorizer

foo.body:

```
...  
%0 = load i32, %arrayidx  
%mul1 = mul %0, 3  
%1 = load i32, %arrayidx3  
%add4 = add %mul1, %1  
store %add4, %arrayidx5  
...
```

VPlan for VF=4



Ingredients

VPlan Execution

IR After Vectorizing for VF=4

vector.body:

```
...  
%all = load <8 x i32>, %5  
%even = shufflevector %all, <0,2,4,6>  
%odd = shufflevector %all, <1,3,5,7>  
%6 = mul %odd, <3,3,3,3>  
%9 = add %6, %even  
store %9, %12  
...
```

Effectively hoists load %1 to join load %0

Taking Decision (2/4): Unravel 1st Order Recurrence

```
void sink_after(short *a, int *b, int n) {  
    for (int i = 0; i < n; ++i)  
        b[i] = (a[i] * a[i+1]);  
}
```

IR Before Vectorizer



```
foo.body:  
    %iv = phi i64 [ 0, %entry ], [ %iv.next, %for.body ]  
    %0 = phi i16 [ %.pre, %entry ], [ %1, %for.body ]  
    %conv = sext i16 %0 to i32  
    %iv.next = add nuw nsw i64 %iv, 1  
    %arrayidx2 = getelementptr i16, i16* %a, i64 %iv.next  
    %1 = load i16, i16* %arrayidx2  
    %conv3 = sext i16 %1 to i32  
    %mul = mul nsw i32 %conv3, %conv  
    %arrayidx5 = getelementptr i32, i32* %b, i64 %iv  
    store i32 %mul, i32* %arrayidx5  
    %exitcond = icmp eq i64 %indvars.iv.next, %n  
    br i1 %exitcond, label %for.end, label %for.body
```

IR After Vectorizer

```
vector.body  
    %iv = phi i64 [ 0, %vec.ph ], [ %iv.next, %vec.body ]  
    %recur = phi <4 x i16> [ %recur.init, %vec.ph ],  
        [ %wide.load, %vec.body ]  
    ...  
    %3 = getelementptr inbounds i16, i16* %a, i64 %2  
    %wide.load = load <4 x i16>, <4 x i16>* %5, align 2  
    %6 = shufflevector <4 x i16> %recur,  
        <4 x i16> %wide.load,  
        <4 x i32> <3, 4, 5, 6>  
    %7 = sext <4 x i16> %6 to <4 x i32>  
    %8 = sext <4 x i16> %wide.load to <4 x i32>  
    %9 = mul nsw <4 x i32> %8, %7
```

Phase-ordering: first sink cast after load, then hoist interleave load [PR34743]

Taking Decision (3/4): Predication

- Must convert divergent branches using masking
- Much more challenging for outer-loop vectorization

- Earlier today: *VPlan + RV: A Proposal* by Simon Moll and Sebastian Hack

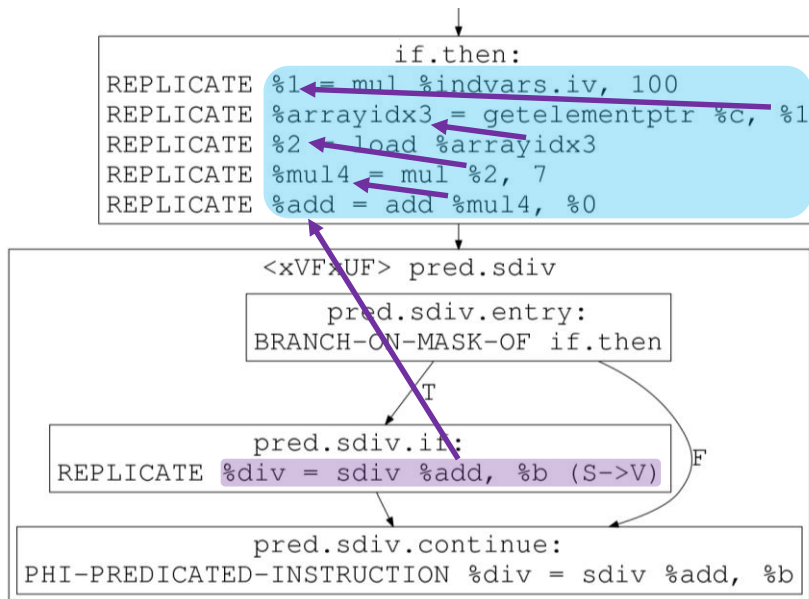
- Last year's *Extending LoopVectorizer*: by Hideki Saito:

```
// vectorize here
for (i = ilb; i < iub; ++i) {
    ...
    for (j = jlb(i); j < jub(i); ++j) {
        while (cond(i, j)) { ... }
        if (...) break;
    }
}
```

```
julb = hmin(jlb(i));
juub = hmax(jub(i));
cont1 = T;
for (j = julb; j < juub; ++j) {
    if (jlb(i) <= j && j < jub(i) && cont1) {
        cont2 = cond(i, j);
        while (hor(cont2)) {
            if (cont2) {
                ...
                cont2 = cond(i, j);
            }
        }
        if (...) cont1 = F;
        if (!hor(cont1)) break;
    }
}
```

Take Predication Decisions by Transforming One VPlan to Another

Taking Decision (4/4): SinkScalarOperands



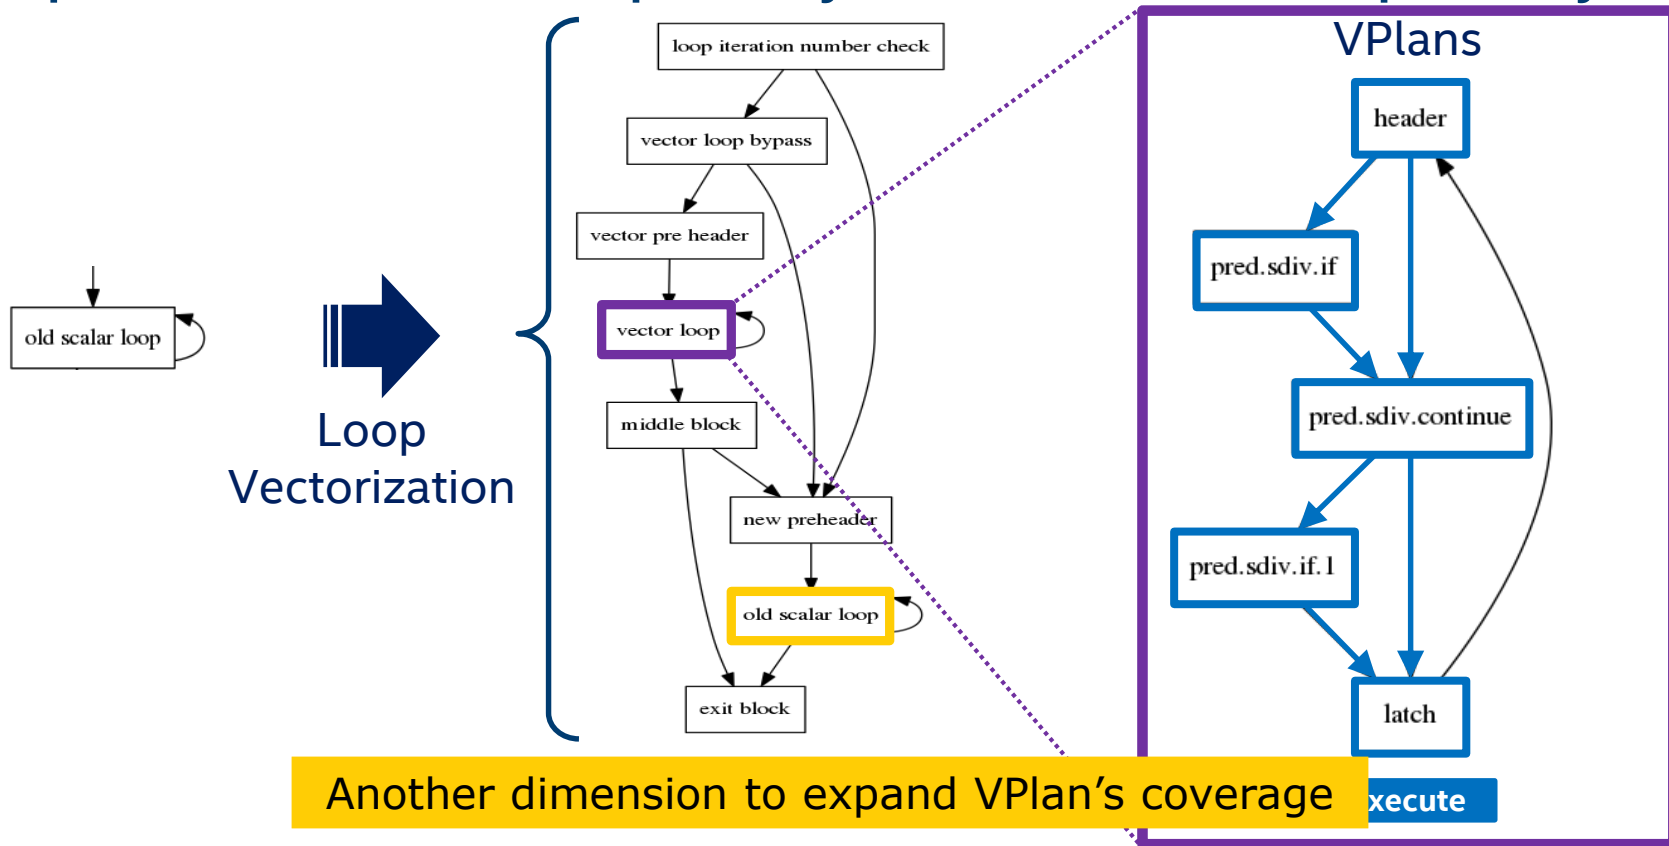
Requires Fine-grain Modeling of Def/Use at instruction-level

B.3. FROM CARRYING OUT DECISIONS TO MAKING THEM

Use VPlan to also Make Vectorization Decisions

- Instead of first making the decisions, and then using VPlan to carry them out
- Run cost-based analyses on VPlan
 - Based on cost estimates computed by VPlan
 - Based on VPIinstruction model
- Apply desired decisions by transforming VPlan, potentially versioning it
 - Based on “what-if” versioning support

Expand VPlan's Scope Beyond Vector Loop Body



Taking Decision (1/4): Interleave Groups – revisit

```
void foo(int *a, int n, int *c) {  
    for (int i = 0; i < n; ++i)  
        a[i] = 3*c[2*i+1] + c[2*i];  
}
```



IR Before Vectorizer

foo.body:

```
...  
%0 = load i32, %arrayidx  
%mul1 = mul %0, 3  
%1 = load i32, %arrayidx3  
%add4 = add %mul1, %1  
store %add4, %arrayidx5  
...
```

VPlan for VF=4

VPInterleaveRecipe:

- %1 = load %arrayidx3
- %0 = load %arrayidx1

VPWidenRecipe:

- %mul1 = mul %0, 3
- %add4 = add %mul1, %1
- store %add4, %arrayidx5

Ingredients

VPlan Execution

IR After Vectorizing for VF=4

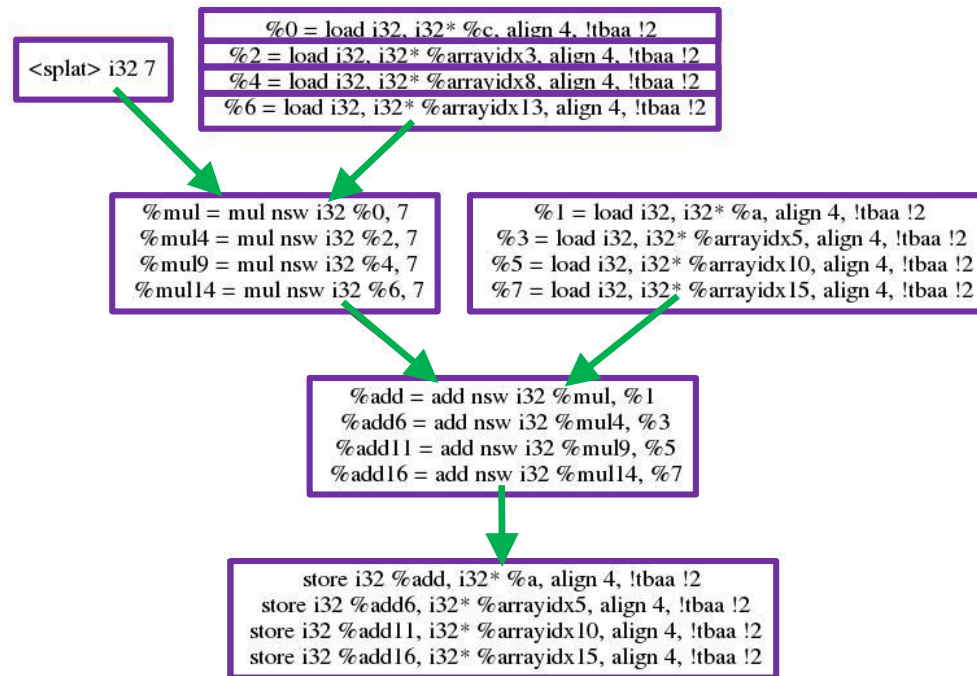
vector.body:

```
...  
%all = load <8 x i32>, %5  
%even = shufflevector %all, <0,2,4,6>  
%odd = shufflevector %all, <1,3,5,7>  
%6 = mul %odd, <3,3,3,3>  
%9 = add %6, %even  
store %9, %12  
...
```

Combining two load %0, %1 into one load %a11 – looks familiar?

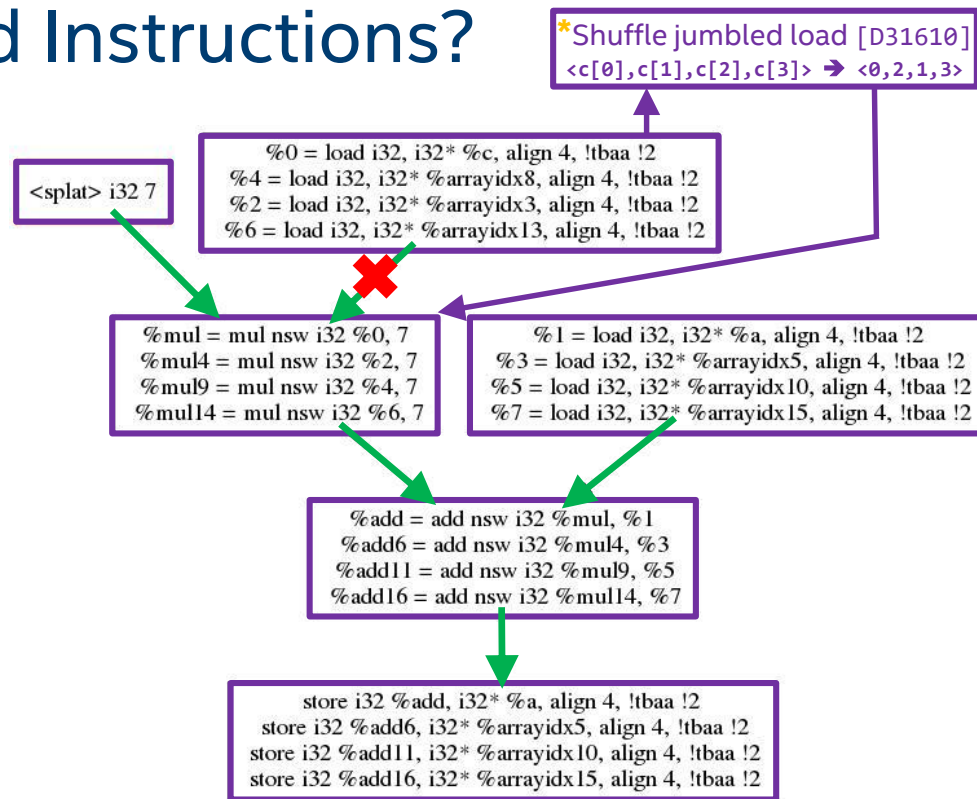
A Model for Vectorized Instructions?

```
void foo(int * restrict a, int b, int *c) {  
    a[0] = c[0] * 7 + a[0];  
    a[1] = c[2] * 7 + a[1];  
    a[2] = c[1] * 7 + a[2];  
    a[3] = c[3] * 7 + a[3];  
}
```



A Model for Vectorized Instructions?

```
void foo(int * restrict a, int b, int *c) {  
    a[0] = c[0] * 7 + a[0];  
    a[1] = c[2] * 7 + a[1];  
    a[2] = c[1] * 7 + a[2];  
    a[3] = c[3] * 7 + a[3];  
}
```



SLP: Spill Cost = 0.

Def/Use Model for New & Ingredient-based Instructions & Dependences

Key Takeaways

- A. Current State: 1st step introducing VPlan to Loop Vectorizer – committed
 - 1. Records vectorization decisions in VPlan
 - 2. Drives vector code generation by executing a VPlan
- B. Going Forward: shift vectorization process to be VPlan-based
 - 1. Refine the model, include masking and break Recipes into VPInstructions
 - 2. Carry out decisions based on VPlan, in addition to recording them
 - 3. Make decisions based on VPlan, including legal and cost-based analyses

