

# Finding Missed Optimizations in LLVM (and other compilers)

Gergö Barany  
Inria Paris, France  
`gergo.barany@inria.fr`

EuroLLVM, Bristol, April 16, 2018

<https://github.com/gergo-/missed-optimizations>

## Main takeaways

Does your compiler *always* optimize well?

## Main takeaways

Does your compiler *always* optimize well?

- ▶ compare compilers' outputs to find missed optimizations
- ▶ automated toolchain finds minimal test cases
- ▶ issues found in GCC, Clang, CompCert:
  - ▶ peephole optimizations, dead stores, useless spills, missed instruction selection patterns, missed copy propagation, ...

## Example: missing range analysis

Generated source code:

```
int f(int p, int q) {  
    return q + (p % 6) / 9;  
}
```

Clang:

GCC:

## Example: missing range analysis

Generated source code:

```
int f(int p, int q) {  
    return q + (p % 6) / 9;  
}
```

Clang:

```
movw r2, #43691  
movt r2, #10922  
smmul r2, r0, r2  
add r2, r2, r2, lsr #31  
add r2, r2, r2, lsl #1  
sub r0, r0, r2, lsl #1  
movw r2, #36409  
movt r2, #14563  
smmul r0, r0, r2  
asr r2, r0, #1  
add r0, r2, r0, lsr #31  
add r0, r0, r1
```

GCC:

## Example: missing range analysis

Generated source code:

```
int f(int p, int q) {  
    return q + (p % 6) / 9;  
}
```

Clang:

```
movw r2, #43691  
movt r2, #10922  
smmul r2, r0, r2  
add r2, r2, r2, lsr #31  
add r2, r2, r2, lsl #1  
sub r0, r0, r2, lsl #1  
movw r2, #36409  
movt r2, #14563  
smmul r0, r0, r2  
asr r2, r0, #1  
add r0, r2, r0, lsr #31  
add r0, r0, r1
```

GCC:

```
mov r0, r1
```

## Example: missing range analysis

Generated source code:

```
int f(int p, int q) {  
    return q + (p % 6) / 9;  
}
```

$(p \% 6 \in [-5, 5],$   
division truncates to 0)

Clang:

```
movw r2, #43691  
movt r2, #10922  
smmul r2, r0, r2  
add r2, r2, r2, lsr #31  
add r2, r2, r2, lsl #1  
sub r0, r0, r2, lsl #1  
movw r2, #36409  
movt r2, #14563  
smmul r0, r0, r2  
asr r2, r0, #1  
add r0, r2, r0, lsr #31  
add r0, r0, r1
```

GCC:

```
mov r0, r1
```

[https://bugs.llvm.org/show\\_bug.cgi?id=34517](https://bugs.llvm.org/show_bug.cgi?id=34517) (fixed)

## Example: redundant code

Source code:

```
int fn3(  
    double c,  
    int *p, int *q)  
{  
    int i = (int)c;  
    *p = i;  
    *q = i;  
    return i;  
}
```

Clang:

GCC:

```
vcvt.s32.f64 s15, d0  
vstr.32 s15, [r0]  
vmov r0, s15  
vstr.32 s15, [r1]
```

## Example: redundant code

Source code:

```
int fn3(  
    double c,  
    int *p, int *q)  
{  
    int i = (int)c;  
    *p = i;  
    *q = i;  
    return i;  
}
```

Clang:

```
vcvt.s32.f64 s2, d0  
vstr s2, [r0]
```

GCC:

```
vcvt.s32.f64 s15, d0  
vstr.32 s15, [r0]  
vmov r0, s15  
vstr.32 s15, [r1]
```

## Example: redundant code

Source code:

```
int fn3(  
    double c,  
    int *p, int *q)  
{  
    int i = (int)c;  
    *p = i;  
    *q = i;  
    return i;  
}
```

Clang:

```
vcvt.s32.f64 s2, d0  
vstr s2, [r0]  
vcvt.s32.f64 s2, d0
```

GCC:

```
vcvt.s32.f64 s15, d0  
vstr.32 s15, [r0]  
vmov r0, s15  
vstr.32 s15, [r1]
```

## Example: redundant code

Source code:

```
int fn3(  
    double c,  
    int *p, int *q)  
{  
    int i = (int)c;  
    *p = i;  
    *q = i;  
    return i;  
}
```

Clang:

```
vcvt.s32.f64 s2, d0  
vstr s2, [r0]  
vcvt.s32.f64 s2, d0  
vcvt.s32.f64 s0, d0
```

GCC:

```
vcvt.s32.f64 s15, d0  
vstr.32 s15, [r0]  
vmov r0, s15  
vstr.32 s15, [r1]
```

## Example: redundant code

Source code:

```
int fn3(  
    double c,  
    int *p, int *q)  
{  
    int i = (int)c;  
    *p = i;  
    *q = i;  
    return i;  
}
```

Clang:

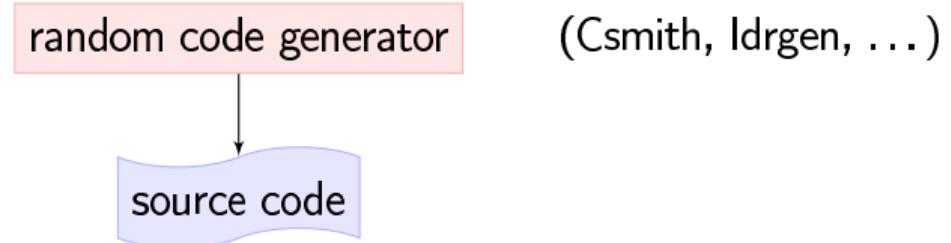
```
vcvt.s32.f64 s2, d0  
vstr s2, [r0]  
vcvt.s32.f64 s2, d0  
vcvt.s32.f64 s0, d0  
vmov r0, s0  
vstr s2, [r1]
```

GCC:

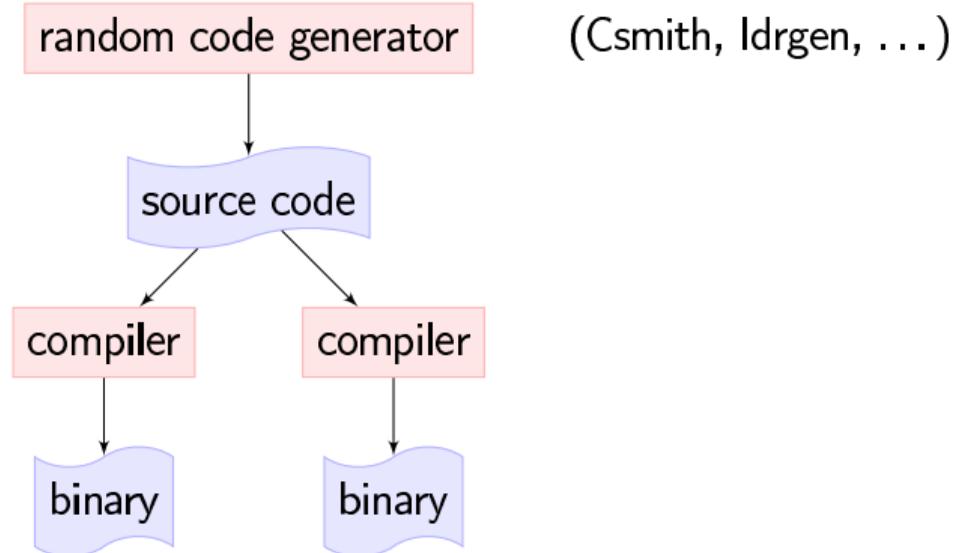
```
vcvt.s32.f64 s15, d0  
vstr.32 s15, [r0]  
vmov r0, s15  
vstr.32 s15, [r1]
```

[https://bugs.llvm.org/show\\_bug.cgi?id=33199](https://bugs.llvm.org/show_bug.cgi?id=33199) (fixed, then reverted)

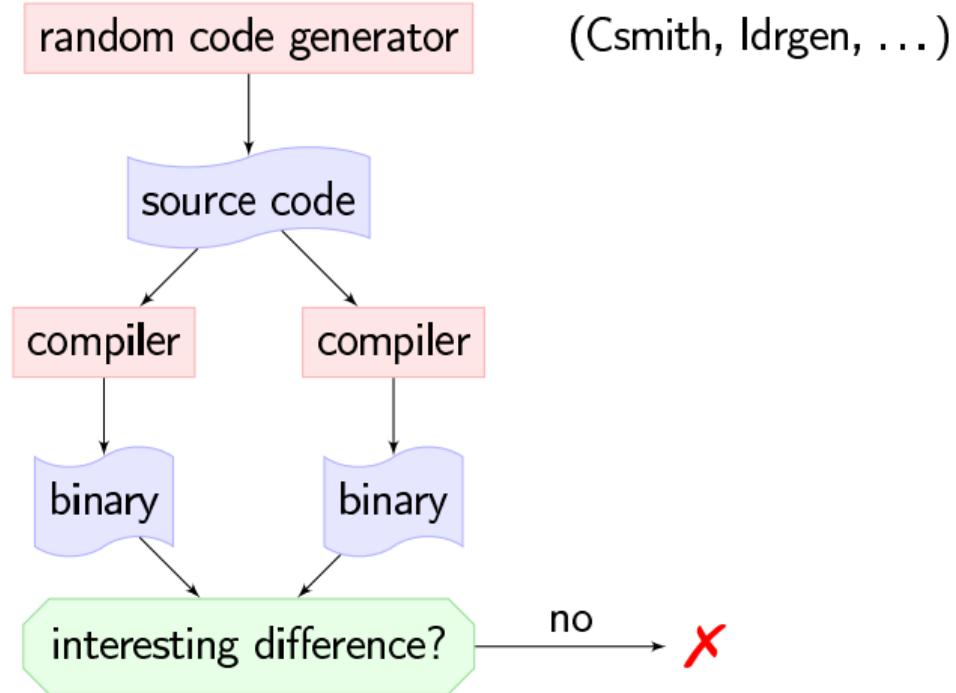
# Randomized differential testing



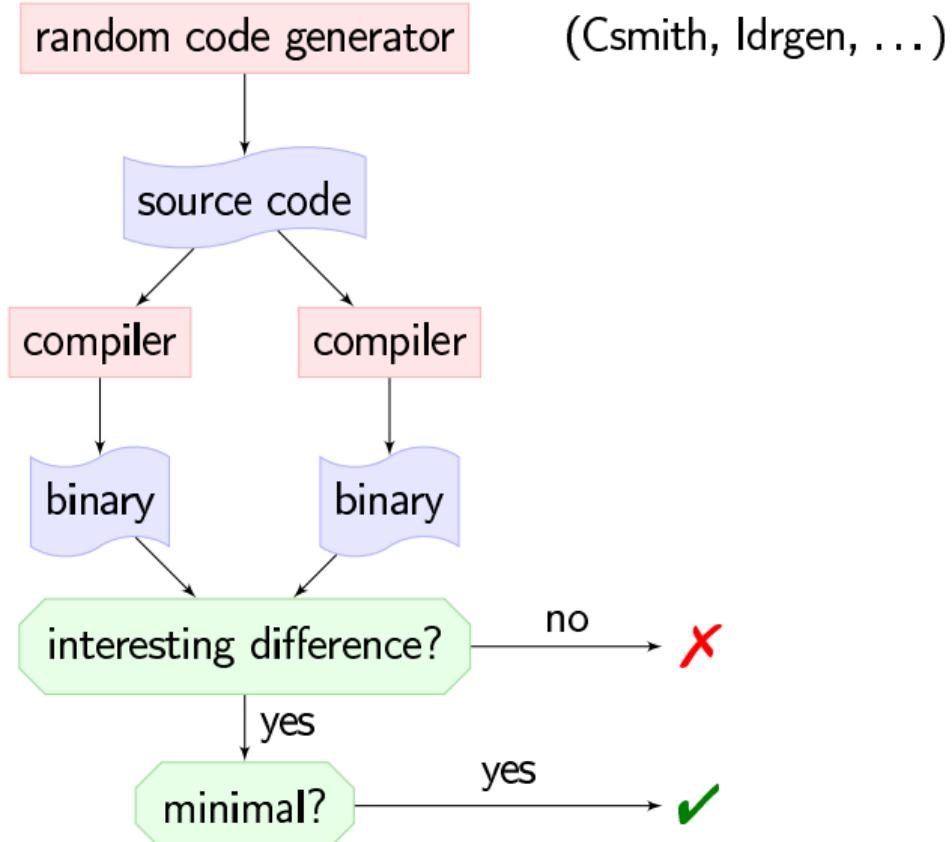
# Randomized differential testing



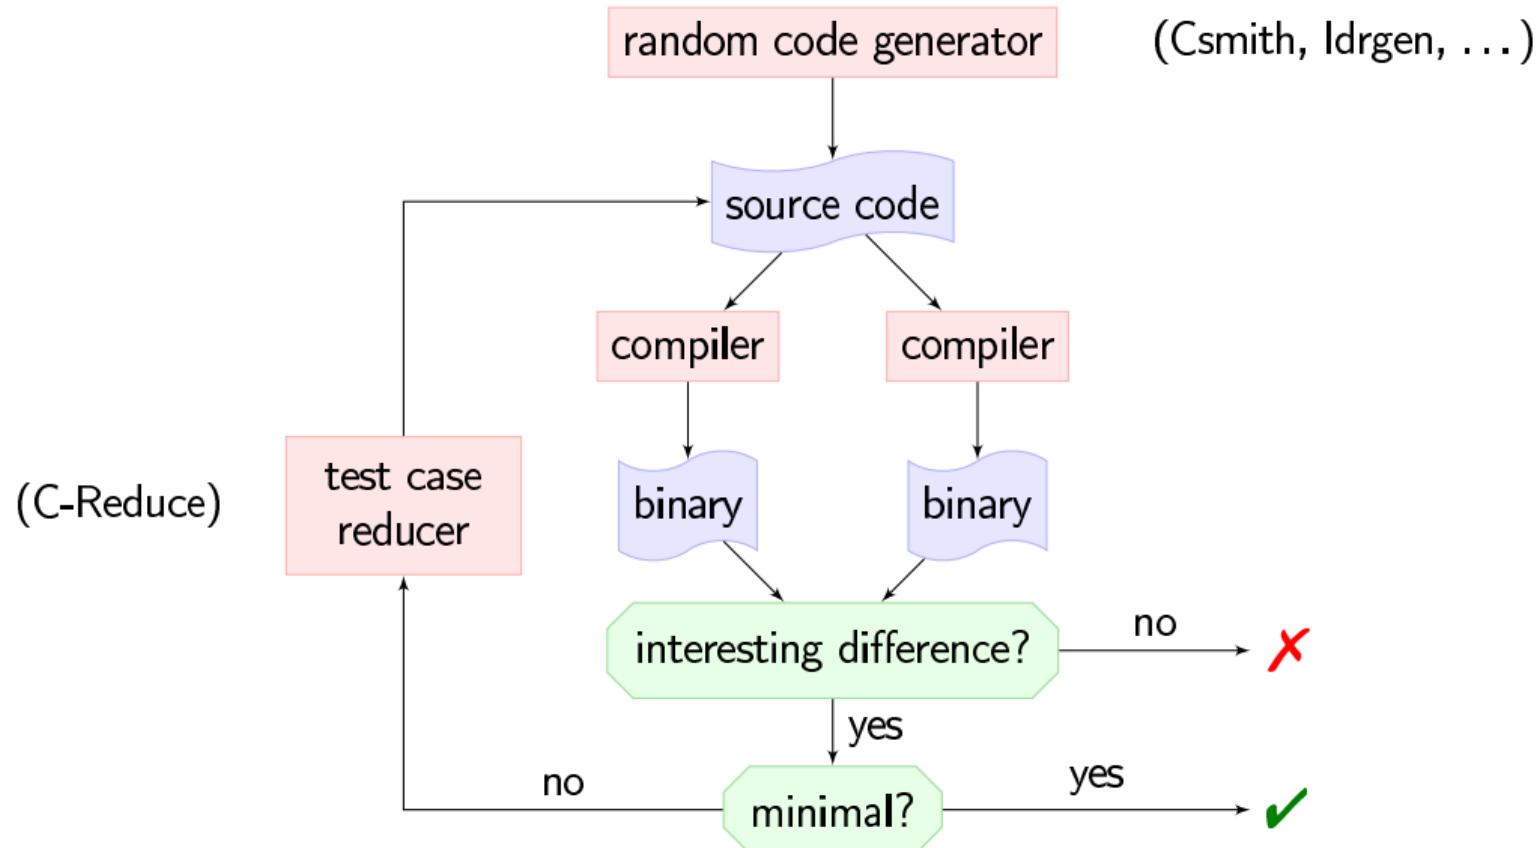
# Randomized differential testing



# Randomized differential testing



# Randomized differential testing



# Randomized differential testing for missed optimizations

interesting difference?

custom tool: optdiff

- ▶ binary analysis to find optimization differences
- ▶ assigns scores to binaries, compares
- ▶ based on `angr` binary analysis framework
  - ▶ multi-platform (x86, x86-64, ARM, PowerPC, ...)
  - ▶ Python API

## optdiff

Preliminaries: load binary, compute CFG, estimate basic block frequencies  $w_b$

Checkers: local scoring functions  $c: \text{instruction} \rightarrow \mathbb{N}$

Total score:  $s = \sum_{b \in f} w_b \cdot \sum_{i \in b} c(i)$

Examples checkers: number of instructions, general memory loads/stores, stack loads/stores, function calls, floating-point arithmetic instructions, vector instructions,

...

# Checker implementation: instructions

## Checkers

- ▶ Python functions with `@checker` decorator
- ▶ inspect one instruction at a time

```
@checker
def instructions(arch, instr):
    """Number of instructions."""
    return 1
```

## Checker implementation: loads

```
@checker
def loads(arch, instr):
    """Number of memory loads."""
    op = instr.insn.mnemonic
    if is_arm(arch):
        if op == 'ldr.d':
            # load doubleword
            return 2
        elif re.match('ldm.*', op):
            # load multiple
            return len(instr.insn.operands) - 1
        return bool(re.match('v?ldr.*', op))
    ... # other architectures
```

## Example: instruction selection

Source code:

```
double fn1(double p1, double p2,
           double p3)
{
    double a = -p1 - p2 * p3;
    return a;
}
```

Clang:

GCC:

```
vnmla.f64 d0, d1, d2
```

## Example: instruction selection

Source code:

```
double fn1(double p1, double p2,
           double p3)
{
    double a = -p1 - p2 * p3;
    return a;
}
```

Clang:

```
vneg.f64 d0, d0
vmls.f64 d0, d1, d2
```

GCC:

```
vnmla.f64 d0, d1, d2
```

Clang selected `vnmla` for `-(a * b) - c` but not for `-c - (a * b)`

<https://reviews.llvm.org/D35911> (merged)

## Example: register allocation

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, v, e;
    a = p4 - 5 + 9 * p1 * 7;
    c = 2 - p6 + 1;
    b = 3 * p6 * 4 - (p2 + 10) * a;
    d = (8 + p3) * p5 * 5 * p7 - p4;
    v = p1 - p8 - 2 -
        (p8 + p2 - b) * ((p3 + 6) * c);
    e = v * p5 + d;
    return e;
}
```

## Example: register allocation

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, v, e;
    a = p4 - 5 + 9 * p1 * 7;
    c = 2 - p6 + 1;
    b = 3 * p6 * 4 - (p2 + 10) * a;
    d = (8 + p3) * p5 * 5 * p7 - p4;
    v = p1 - p8 - 2 -
        (p8 + p2 - b) * ((p3 + 6) * c);
    e = v * p5 + d;
    return e;
}
```

```
...
vmov.f64 d13, #4.000000e+00
vmov.f64 d14, #1.000000e+01
vmul.f64 d10, d10, d13
vadd.f64 d11, d1, d14
vmov.f64 d12, #2.000000e+00
vmov.f64 d15, #8.000000e+00
vsub.f64 d5, d12, d5
vadd.f64 d12, d2, d15
vmls.f64 d10, d11, d9
vstr d6, [sp]
vmov.f64 d6, #6.000000e+00
...
```

## Example: register allocation

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, v, e;
    a = p4 - 5 + 9 * p1 * 7;
    c = 2 - p6 + 1;
    b = 3 * p6 * 4 - (p2 + 10) * a;
    d = (8 + p3) * p5 * 5 * p7 - p4;
    v = p1 - p8 - 2 -
        (p8 + p2 - b) * ((p3 + 6) * c);
    e = v * p5 + d;
    return e;
}
```

```
...
vmov.f64 d13, #4.000000e+00
vmov.f64 d14, #1.000000e+01
vmul.f64 d10, d10, d13
vadd.f64 d11, d1, d14
vmov.f64 d12, #2.000000e+00
vmov.f64 d15, #8.000000e+00
vsub.f64 d5, d12, d5
vadd.f64 d12, d2, d15
vmls.f64 d10, d11, d9
vstr d6, [sp] ← spill d6
vmov.f64 d6, #6.000000e+00
...
```

## Example: register allocation

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, v, e;
    a = p4 - 5 + 9 * p1 * 7;
    c = 2 - p6 + 1;
    b = 3 * p6 * 4 - (p2 + 10) * a;
    d = (8 + p3) * p5 * 5 * p7 - p4;
    v = p1 - p8 - 2 -
        (p8 + p2 - b) * ((p3 + 6) * c);
    e = v * p5 + d;
    return e;
}
```

spill of register **d6** instead of using a free register (**d13**, **d14**, **d15**)

```
...
vmov.f64 d13, #4.000000e+00 | d13
vmov.f64 d14, #1.000000e+01 | d14
vmul.f64 d10, d10, d13
vadd.f64 d11, d1, d14
vmov.f64 d12, #2.000000e+00
vmov.f64 d15, #8.000000e+00 | d15
vsub.f64 d5, d12, d5
vadd.f64 d12, d2, d15
vmls.f64 d10, d11, d9
vstr d6, [sp] ← spill d6
vmov.f64 d6, #6.000000e+00
...
```

## Example: register allocation

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, v, e;
    a = p4 - 5 + 9 * p1 * 7;
    c = 2 - p6 + 1;
    b = 3 * p6 * 4 - (p2 + 10) * a;
    d = (8 + p3) * p5 * 5 * p7 - p4;
    v = p1 - p8 - 2 -
        (p8 + p2 - b) * ((p3 + 6) * c);
    e = v * p5 + d;
    return e;
}
```

spill of register **d6** instead of using a free register (**d13**, **d14**, **d15**)

[https://bugs.llvm.org/show\\_bug.cgi?id=37073](https://bugs.llvm.org/show_bug.cgi?id=37073) (reported, caused by scheduling)

```
...
vmov.f64 d13, #4.000000e+00 | d13
vmov.f64 d14, #1.000000e+01 | d14
vmul.f64 d10, d10, d13
vadd.f64 d11, d1, d14
vmov.f64 d12, #2.000000e+00
vmov.f64 d15, #8.000000e+00 | d15
vsub.f64 d5, d12, d5
vadd.f64 d12, d2, d15
vmls.f64 d10, d11, d9
vstr d6, [sp] ← spill d6
vmov.f64 d6, #6.000000e+00
...
```

## Example: excessive copying

Clang:

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, e, f;
    a = d = b = p1 + p3;
    e = p3 + p2 * 8.7 +
        ((p4 - p4) * (a * p7) - 1.7);
    c = b - d - (p2 - 7.4) * 8.1;
    f = e - p2 * p3 -
        d * (4.6 + c) * p4 +
        p5 * p6 - p7 - p8 + p1;
    return f;
}
```

GCC:

## Example: excessive copying

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, e, f;
    a = d = b = p1 + p3;
    e = p3 + p2 * 8.7 +
        ((p4 - p4) * (a * p7) - 1.7);
    c = b - d - (p2 - 7.4) * 8.1;
    f = e - p2 * p3 -
        d * (4.6 + c) * p4 +
        p5 * p6 - p7 - p8 + p1;
    return f;
}
```

Clang:

GCC:

```
vadd.f64 d8, d0, d2
vmul.f64 d9, d8, d6
vsub.f64 d10, d3, d3
...
```

## Example: excessive copying

```
double fn1(double p1, double p2,
           double p3, double p4, double p5,
           double p6, double p7, double p8)
{
    double a, b, c, d, e, f;
    a = d = b = p1 + p3;
    e = p3 + p2 * 8.7 +
        ((p4 - p4) * (a * p7) - 1.7);
    c = b - d - (p2 - 7.4) * 8.1;
    f = e - p2 * p3 -
        d * (4.6 + c) * p4 +
        p5 * p6 - p7 - p8 + p1;
    return f;
}
```

caused by scheduling?

Clang:

```
vadd.f64 d8, d0, d2
vmov.f64 d10, d7
vmov.f64 d7, d5
vmov.f64 d5, d4
vmov.f64 d4, d3
vmul.f64 d11, d8, d6
vsub.f64 d12, d4, d4
...
```

GCC:

```
vadd.f64 d8, d0, d2
vmul.f64 d9, d8, d6
vsub.f64 d10, d3, d3
...
```

5 copies in 30 instrs

1 copy in 26 instrs

## A GCC example

```
int N;
long fn1(void) {
    short i;
    long a;
    i = a = 0;
    while (i < N)
        a -= i++;
    return a;
}
```

GCC:

Loop vectorized...

## A GCC example

```
int N;
long fn1(void) {
    short i;
    long a;
    i = a = 0;
    while (i < N)
        a -= i++;
    return a;
}
```

GCC:

Loop vectorized...

late Feb 2018 no longer vectorized, 1.8 × slowdown

## A GCC example

```
int N;
long fn1(void) {
    short i;
    long a;
    i = a = 0;
    while (i < N)
        a -= i++;
    return a;
}
```

GCC:

Loop vectorized...

late Feb 2018 no longer vectorized, 1.8 × slowdown

[https://gcc.gnu.org/bugzilla/show\\_bug.cgi?id=84986](https://gcc.gnu.org/bugzilla/show_bug.cgi?id=84986) (fixed)

## A CompCert example

```
int f(int a) {          Clang:  
    if (a + 42)  
        ;  
    else  
        ;  
    return 1;  
}
```

# A CompCert example

```
int f(int a) {  
    if (a + 42)  
        ;  
    else  
        ;  
    return 1;  
}
```

Clang:

```
    mov r0, #1
```

CompCert:

```
    add r0, r0, #42  
    mov r0, #1
```

(fixed)

# Results

<https://github.com/gergo-/missed-optimizations>

## Some missed optimizations found

	total	reported	fixed
GCC	18	7	2
Clang	9	4	2
CompCert	9	3	3

- ▶ generally treated as low priority by developers
- ▶ many duplicates
- ▶ **main limitation:** loops (C-Reduce generates infinite loops)

## Summary

- ▶ compare compilers' outputs to find missed optimizations
- ▶ automated toolchain finds minimal test cases
- ▶ issues found in GCC, Clang, CompCert

<https://github.com/gergo-/missed-optimizations>

## Summary

- ▶ compare compilers' outputs to find missed optimizations
- ▶ automated toolchain finds minimal test cases
- ▶ issues found in GCC, Clang, CompCert

<https://github.com/gergo-/missed-optimizations>

Thank you for your attention

This research was partially supported by ITEA 3 project no. 14014, ASSUME.