Hardware Memory Tagging makes C++ Memory-Safer

C/C++ Memory Safety is a mess:

- > 20% of all C/C++ bugs and
- > 50% of all vulnerabilities are use-after-free or buffer-overflow (heap,stack,globals)

AddressSanitizer (ASAN) is not enough

- Hard to use in production:
 2x overhead in CPU, RAM, Code Size
- Weak as a mitigation:
 easy to bypass redzones & quarantine

A better scheme: Memory Tagging

- Two types of tags
- Every aligned 16 bytes of memory have a 4-bit tag
- Every 64-bit pointer has a 4-bit tag (in the top byte)
- LD/ST instructions SEGV on tag mismatch
- New instructions to set the tags

Existing Implementations

SPARC ADI: 4-bit tag per 64-bytes

Available in *hardware* since ~2016, SPARC M7/M8 Heap-only; 5%-20% RAM overhead due to alignment

HWASAN (Hardware ASAN): 8-bit tags per 16 bytes Software (compiler) instrumentation, similar to ASAN Uses AArch64-only feature "top-byte-ignore"
Limited applicability on x86_64 (needs to see all loads/stores)

ARM Memory Tagging Extension (MTE)

Announced by ARM on 2018-09-17

```
char *P = new char[20]; //0xa0007ffffffff1240
// Pointer P & memory P[0:31] are tagged with □
-32:-17 -16:-1 0:15 16:31 32:47 48:64

P[32] = ... // heap-buffer-overflow □ ≠ □
delete [] P; // Memory is retagged □ ⇒ □
-32:-17 -16:-1 0:15 16:31 32:47 48:64

P[0] = ... // heap-use-after-free □ ≠ □
```

MTE is better than ASAN

- RAM overhead: 3%-5% (vs ~ 2x)
- CPU overhead: expected low-single-digit %
- Detects buffer overflows far from bounds
- Detects use-after-free long after deallocation

Usage Models:

- Regular testing/fuzzing (faster ASAN)
- Crowdsourced testing in production
- Always-on security mitigation

Desired Compiler Optimizations

Malloc now zero-fills and sets memory tags

```
struct S { int64_t a, b; };
// Redundant initialization after new()
S *foo() { return new S{0, 0}; }
// Redundant initialization/tagging in new()
S *bar() { return new S{1, 2}; }
```

- Malloc costs more: need malloc-to-stack
- Compact & fast stack instrumentation

```
// x0 := sp w/ random tag
void f() {
             irg x0, sp
                                // Tag the memory
             stg [x0]
  int a;
 bar(&a);
                                // Before exit, restore the default tag
             stg [sp], #16
int f() {
             irg x0, sp
 int a = 42; mov w8, #42
 bar(&a); stgp x8, xzr, [x0] // Initialize & tag
  return a;
                                // ld/st [SP, #imm] does not check tags
             ldr w0, [sp]
             stg [sp], #16
```

- Static inter-procedural stack safety analysis
- O Do we need to tag a local variable?
- o Is buffer overflow possible?
- o Is use-after-return possible?
- Optional): is use of uninitialized value possible?

Homework: Ask your CPU vendor to implement Memory Tagging

Kostya Serebryany kcc@google.com kayseesee@ Evgenii Stepanov eugenis@google.com
Vlad Tsyrklevich vtsyrklevich@google.com vlad902@ E
See also: https://arxiv.org/pdf/1802.09517.pdf