

# Optimization Remarks Update

Francis Visoiu Mistrih

# Optimization Remarks

18			<code>int main(int argc, const char *argv[]) {</code>
	prologue		4120 stack bytes in function
	epilog		
	asm-printer		25 instructions in function
19			<code>int result = foo(argc);</code>
		inline	<code>foo</code> inlined into main with (cost=65, threshold=225)
20			<code>return 0;</code>
21			<code>}</code>

opt-viewer.py

# Challenges

- YAML doesn't scale well
- Discovering remarks
- Processing remarks

**\_\_TEXT**  
0.1 GB



**\_\_DWARF**  
8.5 GB

**Remarks (YAML)**  
20.7 GB

**\_\_TEXT**  
0.1 GB



**\_\_DWARF**  
8.5 GB

**Remarks (YAML)**  
20.7 GB


- Significant size impact



- Significant size impact
- Slow to parse: **~4 minutes** for 27 million remarks




- Significant size impact
- Slow to parse: **~4 minutes** for 27 million remarks
- Compile-time impact: **5% slowdown**




- Significant size impact
- Slow to parse: **~4 minutes** for 27 million remarks
- Compile-time impact: **5% slowdown**






- Significant size impact **12x smaller**
- Slow to parse: **~4 minutes** for 27 million remarks
- Compile-time impact: **5% slowdown**



- Significant size impact **12x smaller**
- Slow to parse: ~~~4 minutes~~ **27 seconds (8x faster)**
- Compile-time impact: **5% slowdown**



- Significant size impact **12x smaller**
- Slow to parse: ~~~4 minutes~~ **27 seconds (8x faster)**
- Compile-time impact: ~~5% slowdown~~ **2% slowdown**

# Extra Benefits

- Well-known format in the community
- Testing and tooling like `llvm-bc-analyzer`
- Versioning
- Flexible metadata

# libRemarks


```
#include "llvm-c/Remarks.h"


LLVMRemarkParserRef Parser = LLVMRemarkParserCreateYAML(Buf, Size);
LLVMRemarkEntryRef Remark = NULL;
while ((Remark = LLVMRemarkParserGetNext(Parser))) {
    // use Remark
}
```

# libRemarks


```
#include "llvm-c/Remarks.h"

LLVMRemarkParserRef Parser = LLVMRemarkParserCreateBitstream(Buf, Size);
LLVMRemarkEntryRef Remark = NULL;
while ((Remark = LLVMRemarkParserGetNext(Parser))) {
    // use Remark
}
```

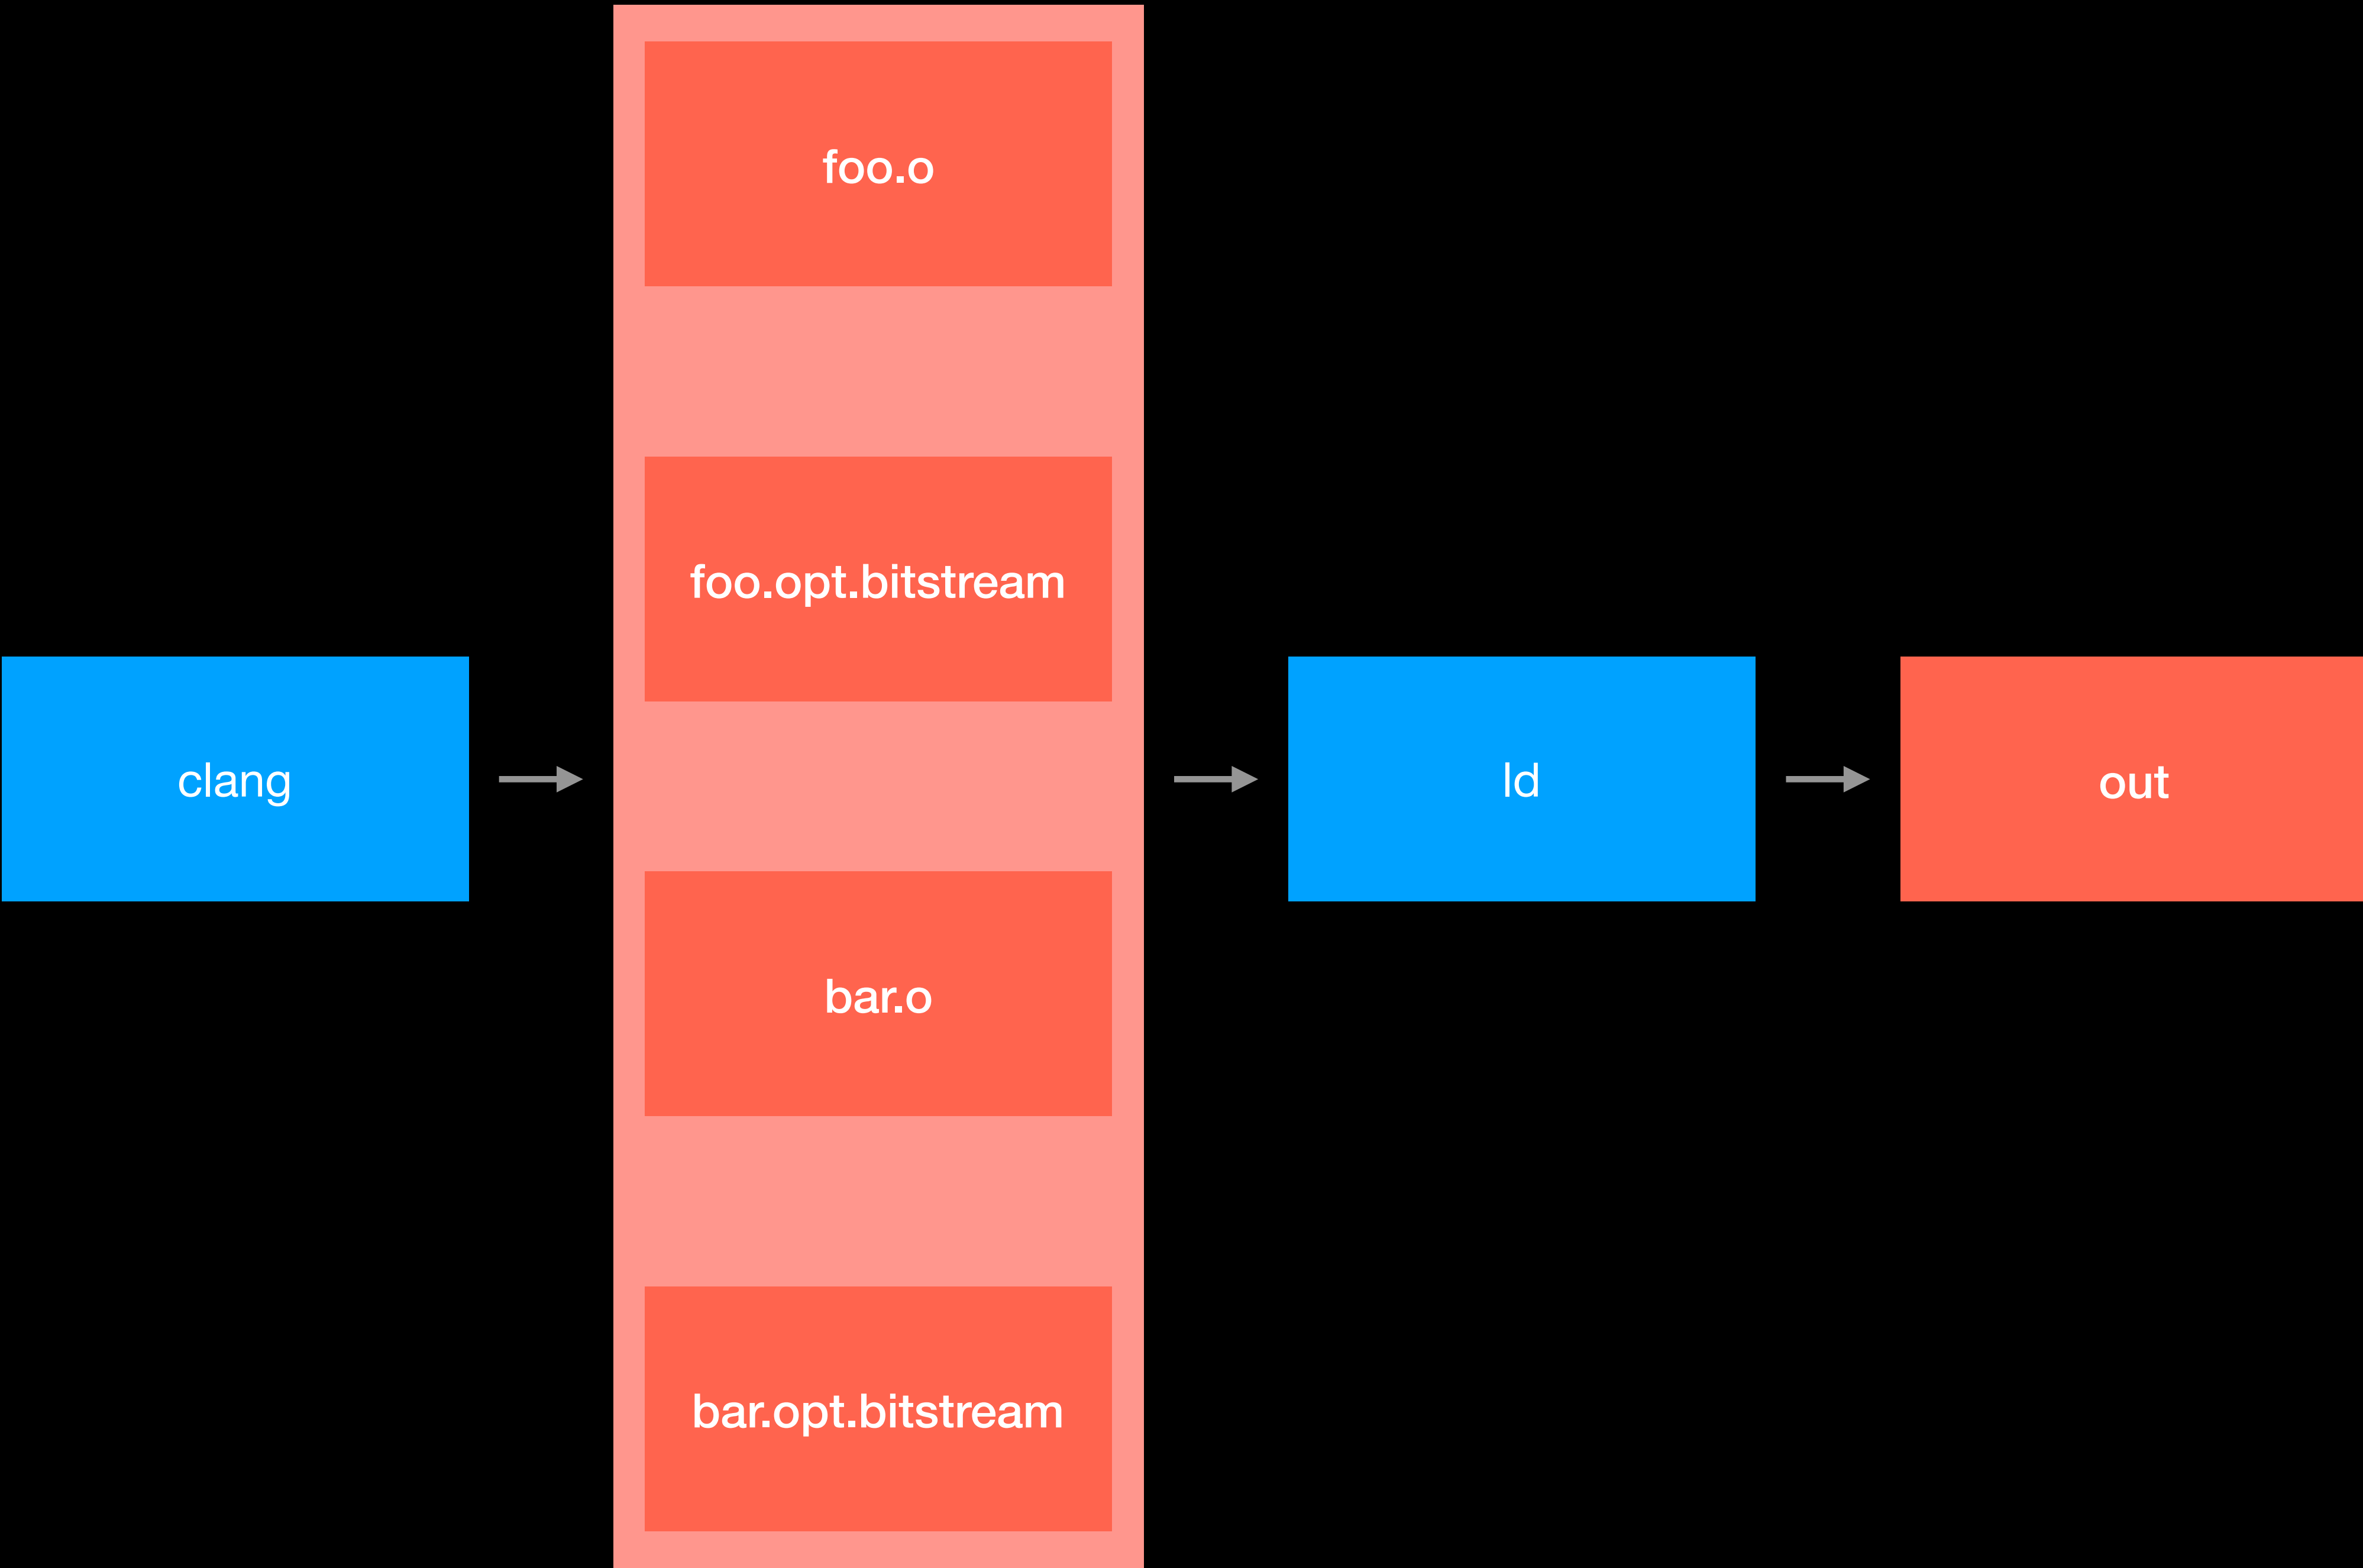


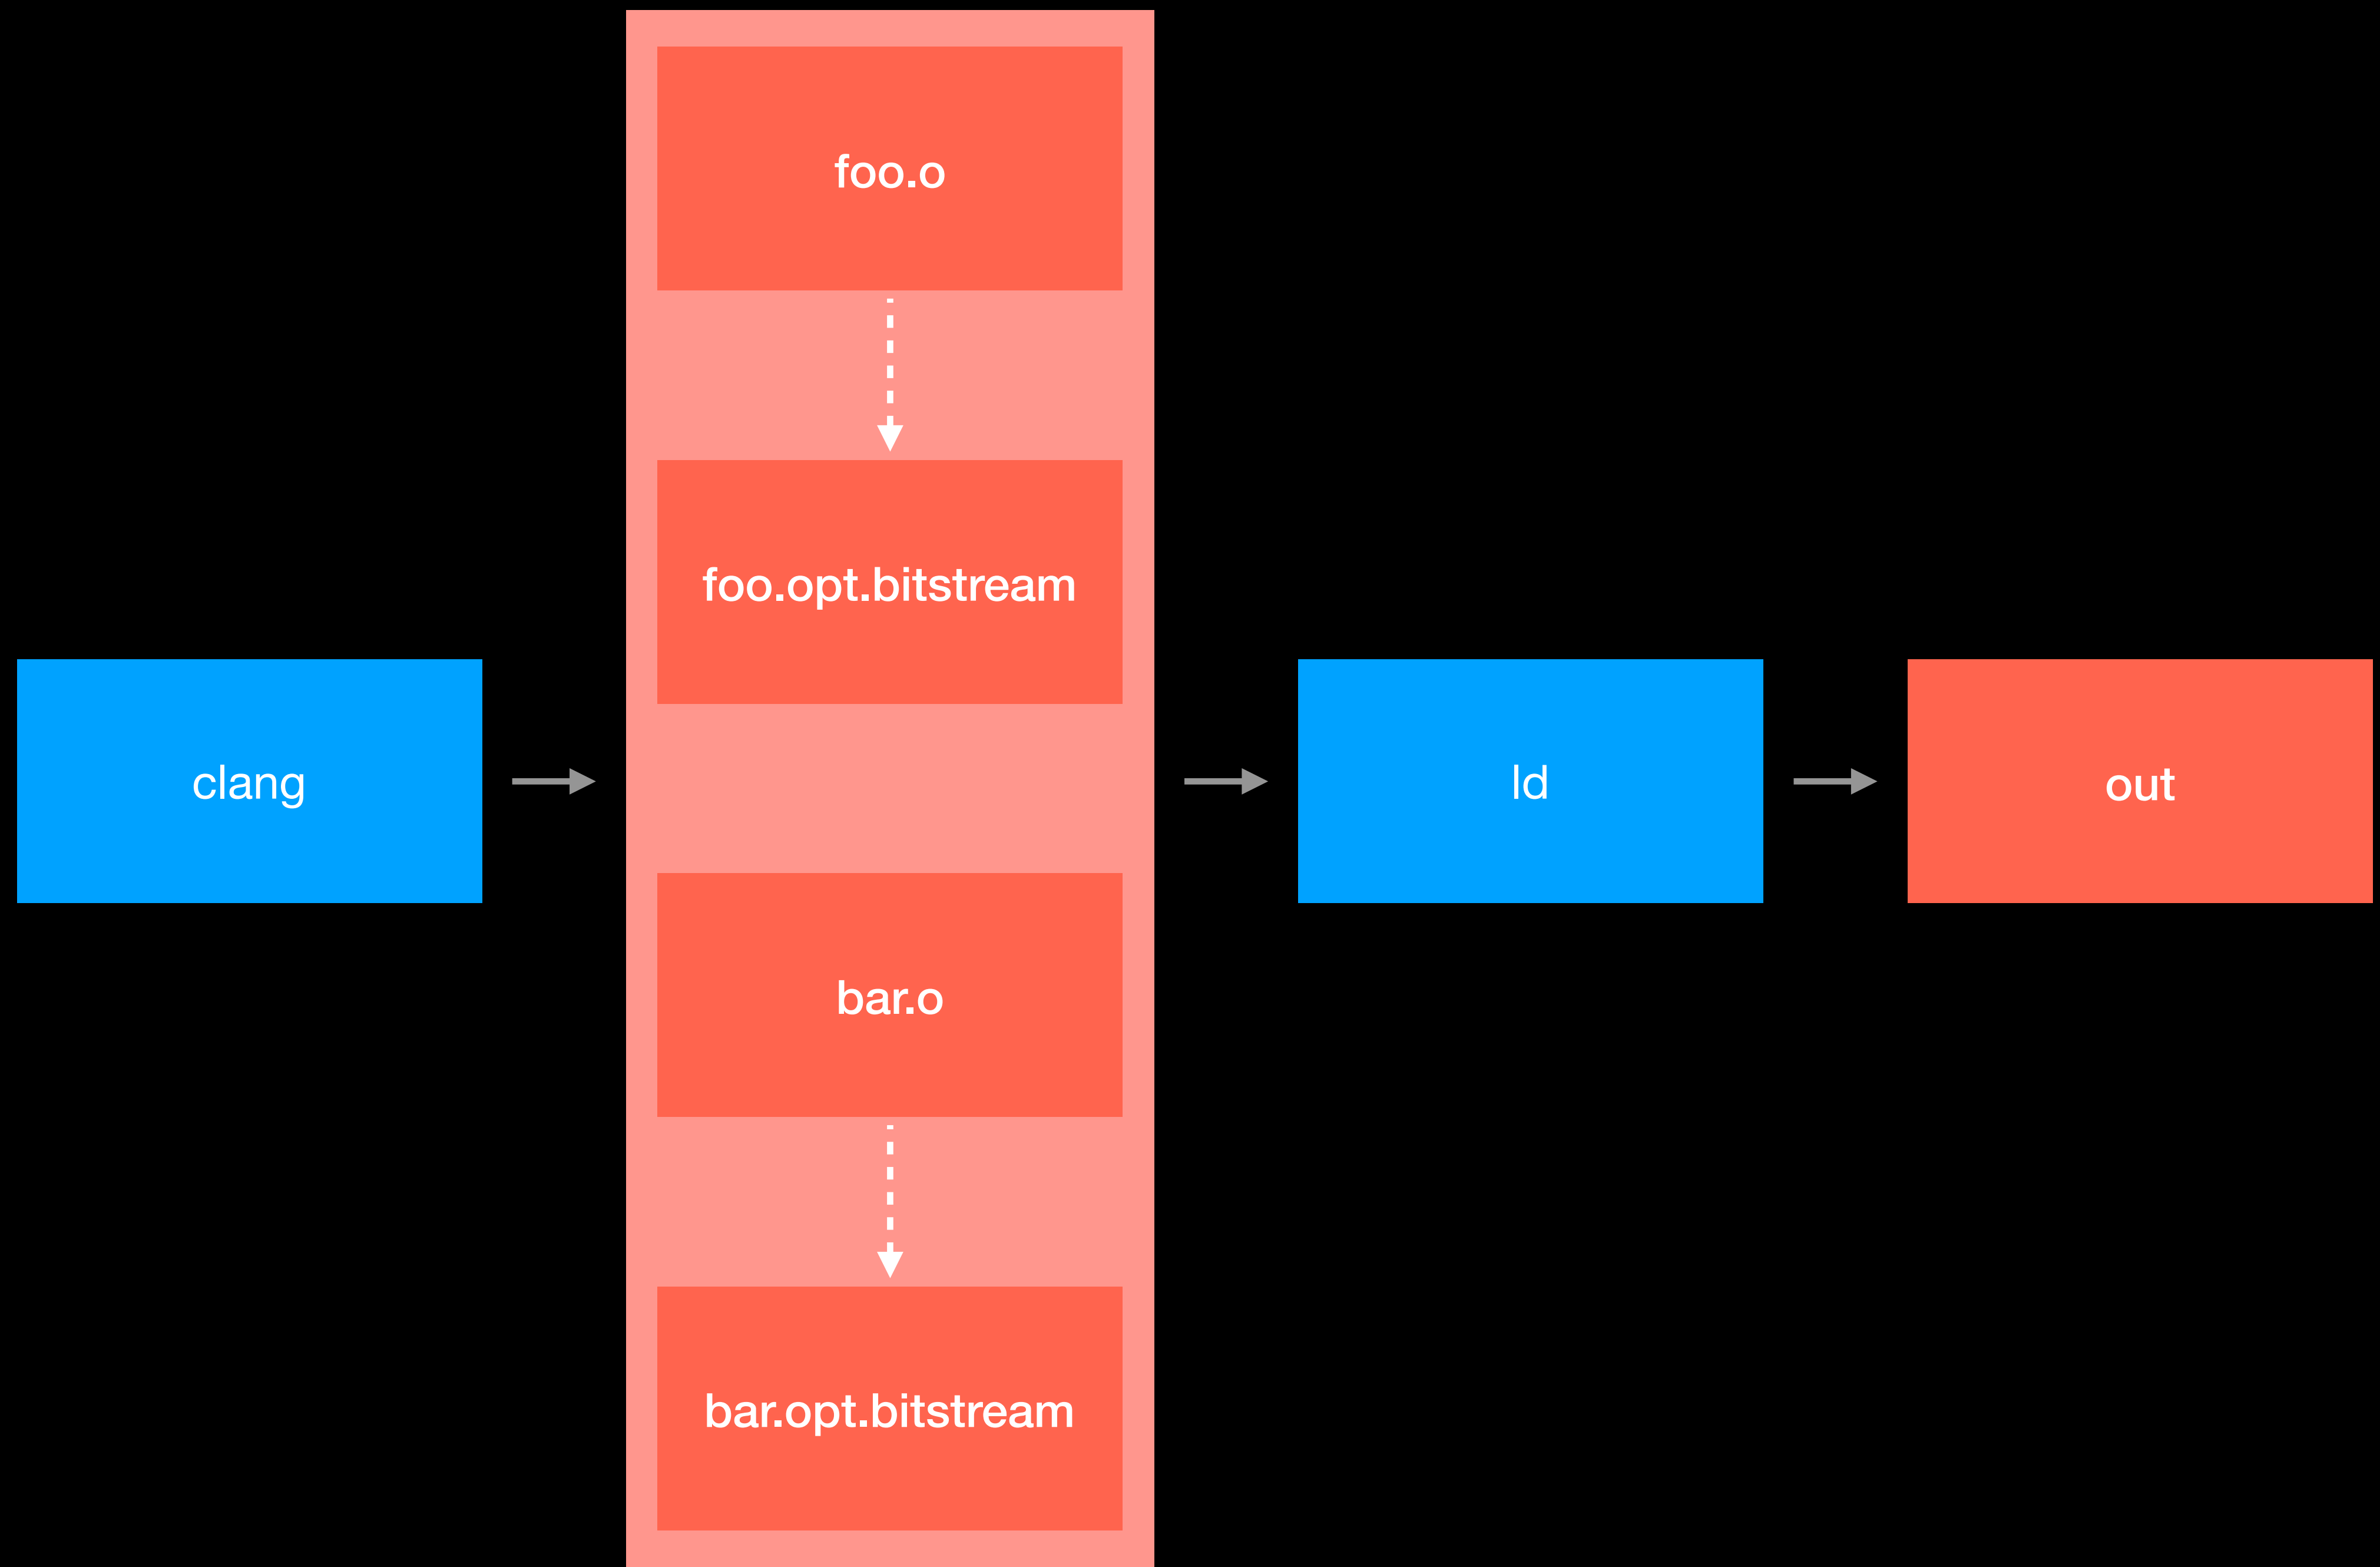


?









foo.o



foo.opt.bitstream

bar.o



bar.opt.bitstream

foo.o

\_\_remarks

52	4d	52	4b	01
08	00	00	17	00
00	00	07	01	b2
40	b4	42	39	d0
43	38	3c	3c	c1
28	bc	83	3b	d0
43	38	a4	83	3b
94	83	3c	80	41
3a	b8	83	39	bc
c3	41	80	38	06
08	14	22	1e	9a
61	16	e8	41	1e
d2	c1	1d	ce	01
0c	e8	21	1c	c4
81	1d	ca	41	71
40	1f	1c	a2	14
f0	81	1e	ca	41
1e	dc	21	1c	d8



foo.opt.bitstream

foo.o

\_\_remarks

52	4d	52	4b	01
08	00	00	17	00
00	00	07	01	b2
40	b4	42	39	d0
43	38	3c	3c	c1
28	bc	83	3b	d0
43	38	a4	83	3b
94	83	3c	80	41

<file path>


d2	c1	1d	ce	01
0c	e8	21	1c	c4
81	1d	ca	41	71
40	1f	1c	a2	14
f0	81	1e	ca	41
1e	dc	21	1c	d8

foo.opt.bitstream

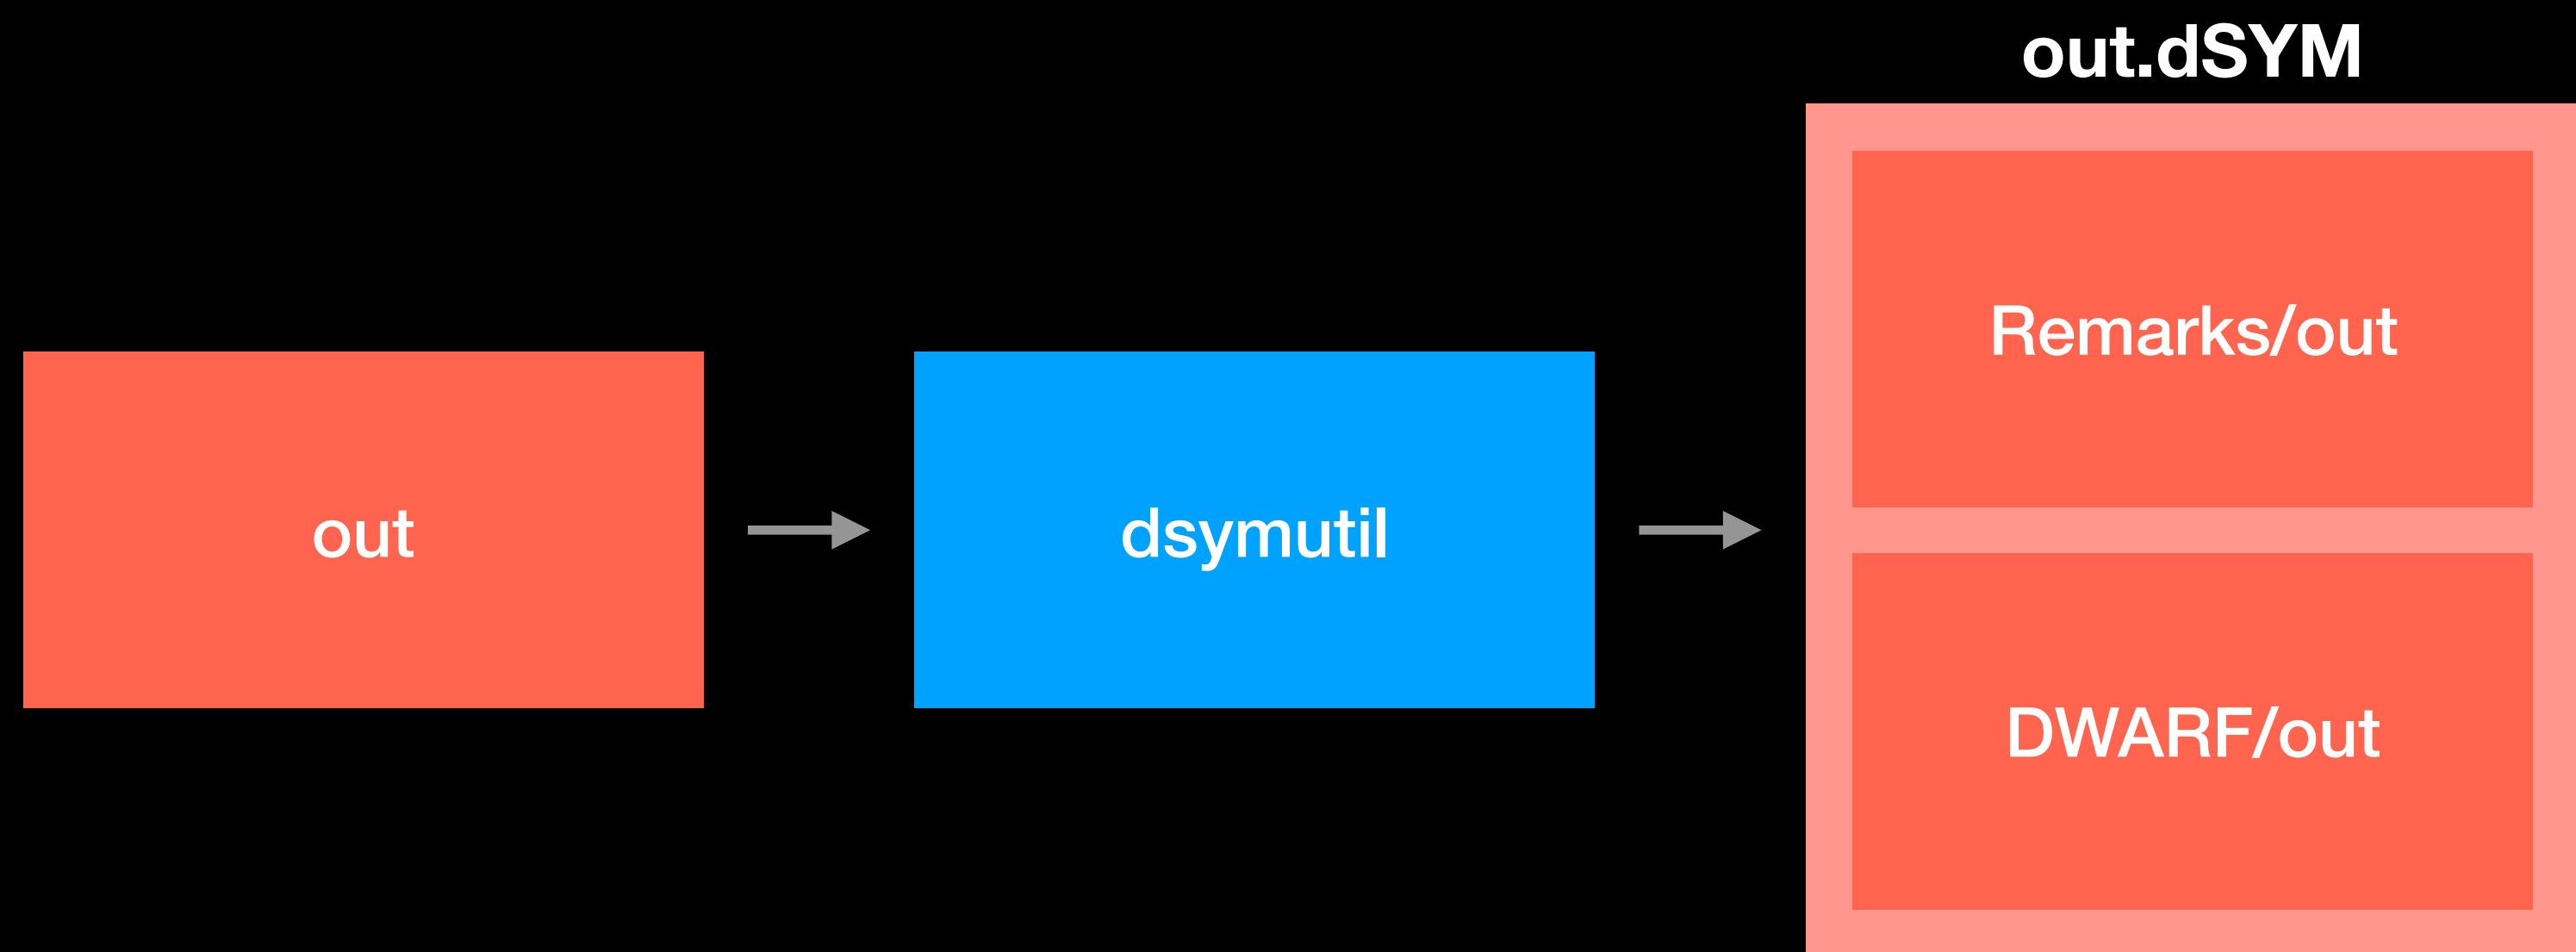
# Remarks Section

- Only in object files
- Not included in the linked binary
- Accessible through the Debug Map
- clang's remarks section: 0.27GB









# Get Involved

- Python bindings for integration with `opt-viewer.py`
- `llvm-remarkutil`: convert, merge, extract, etc.
- Faster lookup: build an index
- Remark classes: verbosity, accuracy, etc.
- Tablegen: better remark structure
- Per-remark documentation: possible actions, reasons, etc.

**Thanks!**