


Memoro

Scaling an LLVM-Based Heap Profiler



 **Thierry Treyer**
Performance &
Capacity Intern

 **Mark Santaniello**
Performance &
Capacity Engineer

James Larus
EPFL IC School Dean

EPFL


```
vector<BigT> getValues(  
    map<Id, BigT>& largeMap,  
    vector<Id>& keys) {  
  
    vector<BigT> values;  
    values.reserve(largeMap.size());  
  
    for (const auto& key: keys)  
        values.emplace_back(largeMap[key]);  
  
    return values;  
}
```

40 GiB


of DRAM wasted per server

```
vector<BigT> getValues(  
    map<Id, BigT>& largeMap,  
    vector<Id>& keys) {  
  
    vector<BigT> values;  
    values.reserve(largeMap.size());  
  
    for (const auto& key: keys)  
        values.emplace_back(largeMap[key]);  
  
    return values;  
}
```

```
vector<BigT> getValues(  
    map<Id, BigT>& largeMap,  
    vector<Id>& keys) {  
  
    vector<BigT> values;  
    values.reserve(largeMap.size());  
  
    for (const auto& key: keys)  
        values.emplace_back(largeMap[key]);  
  
    return values;  
}
```


```
vector<BigT> getValues(  
    map<Id, BigT>& largeMap,  
    vector<Id>& keys) {  
  
    vector<BigT> values;  
    values.reserve(keys.size());  
  
    for (const auto& key: keys)  
        values.emplace_back(largeMap[key]);  
  
    return values;  
}
```

LLVM-Based Profiler



LLVM-Based Profiler


Manipulate the IR



LLVM-Based Profiler

Infrastructure

Manipulate the IR




LLVM-Based Profiler

Collecting and Displaying data

Infrastructure

Manipulate the IR



Memoro +

Run-Time Overhead ♦


Visualizer ♦

Open Challenges ♦


Overview




Overview



Source Code



Compile



Run



Analyze


○ No modification

Overview


INSTRUMENTATION PASS (LLVM)



Source Code



Compile



Run



Analyze

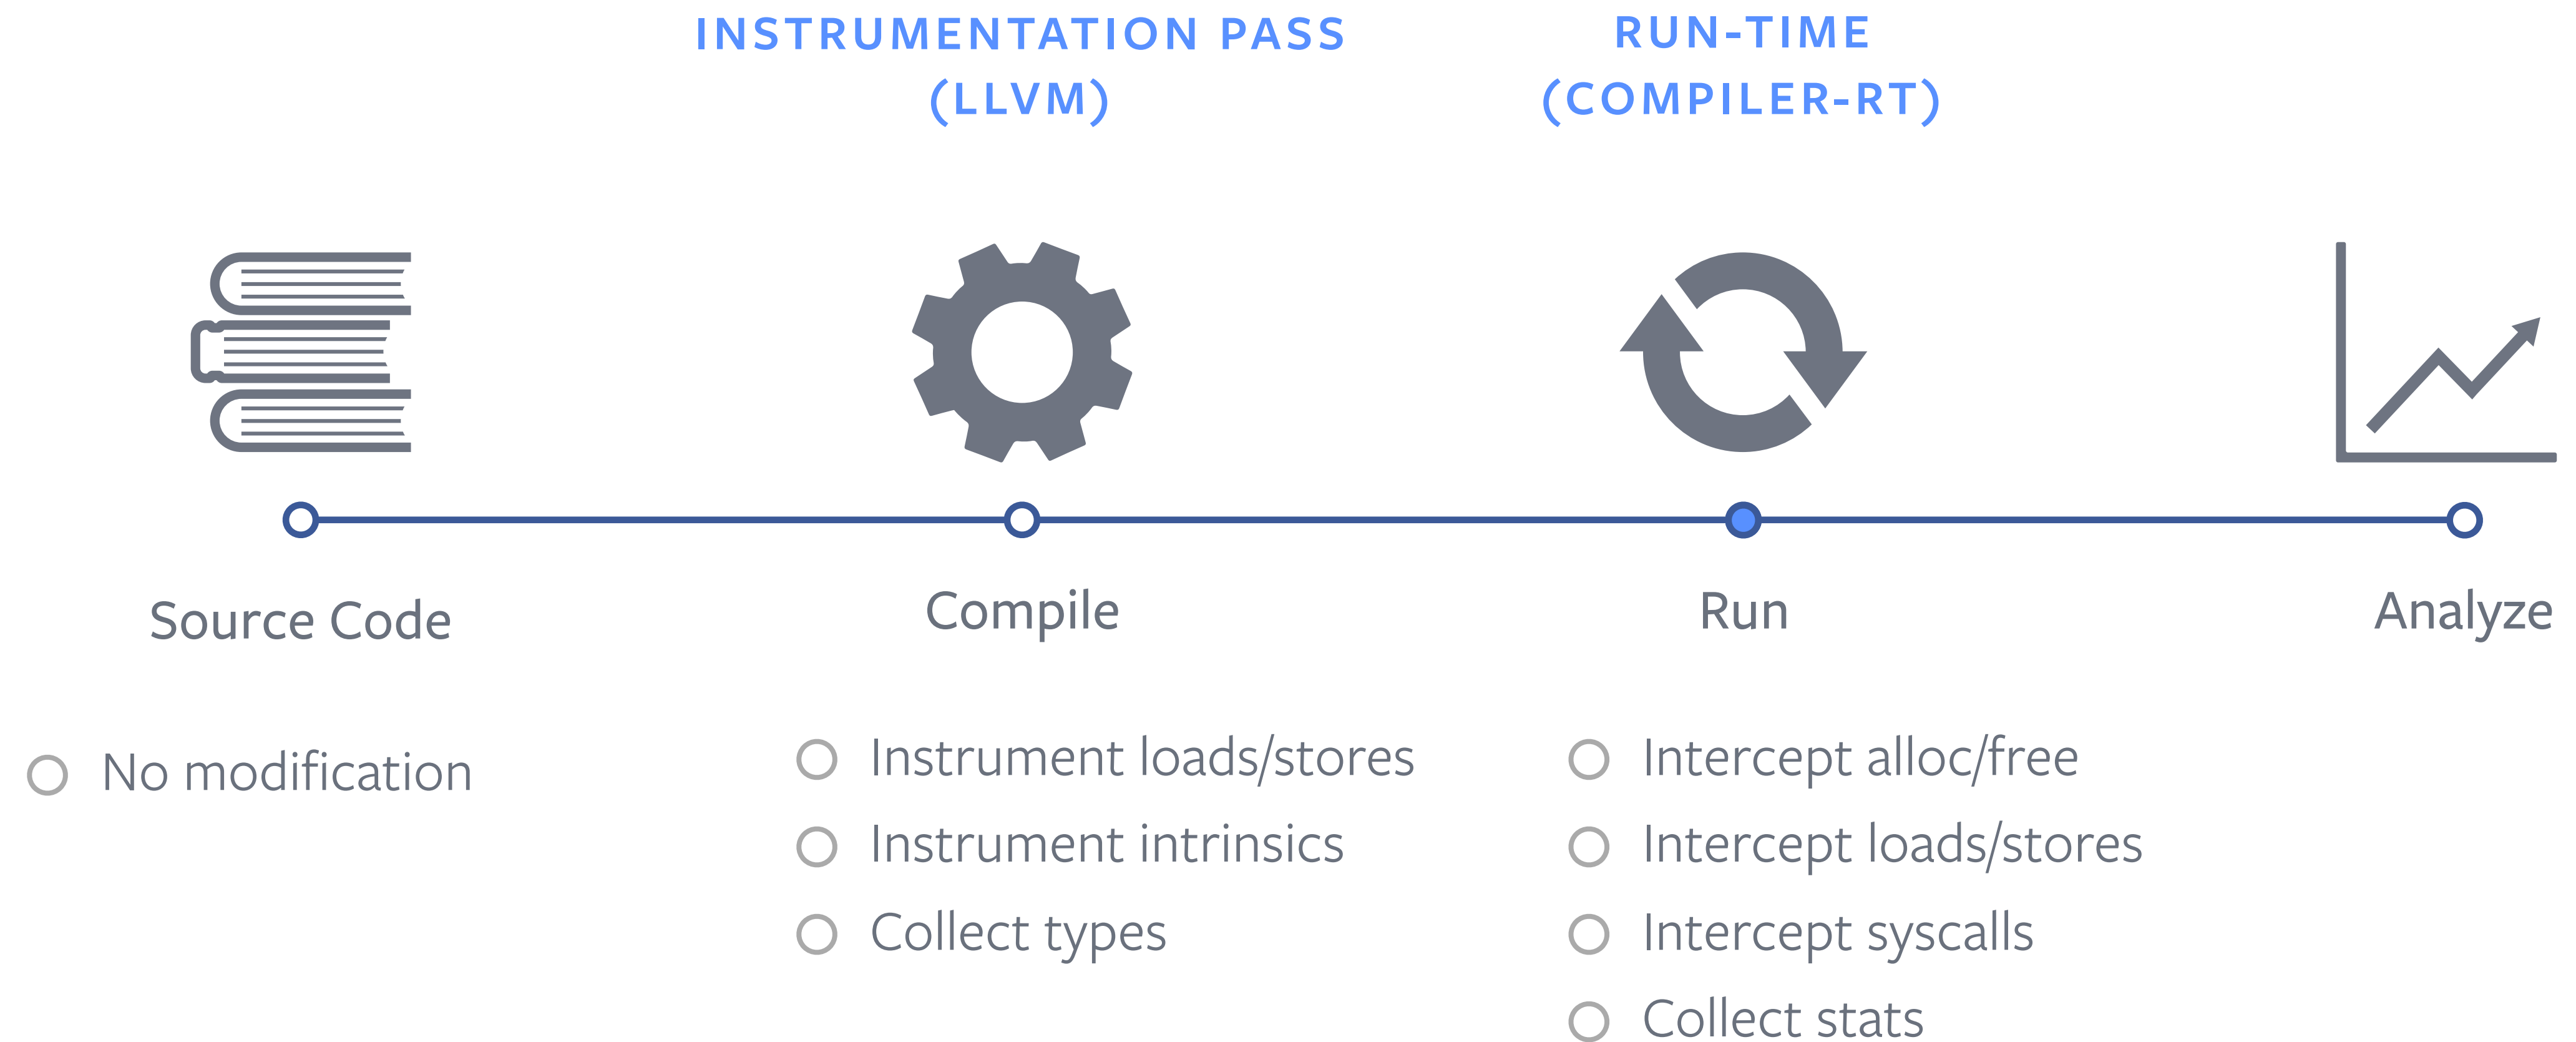
○ No modification

○ Instrument loads/stores

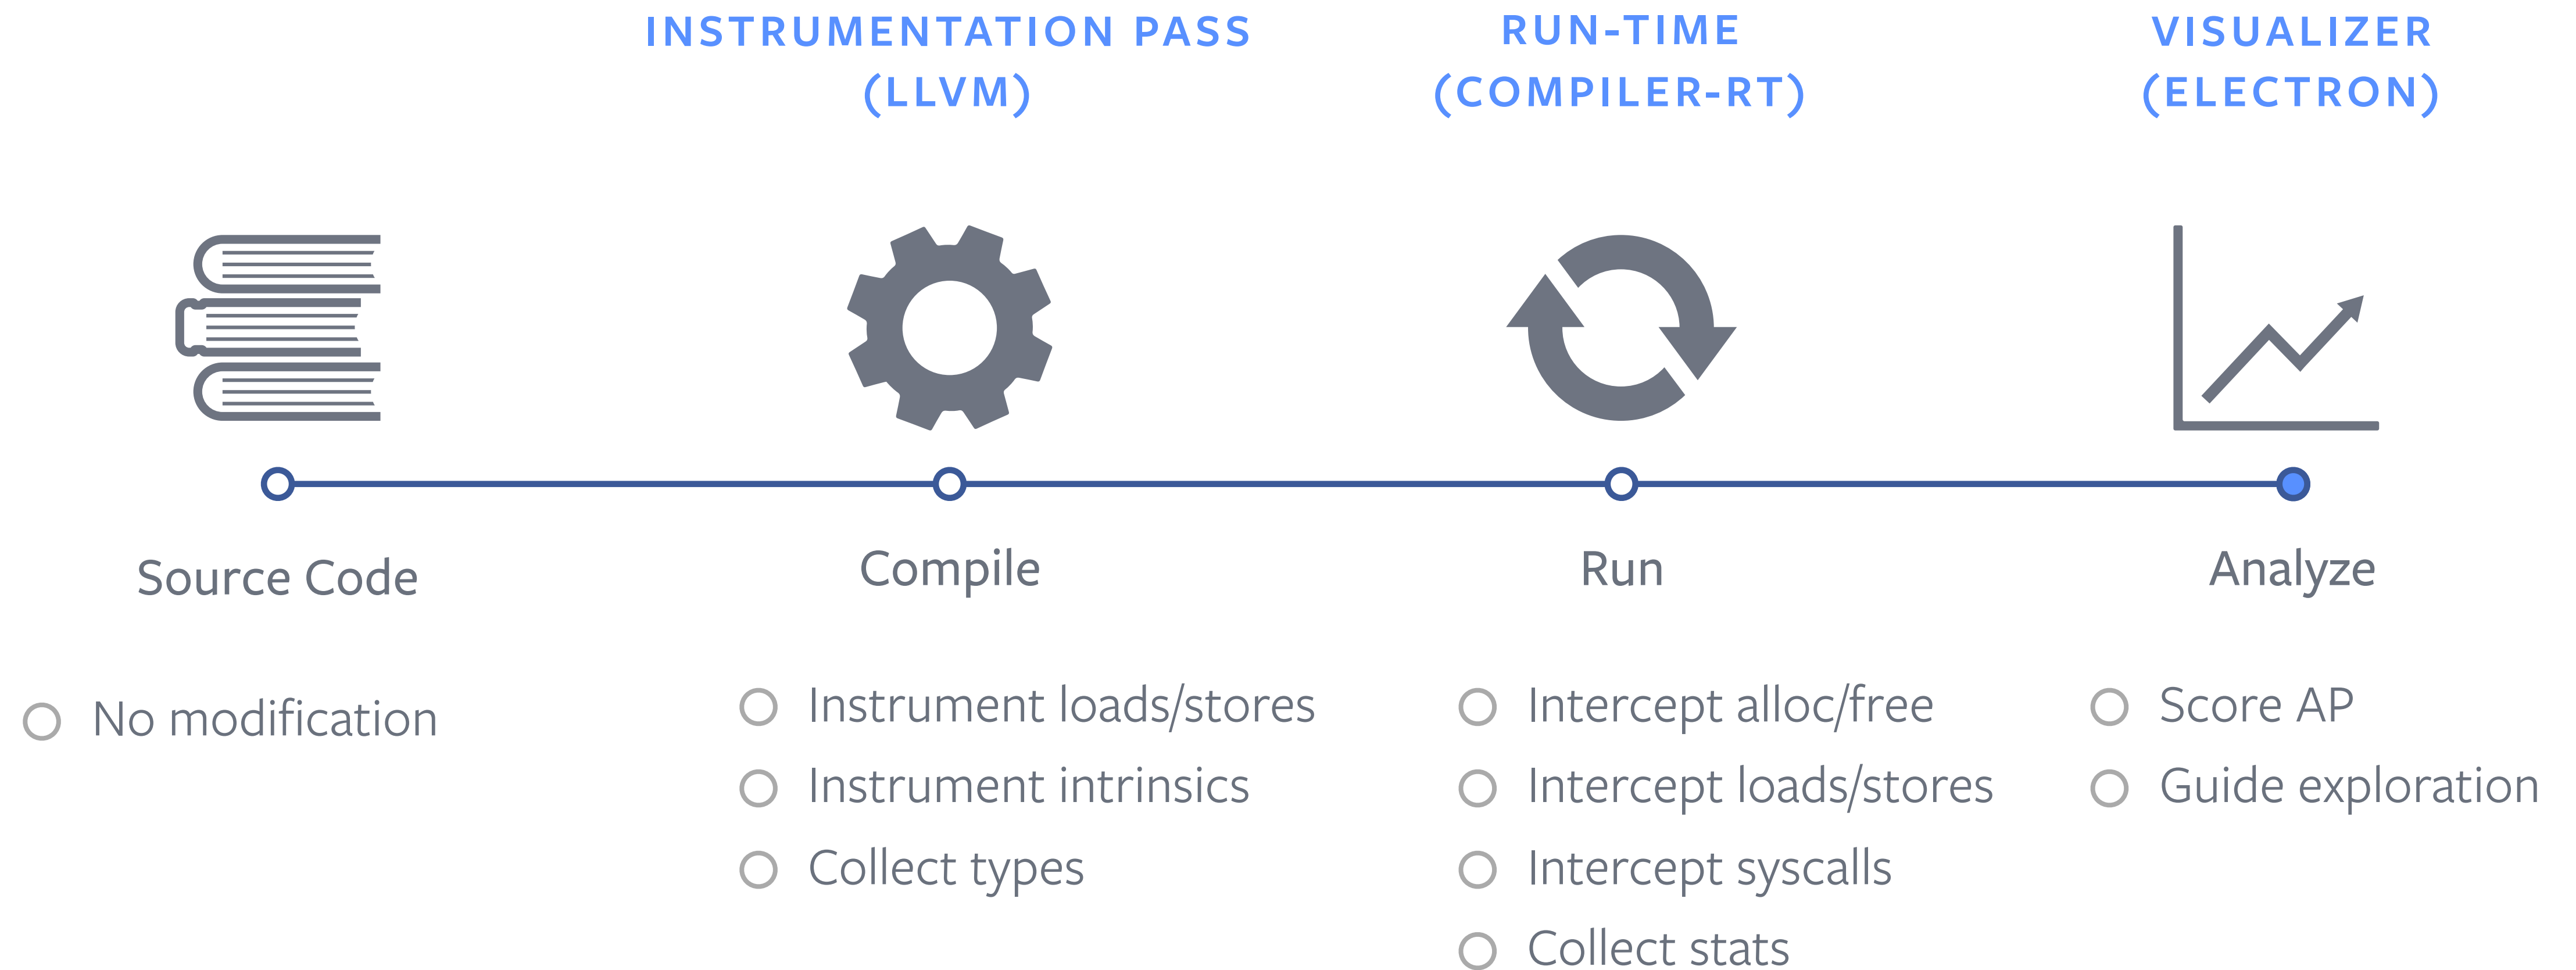
○ Instrument intrinsics

○ Collect types

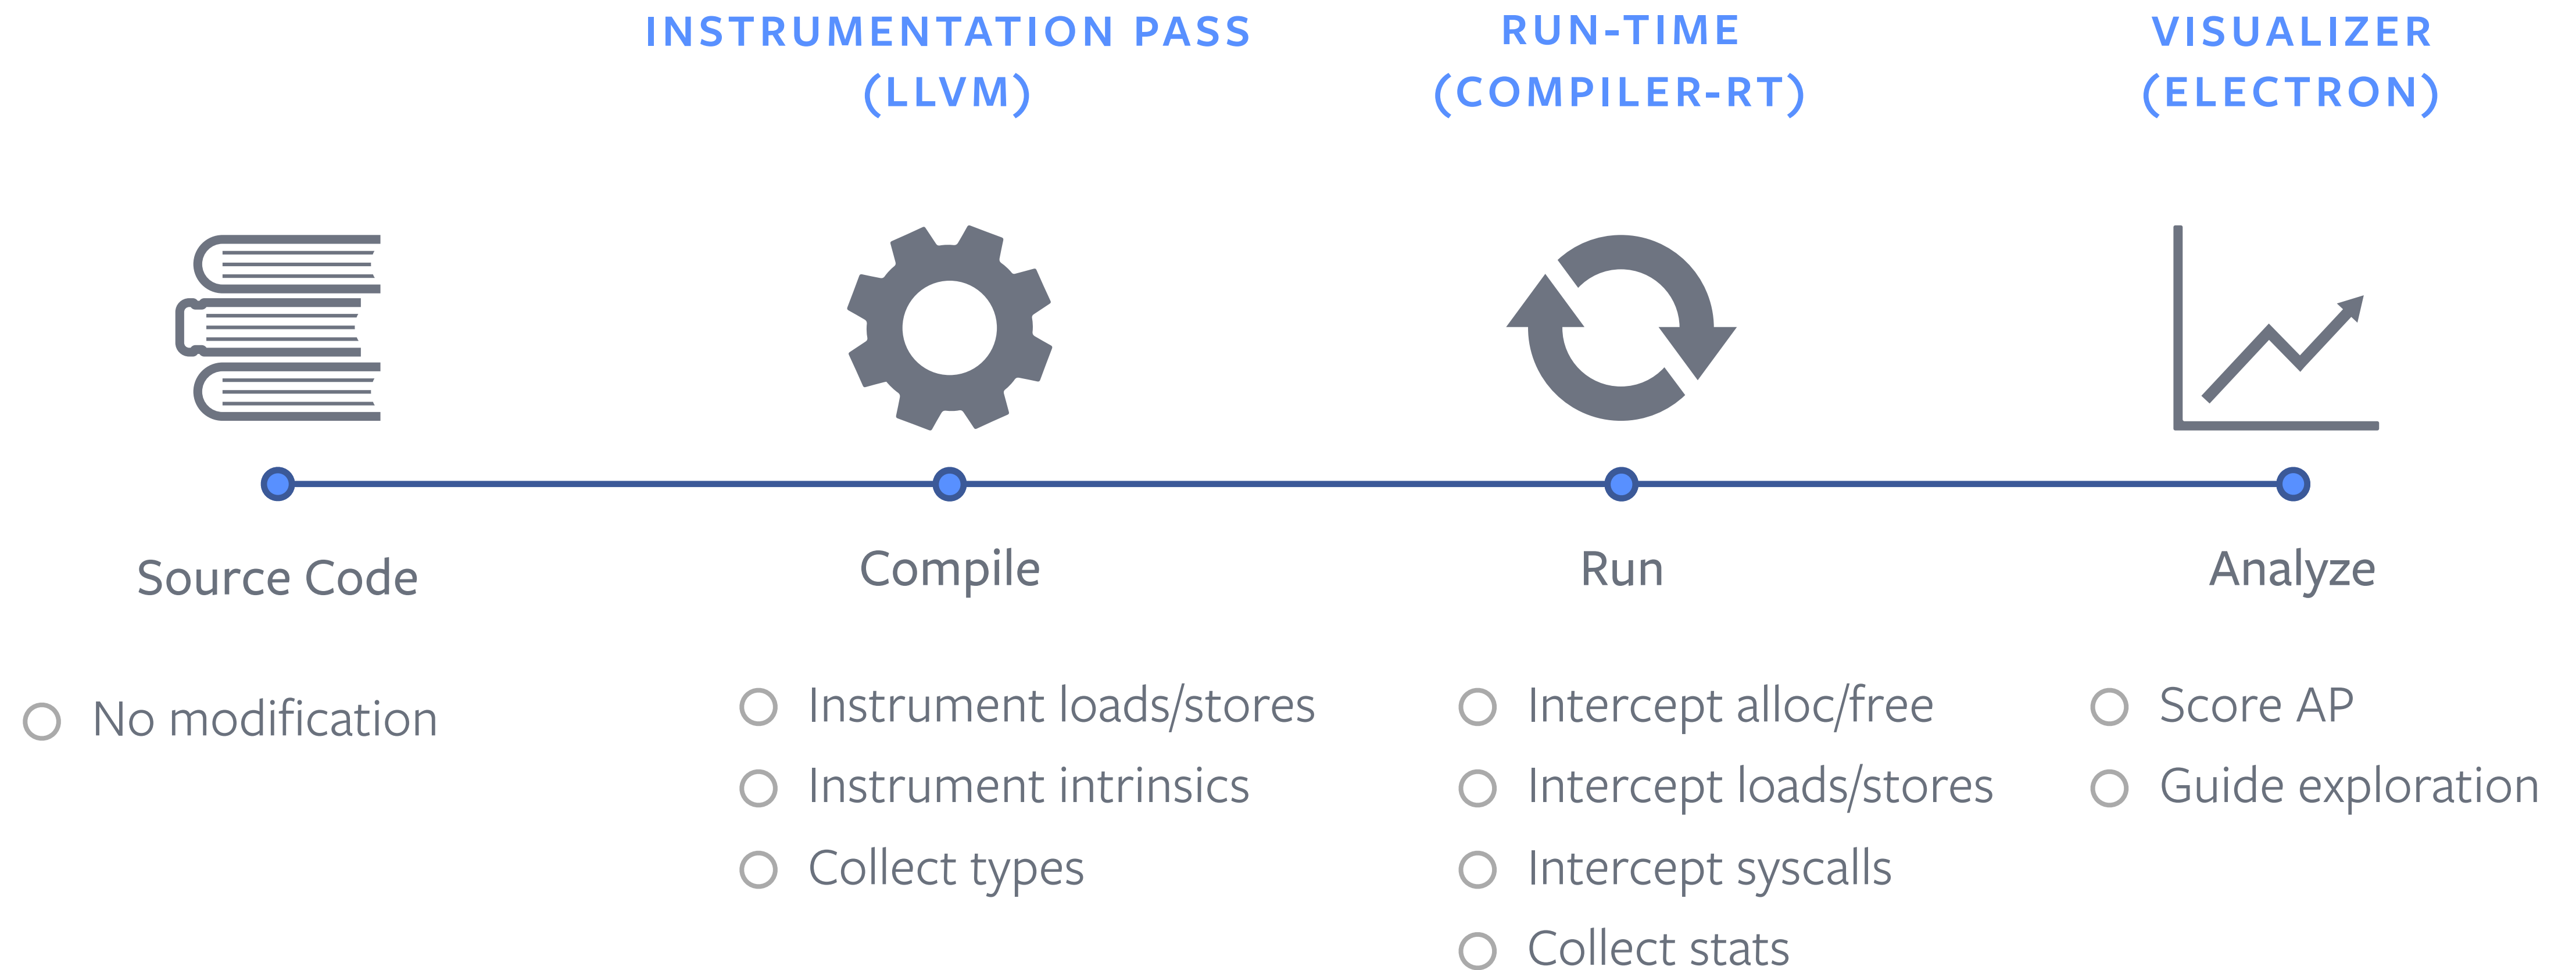
Overview



Overview



Overview



Memoro +

Run-Time Overhead ♦

Visualizer ♦

Open Challenges ♦

Memoro +

Run-Time Overhead ♦

Visualizer ♦

Open Challenges ♦

1,000x

slowdown due to Memoro's run-time

Run-Time Sampling

```
int sample_count = 0;

void interceptLoadStore(...) {
    // Sample accesses
    if (sample_count++ % access_sampling_rate != 0)
        return;

    /* Process access... */
}
```

Run-Time Sampling

```
THREADLOCAL int sample_count = 0;
```

```
void interceptLoadStore(...) {  
    // Sample accesses  
    if (sample_count++ % access_sampling_rate != 0)  
        return;  
  
    /* Process access... */  
}
```

Power to the user!

```
MEMORO_OPTIONS="..." ./myapp
```

```
– access_sampling_rate
```

```
– ...
```

```
// Public API: memoro_interface.h
```

```
#include <memoro_interface.h>
```

```
void foo(...) {
```

```
    MemoroFlags *mflags = memoro::getFlags();
```

```
    mflags->access_sampling_rate = 50;
```


```
    /* ... */
```

```
}
```






Time spent by address type

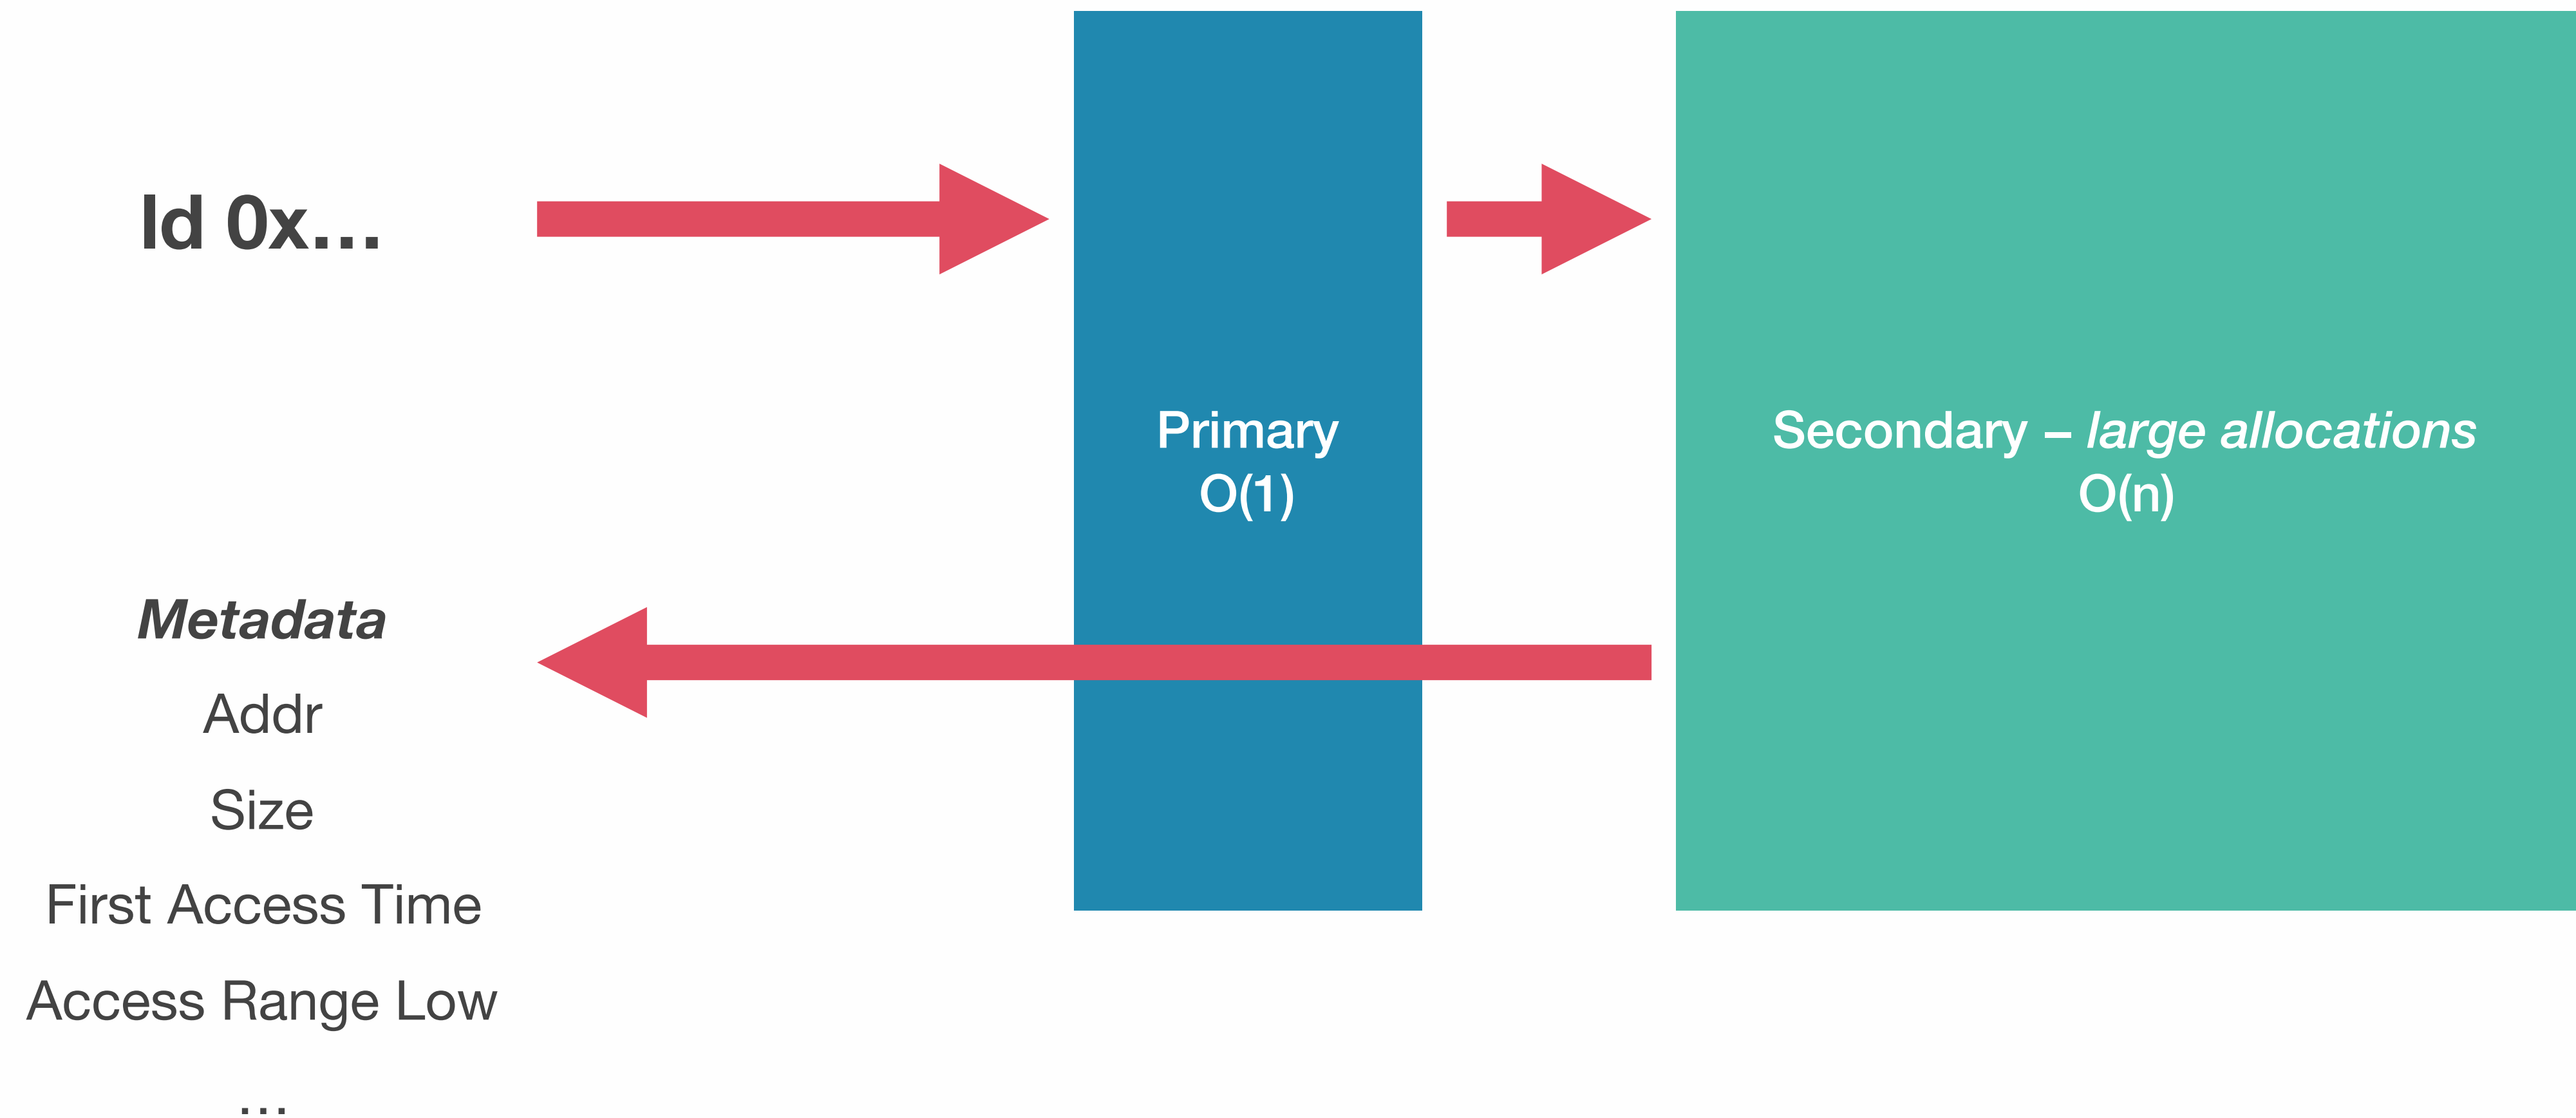




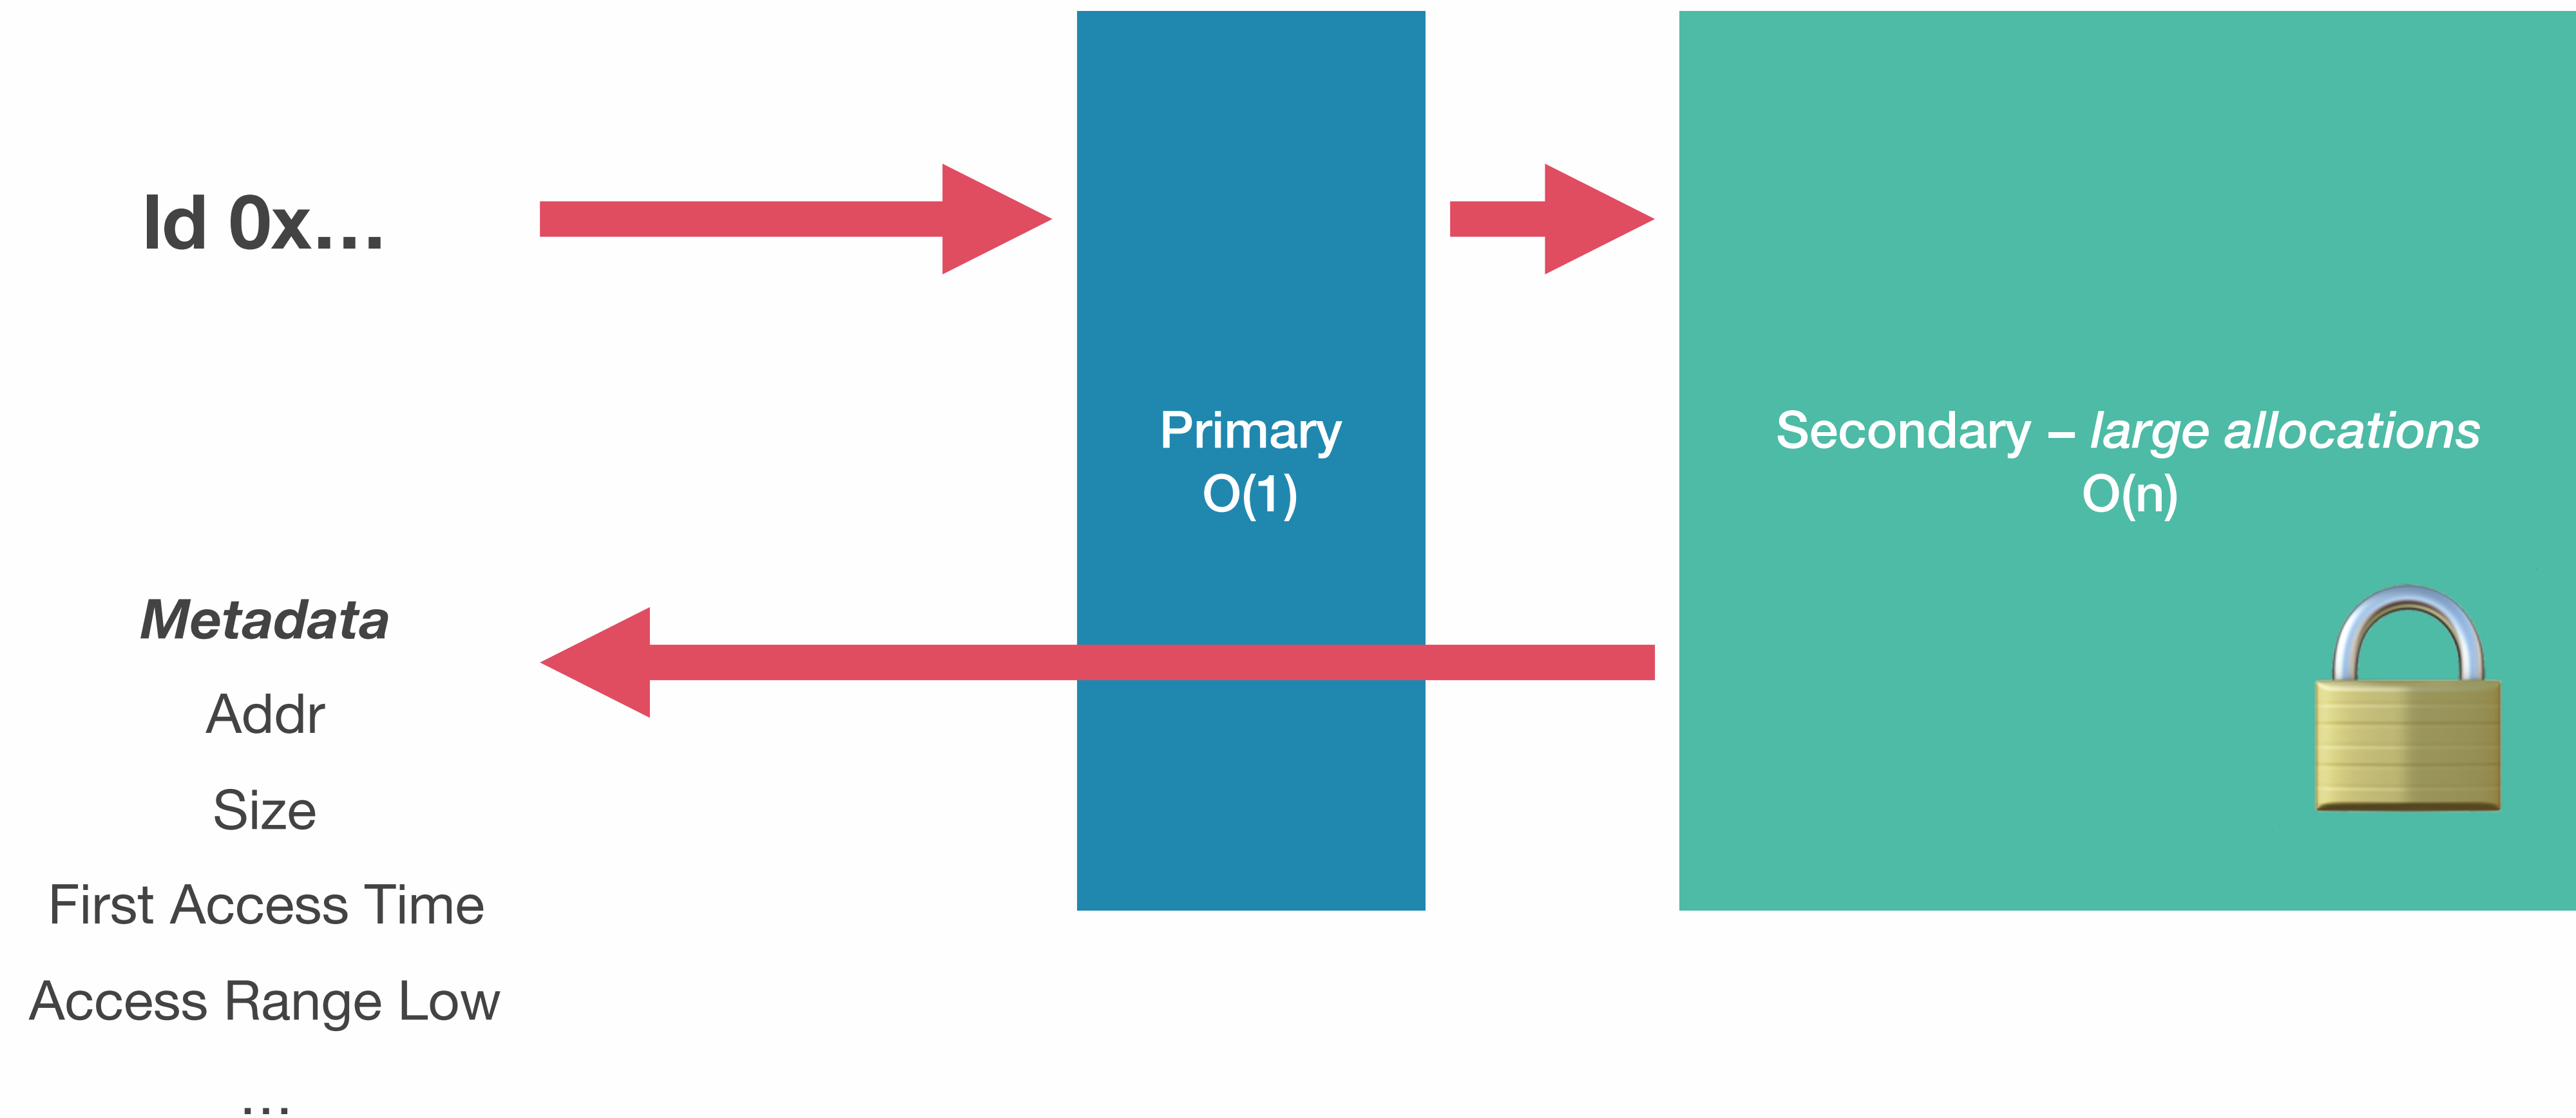
Time spent by address type



The Allocators



The Allocators



Issue with non-heap addresses



Issue with non-heap addresses



1. Allocators only know about heap



Issue with non-heap addresses



1. Allocators only know about heap
2. Traverse all allocations to discard them



Issue with non-heap addresses




1. Allocators only know about heap
2. Traverse all allocations to discard them
3. Takes a global lock



Issue with non-heap addresses



1. Allocators only know about heap
2. Traverse all allocations to discard them
3. Takes a global lock

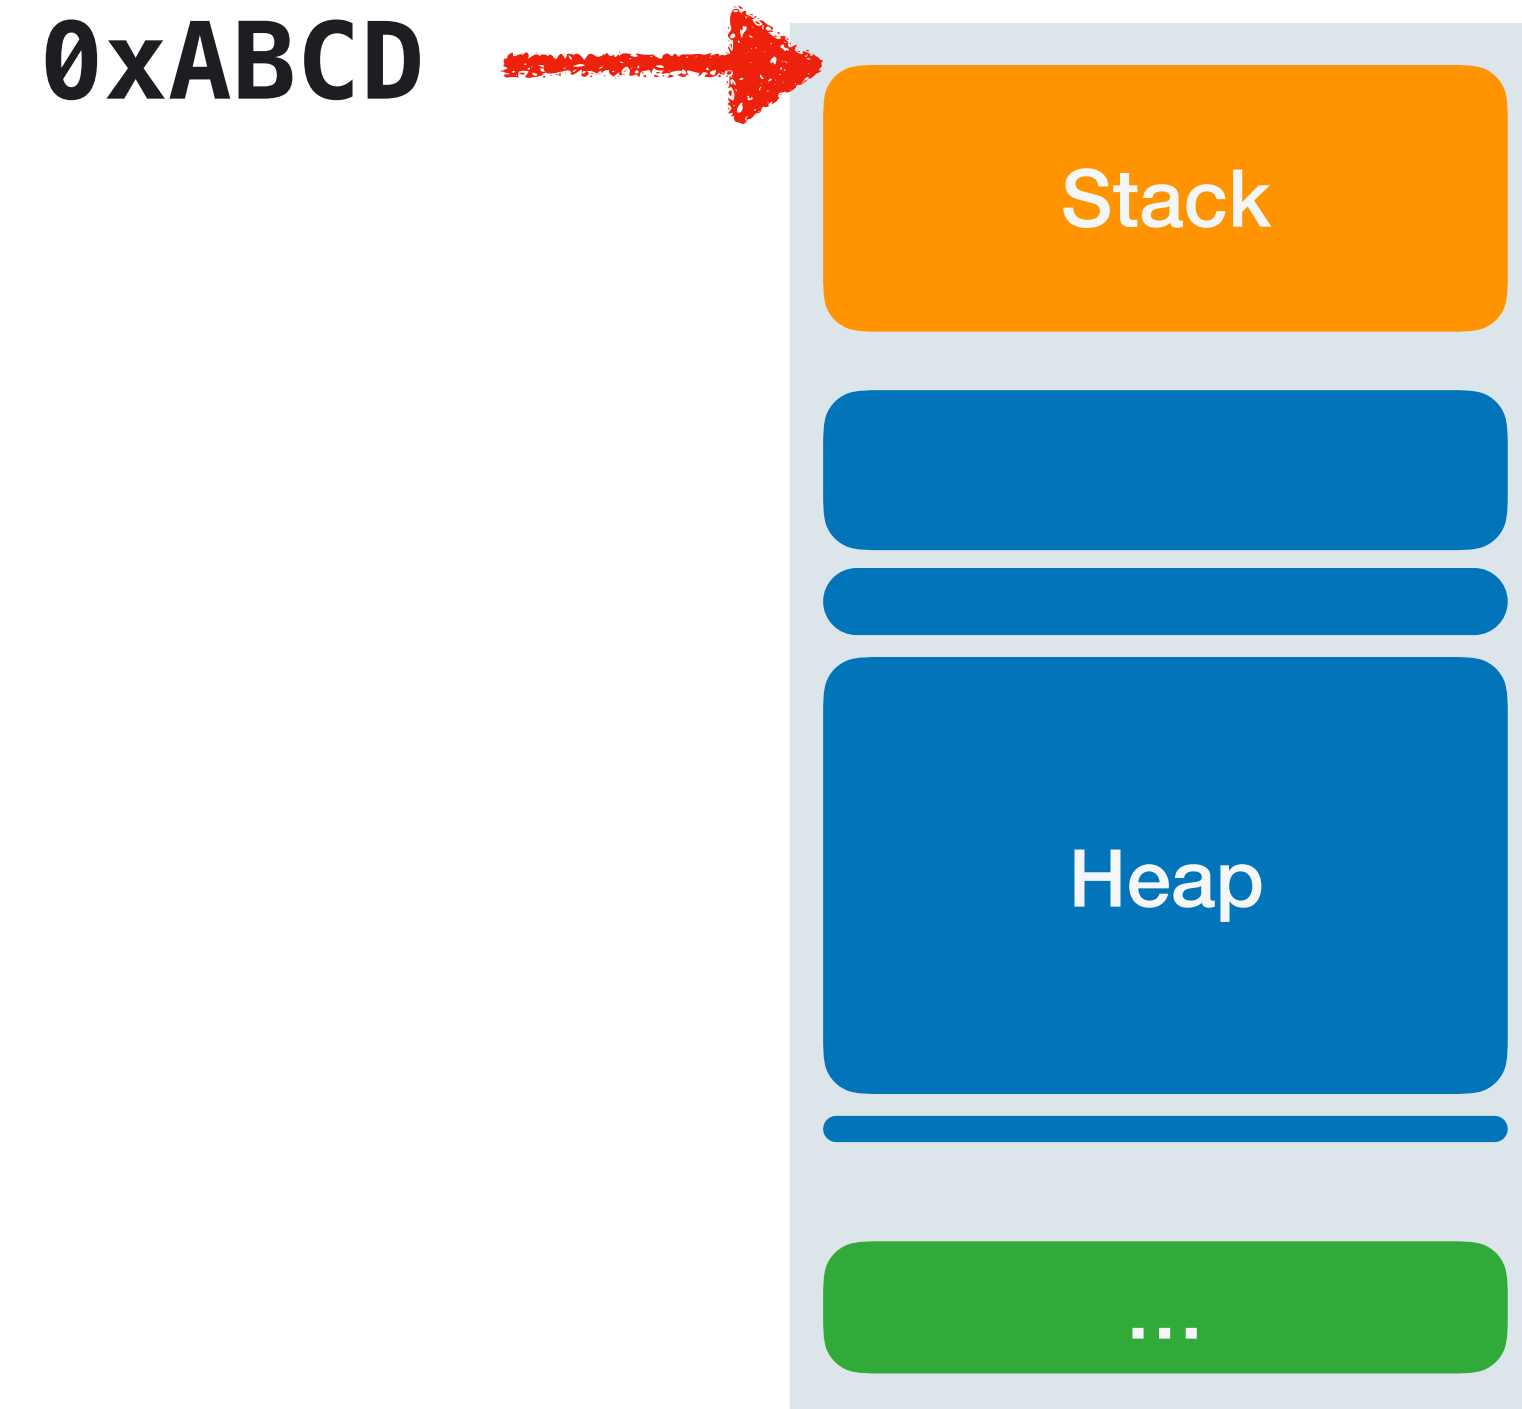


Run-Time Filter



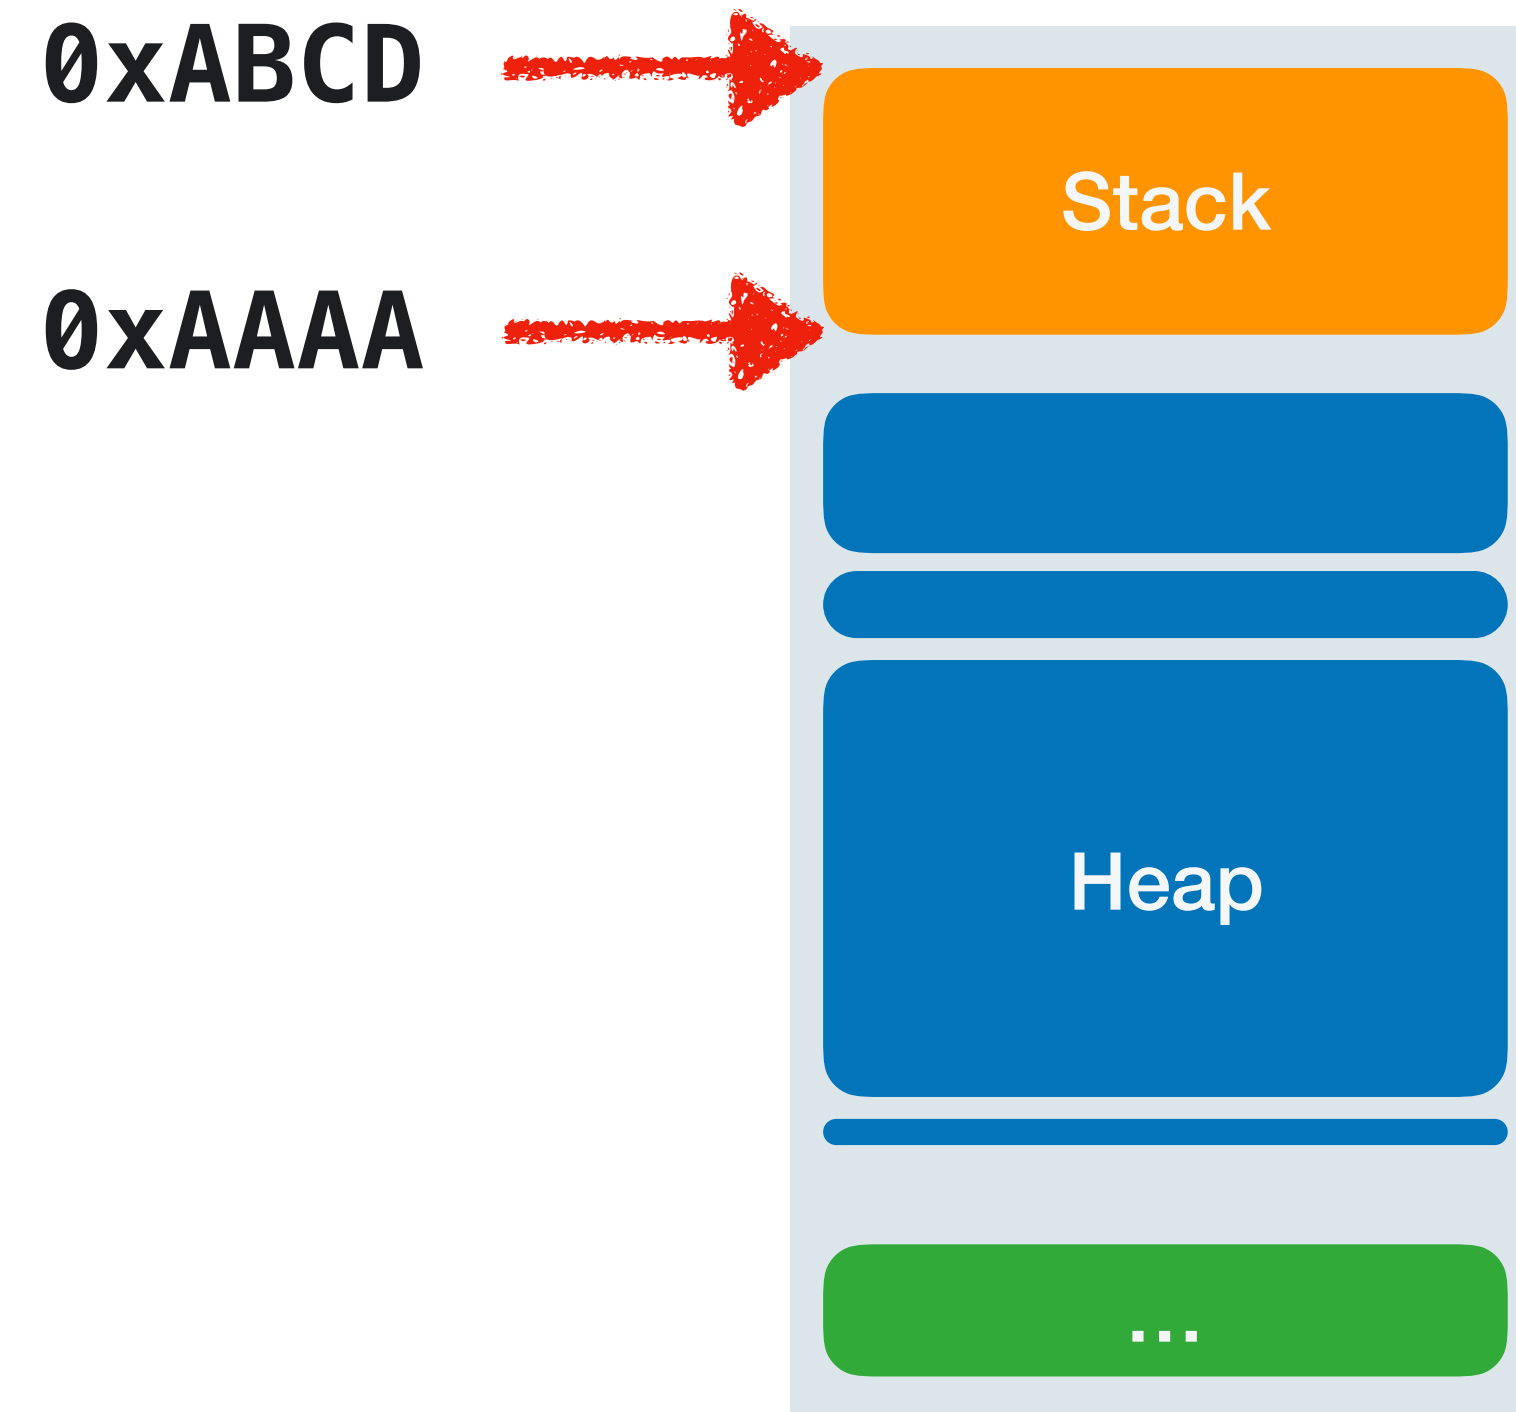
Run-Time Filter

1. *Thread start*: store stack top



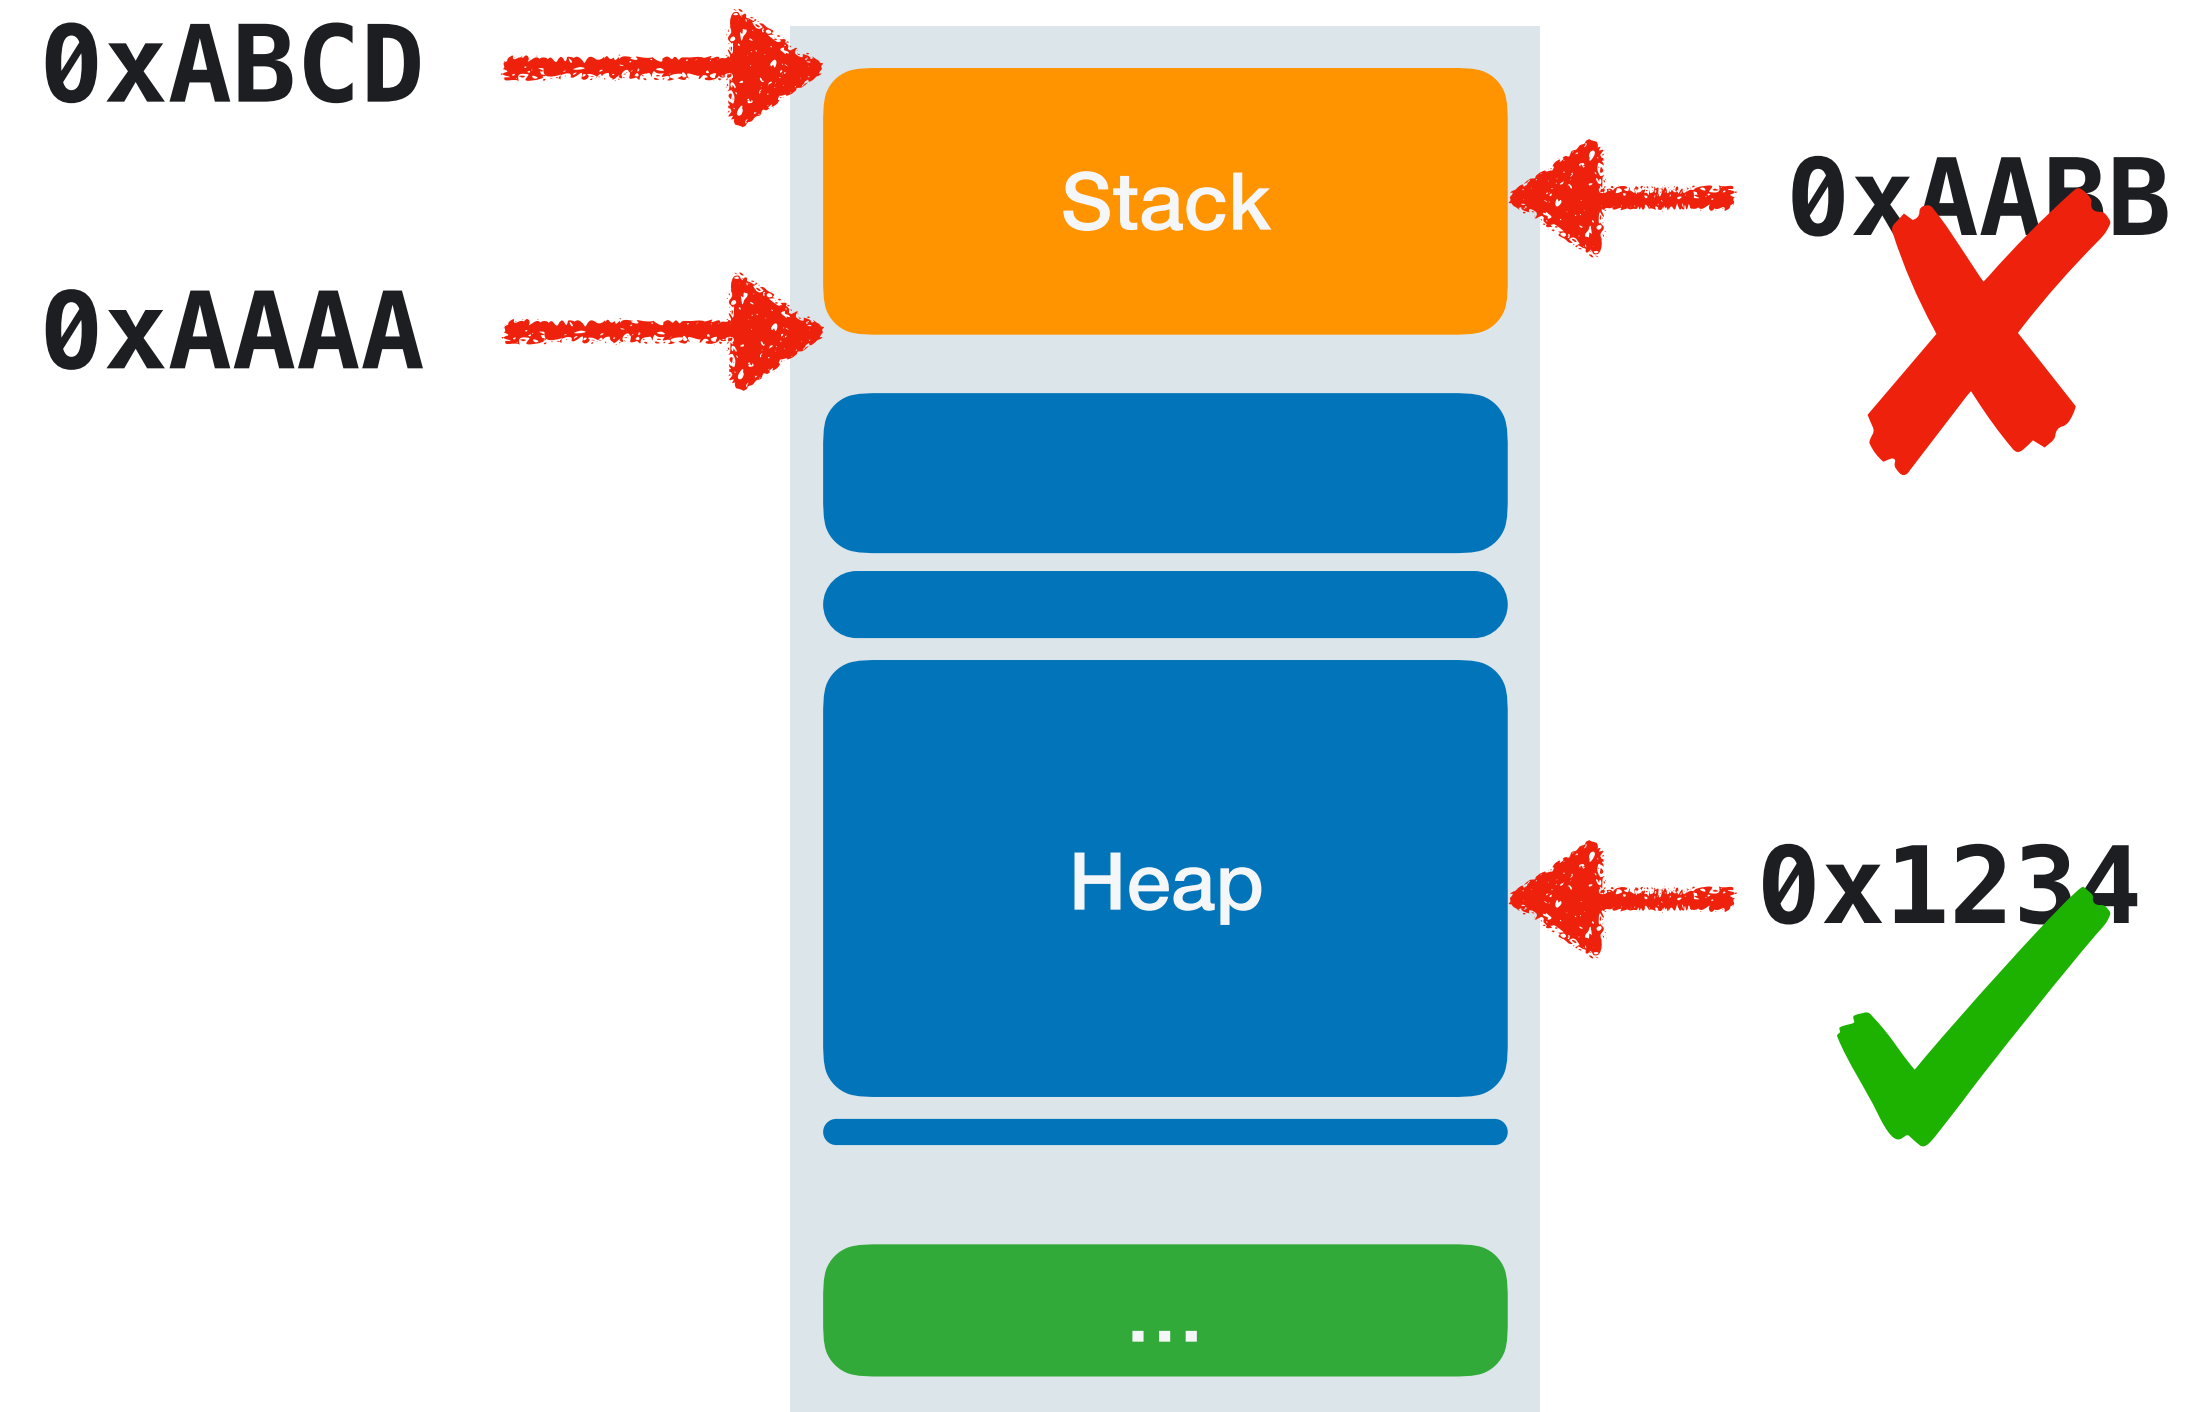
Run-Time Filter

1. *Thread start*: store stack top
2. Get current stack bottom




Run-Time Filter

1. *Thread start*: store stack top
2. Get current stack bottom
3. Discard if Addr. in this range



Time spent by address type



1,000x

slowdown due to Memoro's run-time

5x

slowdown due to Memoro's run-time

Memoro +

Run-Time Overhead ♦

Visualizer ♦


Open Challenges ♦

Memoro +

Run-Time Overhead ♦

Visualizer ♦

Open Challenges ♦



+


100,000

Stack Traces

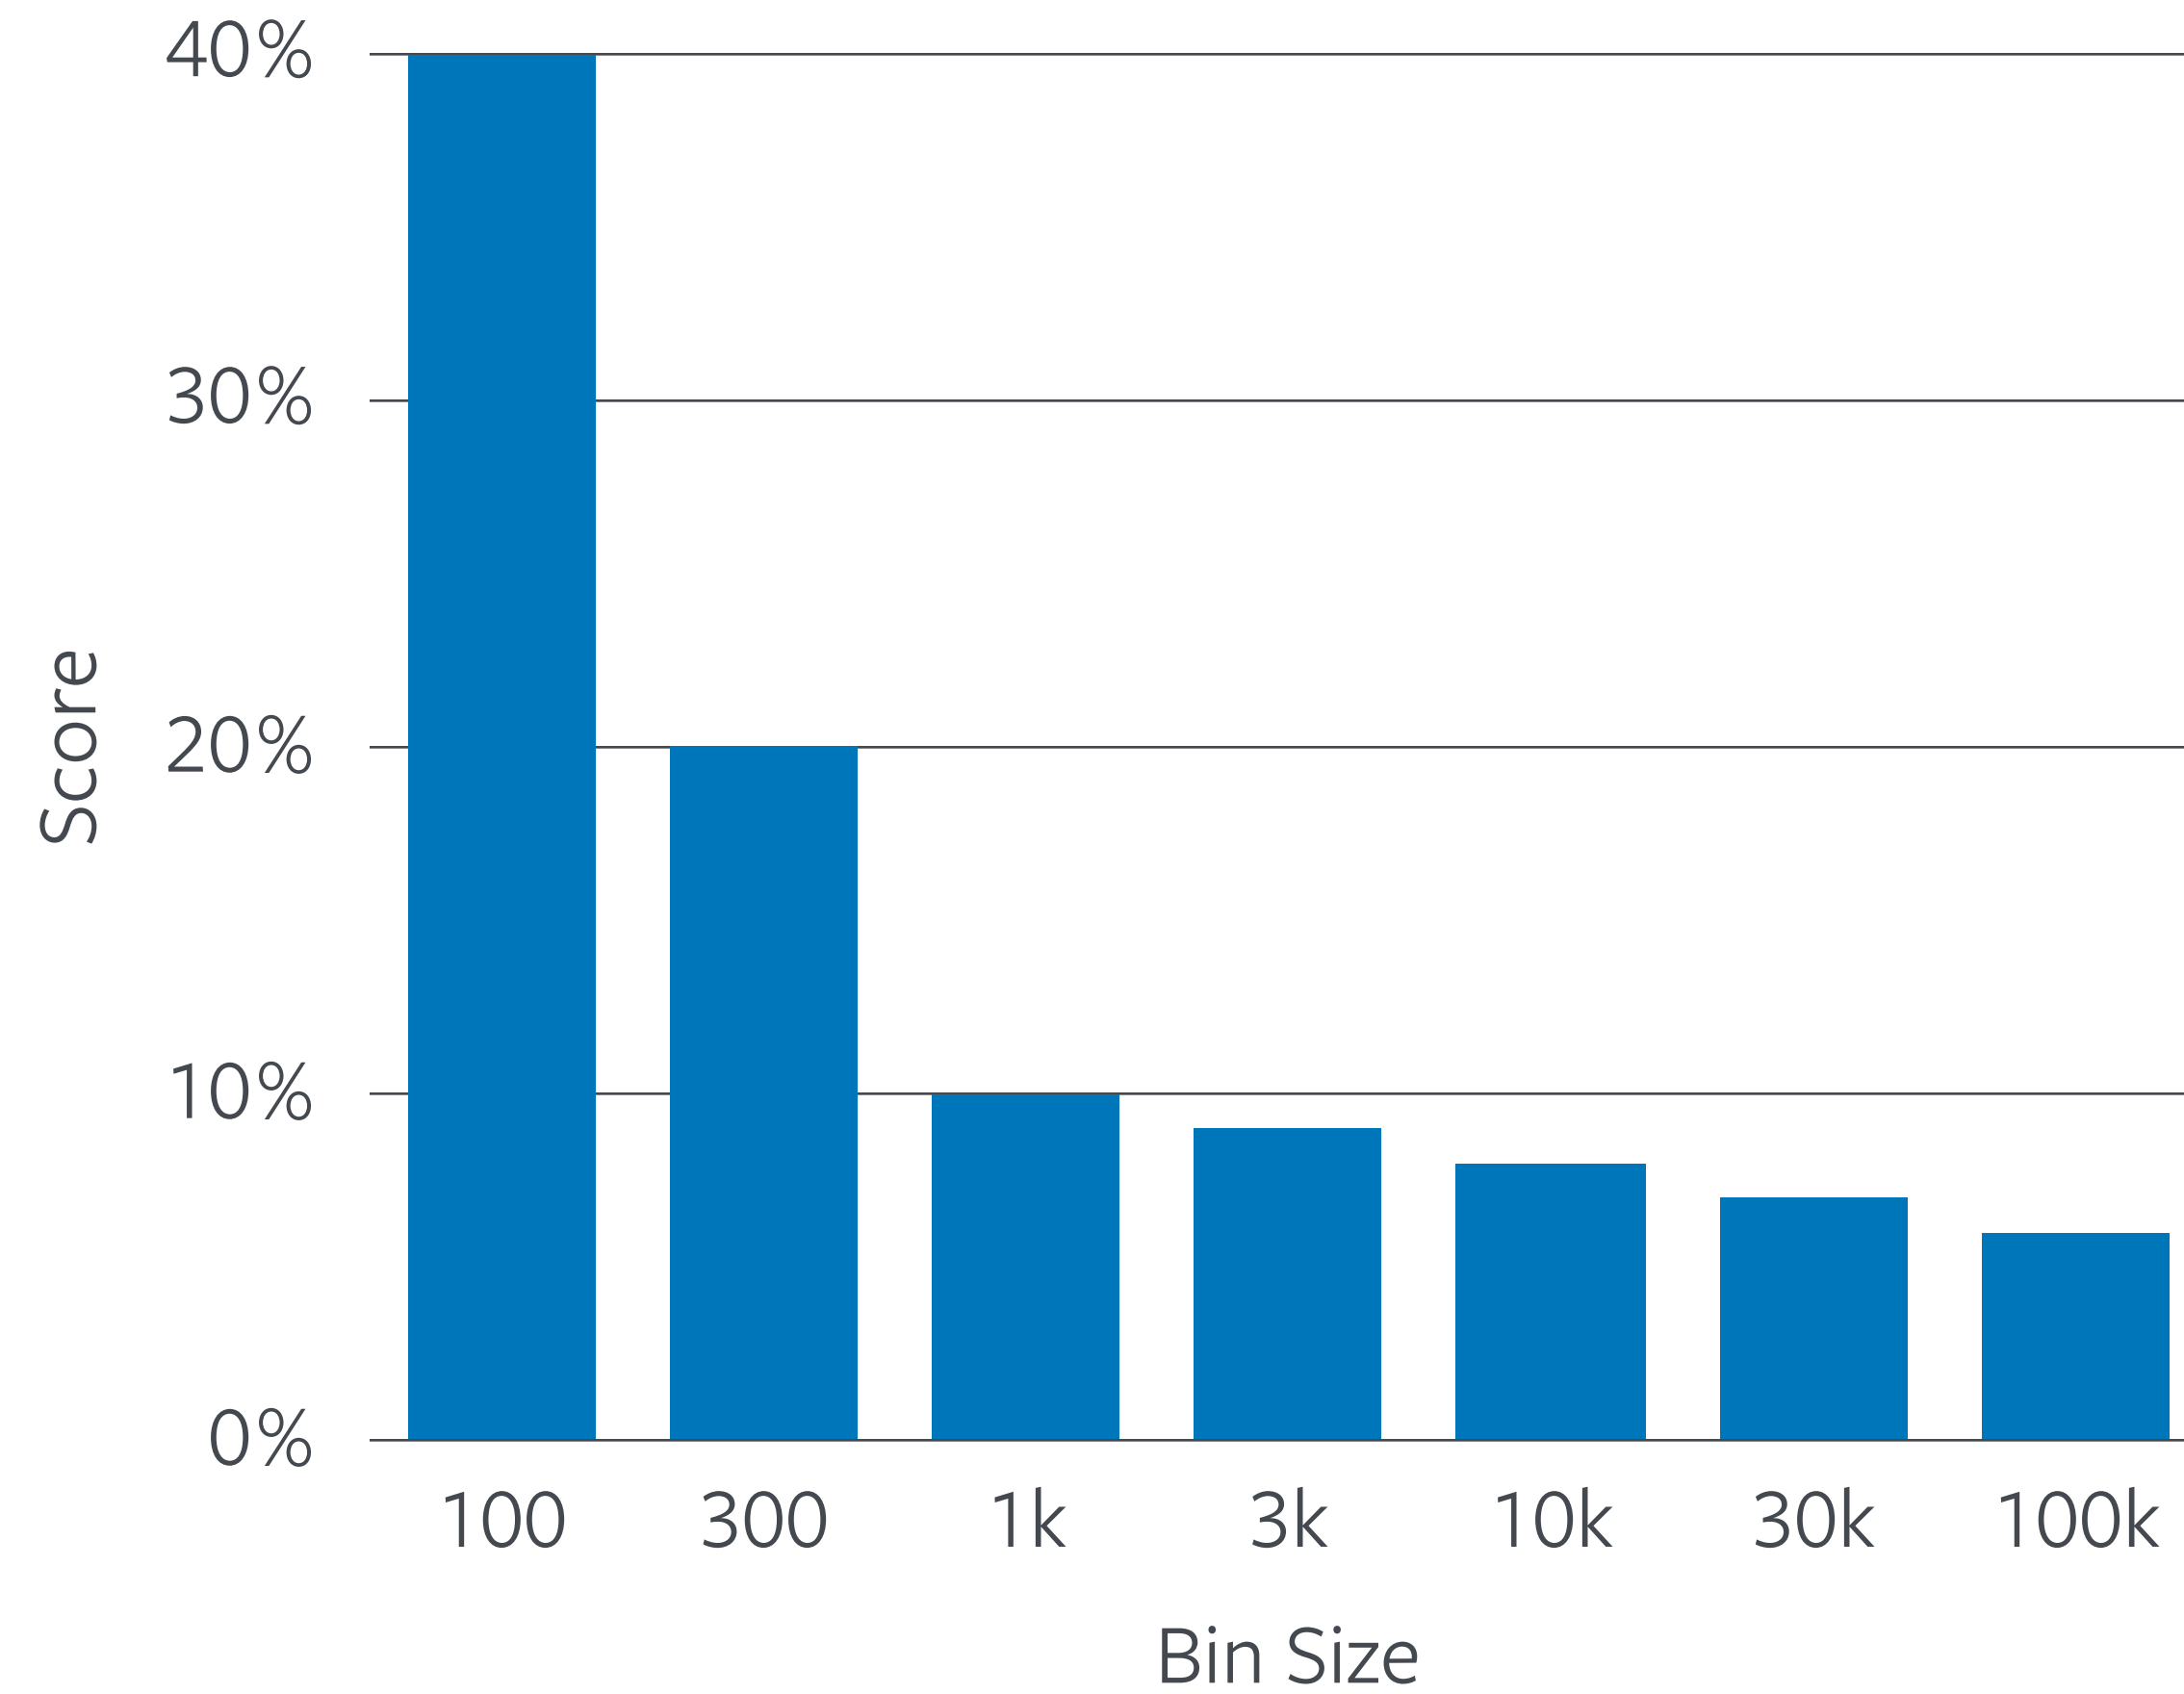
+

1B

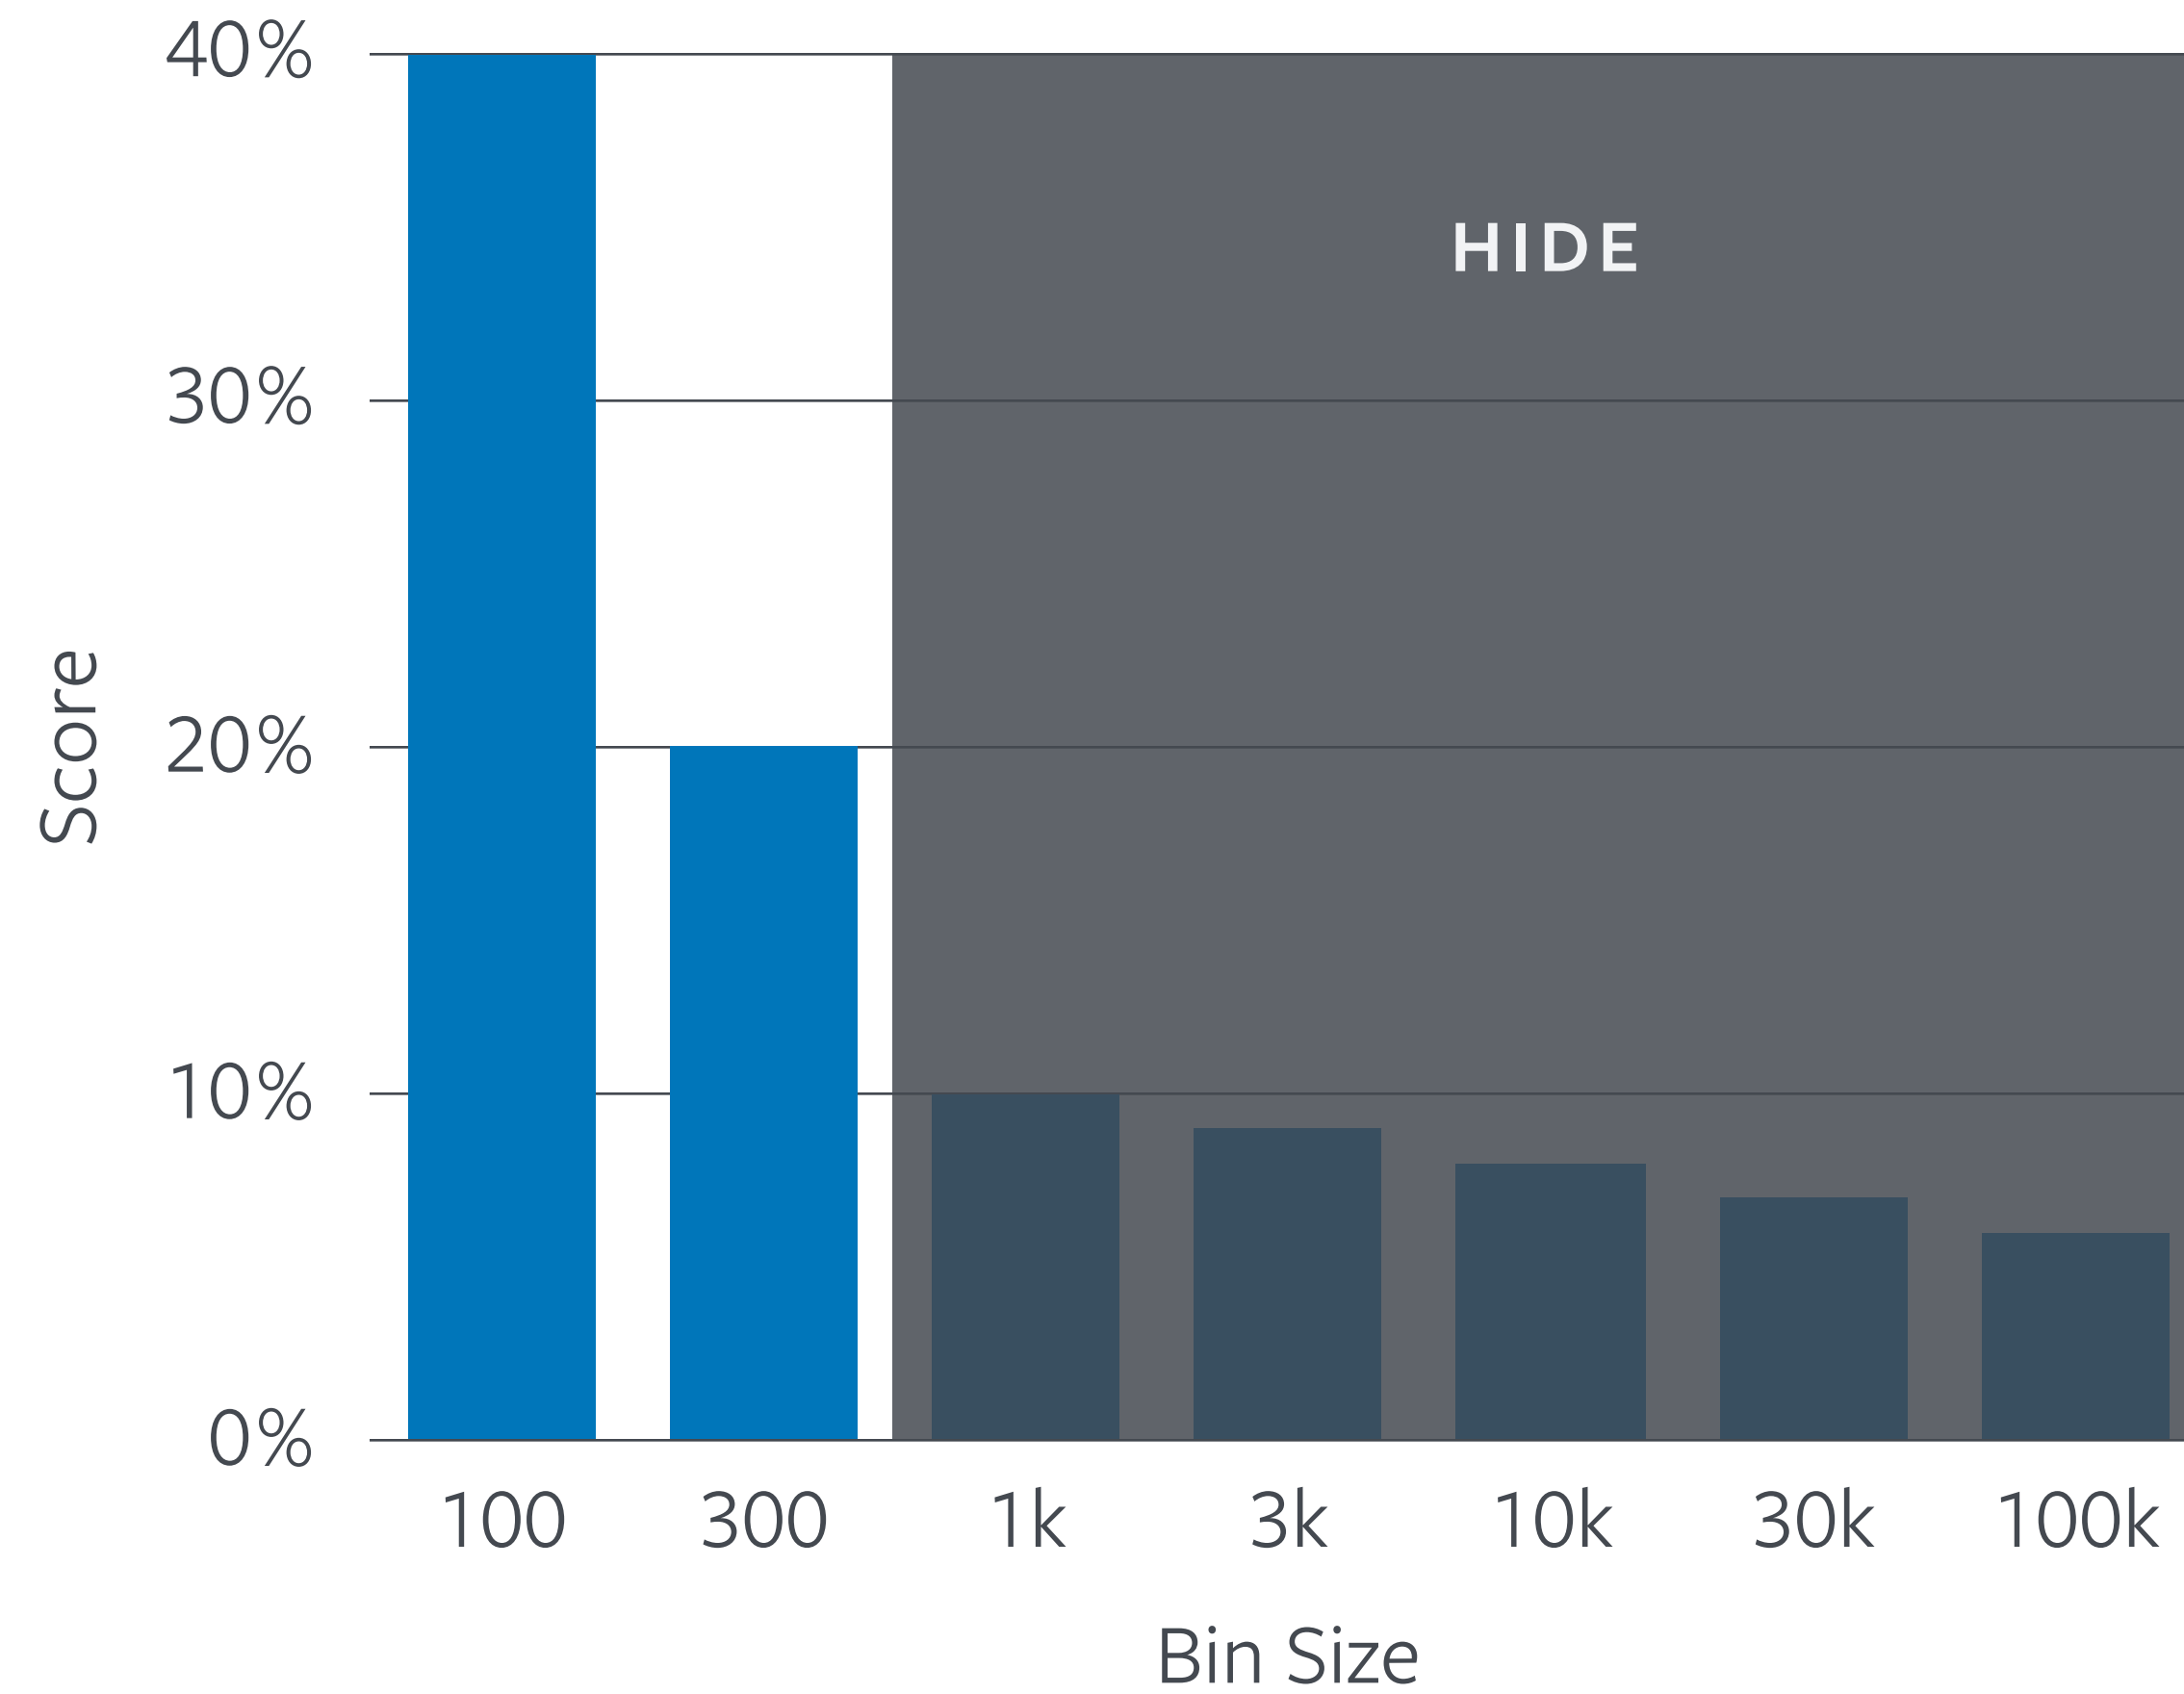
Allocations




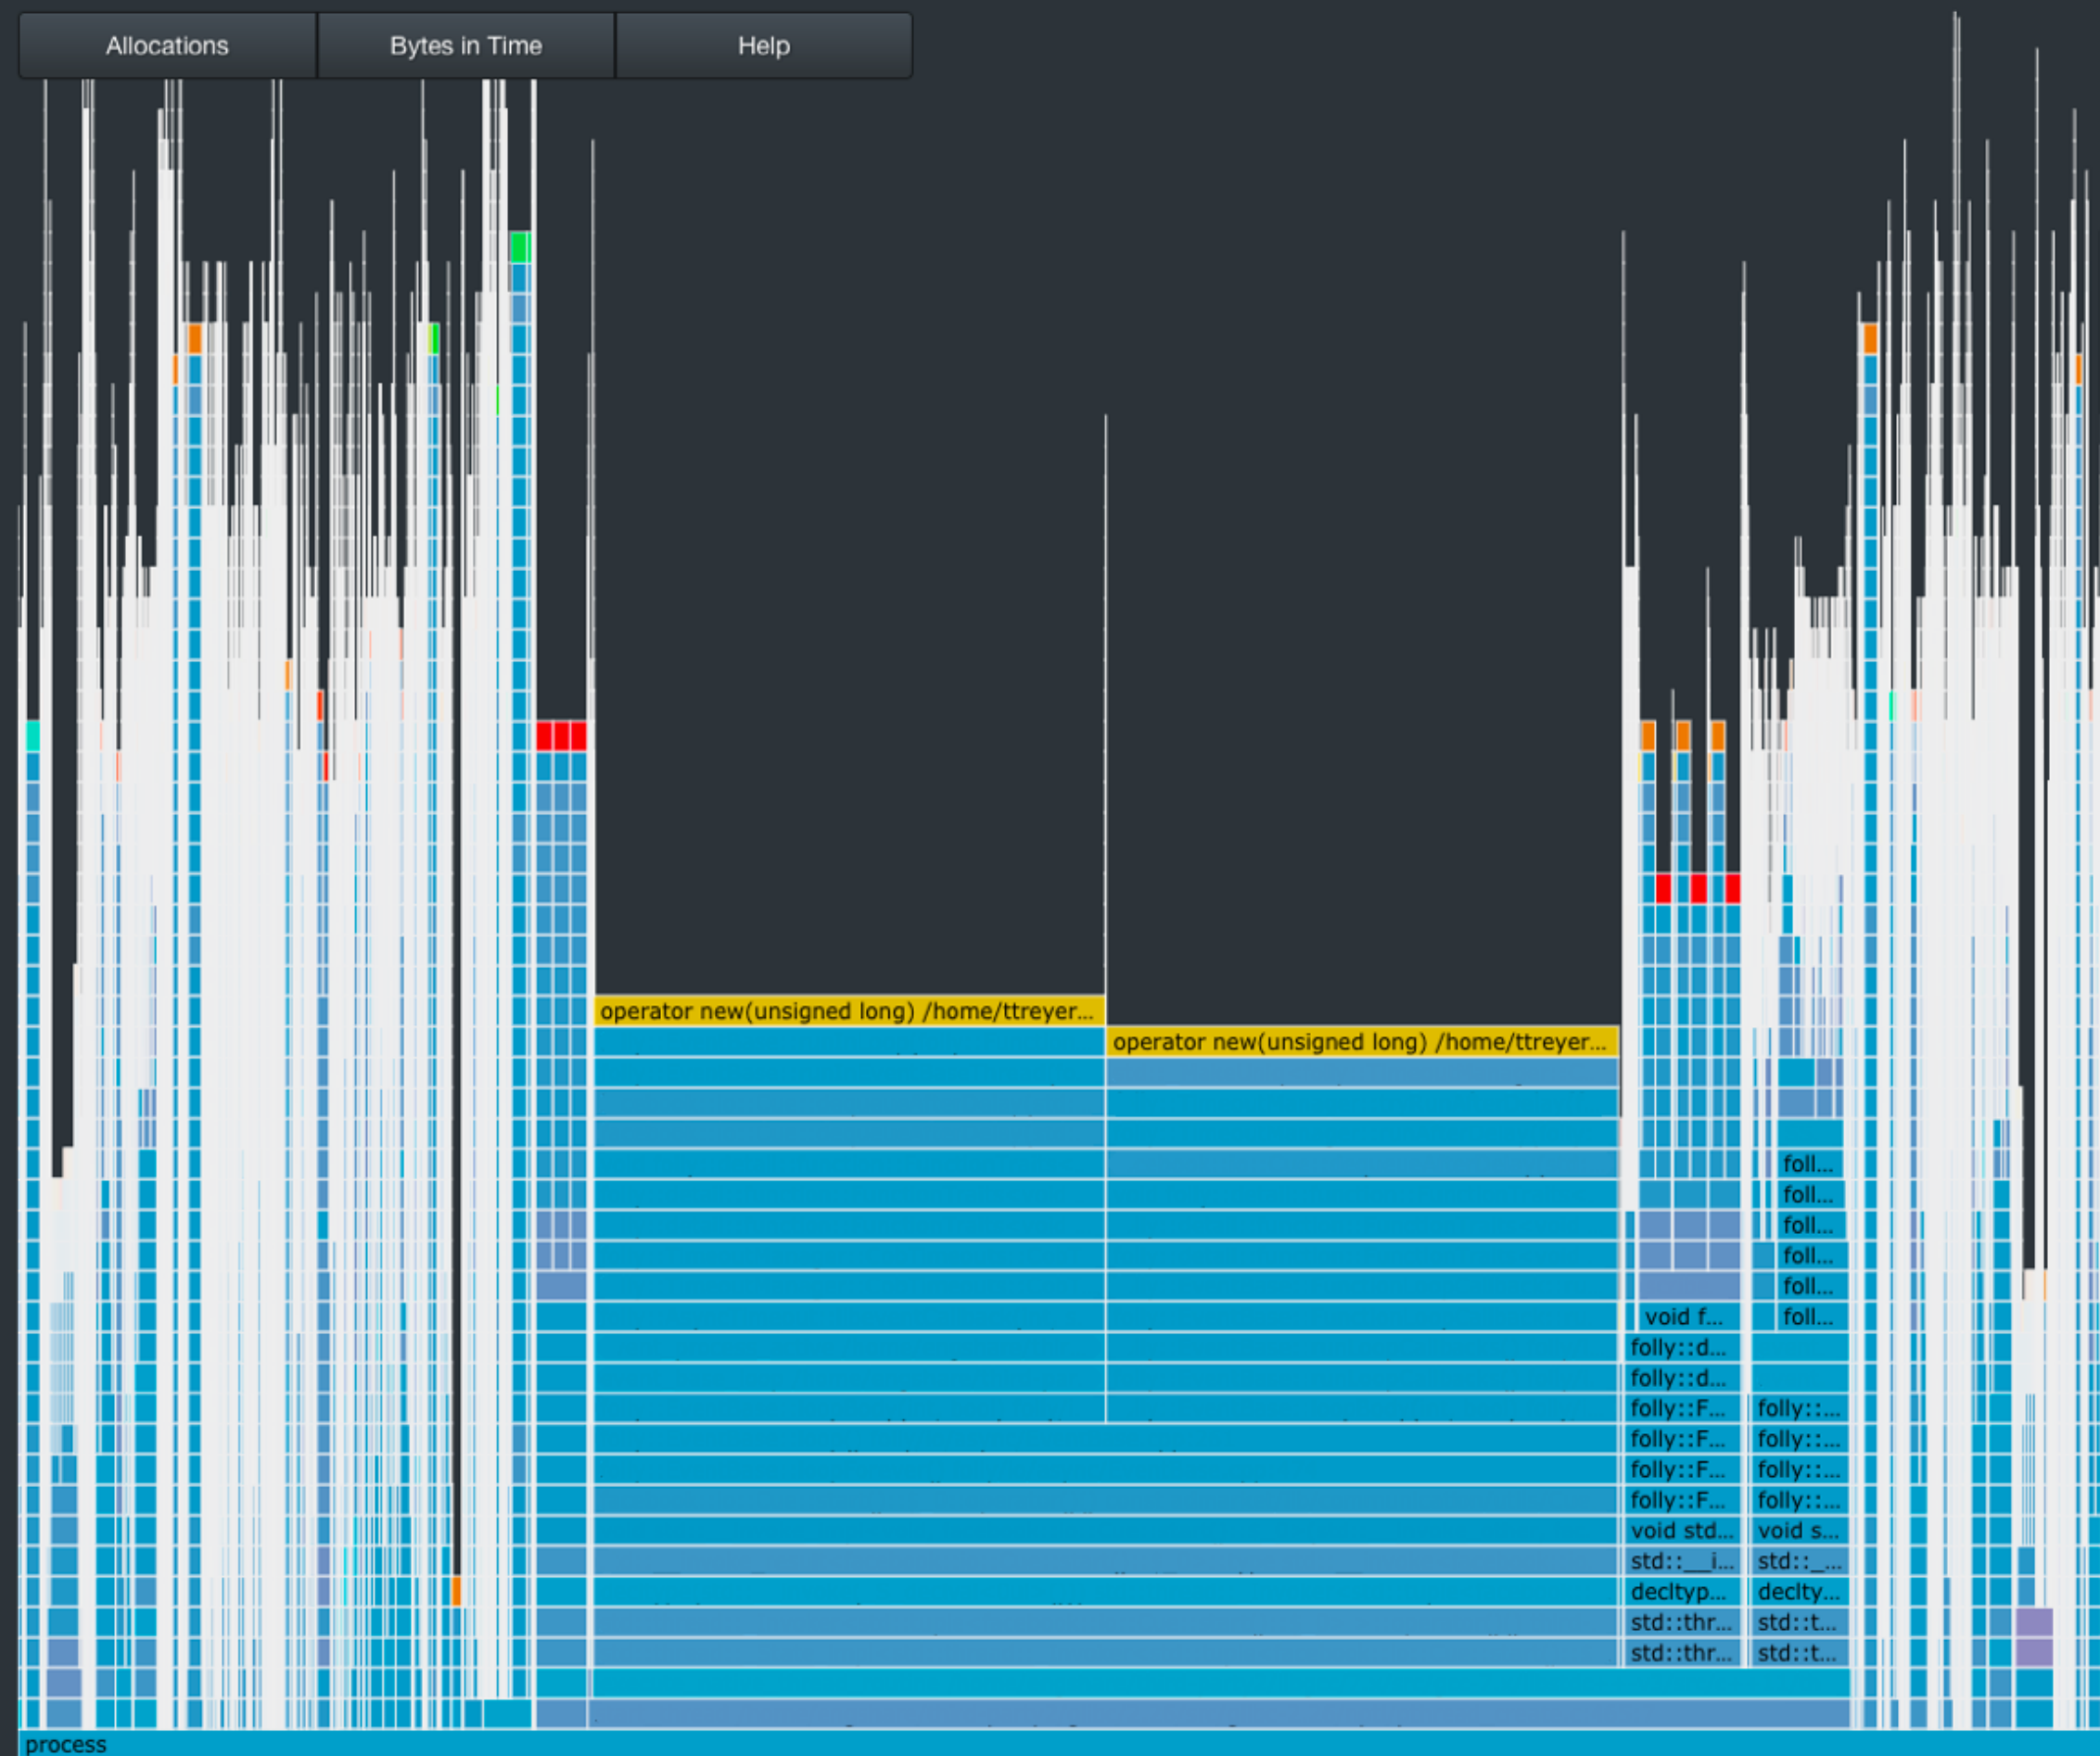
Truncate



Truncate








BEFORE




AFTER


Death by a thousand cuts




Death by a thousand cuts




Death by a thousand cuts




Death by a thousand cuts




Death by a thousand cuts




Death by a thousand cuts




Death by a thousand cuts





Death by a thousand cuts



Death by a thousand cuts



Death by a thousand cuts



Memoro + 

```
vector<BigT> getValues(  
    map<Id, BigT>& largeMap,  
    vector<Id>& keys) {  
  
    vector<BigT> values;  
    values.reserve(largeMap.size());  
  
    for (const auto& key: keys)  
        values.emplace_back(largeMap[key]);  
  
    return values;  
}
```

Demo

Memoro +

Run-Time Overhead ♦

Visualizer ♦

Open Challenges ♦

Dumping Profile

Your regular
service





Dumping Profile


Your regular
service



AtExit()





Dumping Profile



Dumping Profile

Facebook
service

AtExit()






Dumping Profile

Facebook
service



AtExit()

Dumping Profile

Facebook
service



AtExit()




Dumping Profile

Facebook
service




lldb

AtExit()



call
AtExit()




Dumping Profile

Facebook
service





lldb

~~AtExit()~~



call
AtExit()



Dumping Profile



- a. Signal to dump (SIGPROF)

Dumping Profile




- a. Signal to dump (SIGPROF)
- b. Ring buffer + Periodic write

Compile-Time Stack Analysis



Compile-Time Stack Analysis




Compile-Time Stack Analysis





Compile-Time Stack Analysis







Compile-Time Stack Analysis



foo()





bar()



ld/st


llvm::GetUnderlyingObject()

Compile-Time Stack Analysis



Thank you!

github.com/epfl-vlsc/memoro



 **Thierry Treyer**
Performance &
Capacity Intern

 **Mark Santaniello**
Performance &
Capacity Engineer

James Larus
EPFL IC School Dean

EPFL